HSA Core API Programmers Reference Manual

[c1]0.173

17 February 2013

Draft

[c1] *AMD*: 0.172

Contents

1

Introduction

1.1 Overview

Recent system-on-a-chip designs have integrated CPU, GPU, and other accelerator devices onto a single chip with a shared high-bandwidth memory system. In fact, these single-chip designs are now widely used in many computing markets including cellphones, tablets, personal computers, and game consoles. The Heterogeneous System Architecture (HSA) builds on the close physical integration of accelerators that is already occurring in the marketplace, and takes the next step by defining standards for uniting the accelerators architecturally. The HSA specifications includes requirements for virtual memory, memory coherency, architected dispatch mechanisms, and power-efficient signals. HSA refers to these accelerators as components. The system architecture defines a consistent base for building portable applications that access the power and performance benefits of the dedicated hsa components. Many of these components, including GPUs and DSPs, are capable and flexible processors that have been extended with special hardware for accelerating parallel code. Historically these devices have been difficult to program due to a need for specialized or proprietary programming languages. HSA aims to bring the benefits of these components to mainstream programming languages using similar or identical syntax to that which is provided for accessing multi-core CPUs.

In addition to the system architecture, HSAs Programmers Reference Guide defines HSAIL - a portable, low-level, compiler intermediate language designed for parallel computing.

Portable: The HSAIL language is an open-standard, supported by multiple vendors in the HSA Foundation, and is portable across vendors and product generations, so that applications which use HSAIL are guaranteed to run on future hardware that supports the HSAIL standard.

Low-level: HSAILs representation is just above the machine instruction set. Most optimizations including register allocation are intended to be performed by the compiler that generates

HSAIL. HSAIL code is translated to the host instruction set by a tool called the finalizer. Each component will provide its own implementation of the finalizer. The finalization step is intended to be lightweight, fast, and simple. Importantly, the finalizer step does not involve a complex compiler. Applications which contain HSAIL should not see different functional or performance behavior from new finalizer versions that might be deployed in the field after the application ships.

Designed for parallel computing: HSAIL is intended to represent the parallel sections of an application. It complements but does not replace the host code host code still exists and is used for the serial portion of the application. HSAIL represents a single lane of execution, and the parallelism is expressed in the grid dimensions that are specified when the HSAIL kernel is dispatched to a target component. In this way, HSAIL does not encode a specific vector width and can be used to represent a variety of different parallel computing devices.

For more information on HSAIL, refer to the HSA Programmers Reference Guide.

The final piece of the puzzle is the HSA Core Runtime API. The core runtime is a thin, user-mode API that provides the interfaces necessary for the host to launch compute kernels to the available components. This document describes the architecture and APIs for the HSA Core Runtime. Key sections of the runtime API include:

- Error handling

- Runtime initialization (open/close)

- Topology and Component Discovery

- Signals and Synchronization

- Architected Dispatch

- Memory Management

In summary, there are three specifications provided by the HSA Foundation:

HSA System Architecture Requirements: Architectural foundation for how HSA components share memory and communicate work requests.

HSA Programmers Reference Manual: Describes HSAIL, a low-level, portable compiler IR appropriate for use as compiler intermediate language.

HSA Runtime Specification: This document. Describes the host-side API for controlling the launch of HSAIL kernels.

The remainder of this document describes the HSA software architecture and execution model, and includes functional descriptions for all of the HSA APIs and associated data structures.

Figure 1.1 shows how the HSA Core Runtime fits into a typical software architecture stack.
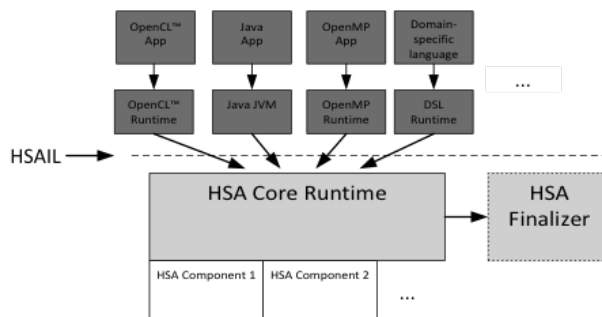
Figure 1.1: HSA Software Architecture

At the top of the stack is a programming model such as OpenCL, Java, OpenMP, or a domain-specific language (DSL). The programming model must include some way to indicate a parallel region that can be accelerated. For example, OpenCL has calls to `clEnqueueNDRangeKernel` with associated kernels and grid ranges. Java has the stream and lambda APIs, which provide support for both multi-core CPUs and HSA Components. OpenMP contains OMP pragmas that mark parallel for loops and control other aspects of the parallel implementation. Other programming models can also build on this same infrastructure.

The language compiler is responsible for generating HSAIL code for the parallel regions. HSA supports several options for when HSAIL is generated and finalized. One possibility is that the HSAIL is generated by a high-level compiler and then embedded in the application binary. In this case, the finalizer is run when the application loads and will convert the HSAIL to machine code for the target machine. Another option is that the HSAIL is finalized when the applications is built, or the machine cod The HSA Finalizer is an optional component of the HSA Core Runtime, which may reduce the footprint of the HSA software on systems where the finalization is done before runtime.

Each language also includes a language runtime that connects the language implementation to the HSA Core Runtime. When the language compiler generates code for a parallel region, it will include calls to the HSA Runtime to set up and dispatch the work to the HSA Component. The language runtime is also responsible for initializing HSA, and may utilize other HSA core runtime features as well.

The API for the HSA core runtime is standard across all HSA vendors, such that languages which use the HSA runtime can run on the different vendors that support the API. Each vendor is responsible for supplying their own implementation which supports the hsa component(s) in the vendors platform. HSA does not provide a mechanism to combine runtimes from different vendors; instead vendors must provide a single runtime which supports all the components in the platform. The implementation of the HSA Runtime may include kernel-level components (typical for hardware components) or may be entirely user-space (simulators or CPU implementations).

1.2   Infrastructure and Execution Flow

1.3   Execution Model

Host/HSAIL brief intro that describes host running host code include HSA runtime, components running HSAIL.

1.3.1   Architected Dispatch

Core runtime exposes several details of the HSA hardware, including architected dispatches and support for execution control. The overall goal of the core runtime design is to provide a high-performance dispatch mechanism that is portable across multiple HSA vendor architectures. Two vendors with the same host ISA but different HSA-compliant GPUs will be able to run the same unmodified binary, because they support the HSA-architected AQL interface and supply a library that implements the architected core runtime API.

In order for user-level applications to use the HSA system and HSA components, they need to write HSAIL programs and compile and execute these programs using user mode queues and AQL commands. The HSA Programmers Reference Manual (PRM) defines HSAIL Virtual ISA and Programming Model, serves as a Compiler Writers Guide, and defines Object Format (BRIG). The HSA runtime helps setup the execution via API calls and data structures to support architected features.

The HSA core runtime realizes architected dispatch. Architected dispatch is the key feature in an HSA system that enables a user-level application to directly issue commands to the HSA Component hardware. Architected dispatch differentiates it from other higher-level runtime systems and programming models: other runtime systems provide software APIs for setting arguments and launching kernels, while HSA architects these at the hardware and specification level. The critical path of the dispatch mechanism is architected at the HSA hardware level and can be done with regular memory operations and runtime provided wrapper API. Fundamentally, the user creates user mode queues and an AQL Packet in memory, and then signals the HSA component to begin executing the packet using light weight operations (which may be wrapped with API calls).

This section describes various features core runtime provides to support architected dispatch as steps that a user needs to take to utilize runtime.

1.3.2   Initial Setup

One of the first steps in the setup is that of device discovery. Device discovery is performed at the initialization of the core runtime and information is made available to the user as data structures. Section 2.3 describes these structures. The next step in the setup is creation of the component queues. Queues are an HSA architected mechanism to submit work to the HSA component HW. The interfaces for queue creation are defined in Section 2.5. Different components may provide implementation-specific code under the core API for these functions. HSA runtime also includes mechanisms to provide implementation-specific data as part of the dispatch, provided such data can be computed at compile time.

1.3.3   Compilation Flow

Once an HSAIL program is written or generated by a higher-level compilation step, it needs to be *assembled* to generate a BRIG. BRIG is the HSAIL object format and is specified in the PRM. HSA runtime defines API call to compile the BRIG and generate a code object that has sufficient information to execute the user program. The details of this compilation process and symbol resolution are discussed in Section A.

1.3.4   Execution of Kernel

The Systems Architecture Requirements (SAR) document specifies the structure of the *packets* (i.e. commands) that can be placed on the HSA user mode queues for the component HW to execute them. The format of the packets is architected and they are referred to as Architected Queuing Language (AQL) packets. One of the types of AQL packets is a dispatch AQL packet. The user can now create an AQL packet and initialize it with the code object obtained from the finalization step, including the allocation of memory to hold the kernel arguments and the spill/arg/private memory. The interface for kernel arguments between the runtime and the kernel ISA (instruction set architecture) is also architected at the HSA level. This is covered in the HSAIL ABI (this is discussed in Section **??**), which specifies the in-memory layout of the kernarg segment. Users can determine the layout of the kernarg memory segment at compile time merely by examining the signature of the HSAIL function. The finalizer is required to support this ABI and thus there is no need for runtime metadata to specify the position or format of arguments. This step can be done once for each AQL packet creation.

Optimized implementations can cache the result of this step and re-use the AQL packet for subsequent launches. Care must be taken to ensure that the AQL Dispatch packet (and the associated kernel and spill/arg/private memory) is not re-used before the launch completes. For simple cases, (that is, a single-thread, synchronous launch, the AQL dispatch packet(s) can be declared as a static variable and initialized at the same time the code is finalized. More advanced cases can create and track several AQL Dispatch packet(s) for a single kernel code object.

HSA HW defines a packet process for processing these packets and a doorbell mechanism to inform the packet processing HW that packets have been written into the queue. The Core runtime defines a structure and update API to inform the HW that the dispatch packet has been written to the queue. Different packet formats and states of a packet are discussed in Section 2.6. Section 2.5 discusses the queue creation and various states the queue can be in, once it is created.

Once the packet is written and the HW is informed by way of the doorbell, the execution can start. The execution happens asynchronously. The user is free to write more packets for executing other kernels in the queue. This activity can overlap the actual execution of the kernel.

1.3.5   Determining Kernel Completion

HSA SAR defines signals as a mechanism for communication between different parts of a HSA system. Signals are defined as opaque objects in the HSA core runtime and APIs have been defined to send a value to the signal and wait for a value at the signal, Section 2.4 discusses signals in detail. The AQL dispatch packet has a provision for the user to pass in an opaque signal. When the HSA Component HW observes a valid signal in the AQL packet, it sends a value to this signal when execution of the kernel is complete (success or error). The user can wait on this signal to determine kernel completion. Errors and their meaning are discussed in Section 2.1.

```
int main(int argc, char **argv)
{
unsigned int count;
hsa_brig_direcve_offset_t code_directive_offset = atoi(argv[2]);
hsa_control_directives_t control = atoi(argv[3]);
uint64_t kernel_input = (uint64_t)atoi(argv[4]);
const hsa_agent_t *component = NULL;
static hsa_aql_dispatch_packet_t dispatch_packet;
hsa_runtime_context_t *runtime_context;
hsa_status_t status;
hsa_queue_t *queue;
hsa_brig_t *brig = (hsa_brig_t *)argv[1];
hsa_code_object_t *code_obj;
hsa_symbol_map_t *symbol_map;
hsa_debug_info_t *debug_info;
hsa_signal_handle_t signal;
hsa_signal_value_t sigval;
hsa_topology_header_t *header;
hsa_platform_t *platform;

uint64_t index;

  /**** this part is the setup for a simple dispatch ****/
  status = hsa_open(&runtime_context);
  assert(status == HSA_STATUS_SUCCESS);

  hsa_context_acquire(&runtime_context);

  hsa_close(&runtime_context);

  status = hsa_topology_table_create(&header);
  if(status != HSA_STATUS_SUCCESS) {
    assert(HSA_STATUS_SUCCESS == hsa_close());
    exit(1);
  }

  platform = (hsa_platform_t *)(header->topology_table_base);

  component = (hsa_agent_t *)((char *)(header->topology_table_base)
    + platform->agent_offset_list[0]);

  status = hsa_queue_create(component, 1024, 0, NONE,
    runtime_context, &queue, NULL);
  if(status != HSA_STATUS_SUCCESS) {
    assert(HSA_STATUS_SUCCESS == hsa_close());
    exit(1);
```

```
  }

  /**** this is the compilation part where the brig is finalized
     ****/
  status = hsa_finalize_brig(component, brig, 1,
     &code_directive_offset, &control, NULL, &code_obj, &debug_info,
     &symbol_map);
  if(status != HSA_STATUS_SUCCESS || symbol_map != NULL) {
    assert(HSA_STATUS_SUCCESS == hsa_queue_destroy(queue));
    assert(HSA_STATUS_SUCCESS == hsa_close());
    exit(1);
  }

  /**** a signal is created for completion detection ****/
  sigval.value64 = 0;
  status = hsa_signal_create(sigval, &signal, runtime_context);
  if(status != HSA_STATUS_SUCCESS) {
    assert(HSA_STATUS_SUCCESS == hsa_queue_destroy(queue));
    assert(HSA_STATUS_SUCCESS == hsa_close());
    exit(1);
  }

  /**** the AQL packet is setup here for the simple kernel ****/
  dispatch_packet.header.format = 2;
  dispatch_packet.header.barrier = 1;
  dispatch_packet.header.acquire_fence_scope = 2;
  dispatch_packet.header.release_fence_scope = 2;
  dispatch_packet.header.dimensions = 1;
  dispatch_packet.workgroup_size_x = 256;
  dispatch_packet.grid_size_x = 256;
  dispatch_packet.kernel_object_address = (uint64_t)code_obj;
  dispatch_packet.kernarg_address = (uint64_t)(&kernel_input);
  dispatch_packet.completion_signal = signal;

  memcpy(queue->base_address, (void *)&dispatch_packet,
     sizeof(dispatch_packet));

  /**** packet processor is informed that a packet is on the queue
     ****/
  index = hsa_queue_set_write_index(queue, 1);
  if(index != 0){
    assert(HSA_STATUS_SUCCESS == hsa_queue_destroy(queue));
    assert(HSA_STATUS_SUCCESS == hsa_close());
    exit(1);
  }
  sigval.value64 = 1;
```

```
  status = hsa_signal_send_release(queue->doorbell_signal, sigval);
  if(status != HSA_STATUS_SUCCESS) {
    assert(HSA_STATUS_SUCCESS == hsa_queue_destroy(queue));
    assert(HSA_STATUS_SUCCESS == hsa_close());
    exit(1);
  }

  /**** await completion ****/
  do {
    status = hsa_signal_wait_acquire(signal,HSA_EQUALS, sigval,
        &sigval);
  } while(status == HSA_STATUS_INFO_SIGNAL_TIMEOUT);
  if(status != HSA_STATUS_SUCCESS) {
    assert(HSA_STATUS_SUCCESS == hsa_queue_destroy(queue));
    assert(HSA_STATUS_SUCCESS == hsa_close());
    exit(1);
  }

  printf("\nkernel successfully executed, value %llu\n",
      kernel_input);

  /**** close up and destroy queue, close the runtime ****/
  status = hsa_queue_destroy(queue);
  if(status != HSA_STATUS_SUCCESS) {
    assert(HSA_STATUS_SUCCESS == hsa_close());
    exit(1);
  }
  status = hsa_close();
  assert(status == HSA_STATUS_SUCCESS);
  return 0;
}
```

2

HSA Core API Specification and Description

This chapter describes HSA Core runtime API by their functional area. Note that except for any setter/getter API, the remainder of the core runtime API may be considered thread-safe. Both the signal update and the queue index update API are setter/getter API and define scope an synchronization that applies to the updates and operate on structure elements.

Several operating systems allow functions to be executed when a DLL or a shared library is loaded (e.g. DLL main in Windows and GCC *constructor/destructor* attributes that allow functions to be executed prior to main in several operating systems). Whether or not the HSA runtime functions are allowed to be invoked in such fashion may be implementation specific and are outside the scope of this specification.

Similarly, any header files distributed by the HSA foundation for this spec may contain calling convention specific prefixes such as ˍˍcdecl or ˍˍstdcall. Such calling conventions are again invocation, usage and implementation specific. Hence, the calling convention specific prefix definition is outside the scope of API definition.

2.1   Synchronous and Asynchronous Errors and Asynchronous Notification

Error handling in the core runtime can broadly be classified into two categories: synchronous error handling and asynchronous error/notification handling.

Synchronous errors are always reported when the call returns. They indicate if the API returned a success or an error.

Asynchronous errors can occur due to various reasons:

(i) Activity in packet processor, executing kernels, their actions and memory accesses. If an error is detected during execution of a kernel, the completion signal (if present) will be signaled with an error indication value.

(ii) To provide *information/warning* (not as an exception in expected behavior but by definition). This information/warning may not necessarily indicate an error. For example, a timeout may be an acceptable response for a wait API but is not indicative of a failure.

2.1.1 Synchronous Errors

When a core runtime API is called by the user and does not execute successfully, the core runtime returns a status that can help determine a cause of the unsuccessful execution. Each API call discussed in this chapter defines what constitutes a successful execution. While a few error conditions can be generalized to a certain degree (e.g. failure in allocating system memory) many errors can have system/implementation specific explanations.

The HSA core runtime API defines an enumeration that captures the result of any API function that has been executed (the only exception to this behavior are setter/getter API that access core runtime structures). This enumeration is of the type `hsa_status_t` and enumerates *success*, *info*, and *error*. The *info* status definition is discussed in Section 2.1.2.

Success status is a single value, `HSA_STATUS_SUCCESS`. Description of every core runtime API call that returns `hsa_status_t` explains the expected successful behavior for that API. The value of `HSA_STATUS_SUCCESS` is always 0.

Error status could be due to user input/actions that are not allowed (e.g. negative value in a size for allocation) or systemic errors (e.g. an asynchronous activity lead to a failure that cascaded into a failure in this API). The constants used for error status are restricted to the negative range of values within the `hsa_status_t` enumeration. Errors must always have a negative value. The Name of any constant that indicates an error status is prefixed by `HSA_STATUS_ERROR`. Errors could potentially be implementation.

While the name of the constant in itself is informative for success, info or error status, there may be scenarios where (i) the user may request more information about the meaning of a particular status, or, (ii) the return status was implementation specific and the user needs to decode it. In the case of implementation specific status, the negative number returned for error may not correspond to a particular enumeration constant. To query additional information on synchronous errors, the core runtime defines the following API:

```
hsa_status_t hsa_status_query_description(hsa_status_t input_status,
    uint64_t *status_info,
    char * const * status_info_string);
```


`input_status`
    is input argument. Any unsuccessful API return status that the user is seeking more information on.

status_info
> user allocated, output. In addition to the string. This value could be 0 and in itself (without status_info_string) may not be independently interpretable by the user.

status_info_string
> output from the API, a ISO/IEC 646 encoded english language string that potentially describes the error status. The string terminates in a ISO 646 defined NUL char.

This API returns HSA_STATUS_SUCCESS if one or both of the *status_info* and *status_info_string* have been successfully updated with information regarding the input *input_status*. Otherwise it returns one of the following errors:

▷ HSA_STATUS_NONE when no additional information is available regarding the status user requested.

▷ HSA_STATUS_ERROR_INVALID_ARGUMENT if a NULL value is passed for either of the arguments

2.1.2 Asynchronous Errors and Notifications

The HSA core runtime supports user-defined callbacks to handle asynchronous errors. There are two different categories of callbacks that can be registered by the user: (i) for asynchronous information or warnings generated when the runtime is executing, or, (ii) for asynchronous errors that get generated in packet processor, or while executing a kernel . The core runtime supports a callback each for asynchronous errors and notifications. The user must use caution when using blocking functions within their callback implementation – a callback that does not return can render the runtime state to be undefined. The user cannot depend on thread local storage within the callbacks implementation and may safely kill the thread that registers the callback. It is the user's responsibility to ensure that the callback function is thread-safe. The runtime does not implement any default callbacks.

Asynchronous Notification of Information or Warning

The information/warning status is represented by a value greater than 0 within the hsa_status_t enumeration. The status is up to user interpretation and the runtime allows the user to register a callback to take necessary action. Consider the example where a user calls the initialize API to initialize the core runtime and the return status is HSA_STATUS_INFO_ALREADY_INITIALIZED (to indicate that the core runtime has already been initialized). This result may be interpreted differently in different usage scenarios. A callback for such notifications may be registered via **hsa_open** API discussed in Section 2.2 or via **hsa_notification_callback_register** API, which is defined as follows:

```
hsa_status_t hsa_notification_callback_register(void (
    *notification_callback)(const hsa_notification_info_t *info),
    void *user_data,
    hsa_runtime_context_t *context);
```

```
notification_callback
     input. the callback that the user is registering, the callback is called with info as a
     parameter. User can read the structure and access its elements.

user_data
     input. the user data to call the callback with. info → user_data will be filled with value
     when the callback is called.

context
     the runtime context that this callback is being registered for.
```

The *context* parameter is used to identify a particular runtime context that this callback is registered for. When a callback is registered for a particular context, it will only be invoked if the notification is for an action in that context. Section 2.2 discusses the context in detail. The **hsa notification callback register** API can return one of the following errors:

▷ HSA_STATUS_ERROR_OUT_OF_RESOURCES if there is a failure in allocation of an internal structure required by the core runtime library in the context of registering a callback. This error may also occur when the core runtime library needs to spawn threads or create internal OS-specific events.

▷ HSA_STATUS_ERROR_INVALID_ARGUMENT if *info* is NULL. @ HSA_STATUS_INVALID_CONTEXT if the context is invalid (e.g. referenced counted to 0).

One of the arguments of the notification callback is a structure that contains notification information. The structure is defined as follows:

```
typedef struct hsa_notification_info_s{
    hsa_status_t    status;
    void           *ptr_info;
    char           *string_info;
    void           *user_data;
} hsa_notification_info_t;
```

```
status
     the info status enum value

ptr_info
     a pointer to more information, this could be pointing to implementation specific details that
     could be useful to some tools or to binary data

string_info
     ISO/IEC 646 character encoding must be used. A string indicating some error
     information. The string should be NUL terminated per ISO 646.
```

> user_data
>> a pointer to user supplied data

Asynchronous Notification of Errors

The HSA system can have several queues in operation and several kernels executing from these queues asynchronously. When any asynchronous activity generates an error, the action that initiated the activity may have concluded. To deal with asynchronous errors, the core runtime supports asynchronous error callbacks. The asynchronous error callback may be registered by means of the **hsa_open** API discussed in Section 2.2 or via **hsa_error_callback_register** API, which is defined as follows:

```
hsa_status_t hsa_error_callback_register(void ( *error_callback)(const
    hsa_async_error_info_t *info),
    void *user_data,
    hsa_runtime_context_t *context);
```

> error_callback
>> input. the callback that the user is registering, the callback is called with info structure. User can read the structure and access its elements.
>
> user_data
>> input. the user data to call the callback with. info-¿user_data will be filled with value when the callback is called.
>
> context
>> the runtime context that this callback is being registered for.

Details on how association of the callback can be done with asynchronous activities are discussed in Sections 2.2 and 2.5. The *context* parameter is used to identify a particular runtime context that this callback is registered for. When a callback is registered for a particular context, it will only be invoked if the notification is for an action in that context. For example, if a queue was created for a runtime context *c1* and a callback registered for a context *c2* but not for *c1*, any error on the queue, such as a packet processing error, will not trigger the execution of asynchronous error callback registered for context *c1*. This API can return one of the following errors:

> ▷ HSA_STATUS_ERROR_OUT_OF_RESOURCES if there is a failure in allocation of an internal structure required by the core runtime library in the context of registering a callback. This error may also occur when the core runtime library needs to spawn threads or create internal OS-specific events.

> ▷ HSA_STATUS_ERROR_INVALID_ARGUMENT if *info* is NULL.

One of the arguments of the notification callback is a structure that contains notification information. The structure is defined as follows:

```
typedef struct hsa_async_error_info_s{
    hsa_status_t  error_type;
    uint32_t      queue_id;
    void          *ptr_info;
    char          *string_info;
    void          *user_data;
    uint64_t      timestamp;
    uint64_t      reserved1;
    uint64_t      reserved2;
    uint64_t      reserved3;
} hsa_async_error_info_t;
```

error_type
    indicates the type of the error, based on this, the user knows if and packet_id is available in one of the reserved words.

queue_id
    the queue that processed the entity that caused the asynchronous error.

ptr_info
    a pointer to more information, this could be pointing to implementation specific details that could be useful to some tools or to binary data

string_info
    ISO/IEC 646 character encoding must be used. A string indicating some error information. The string should be NUL terminated per ISO 646.

user_data
    a pointer to user supplied data

timestamp
    timestamp – system timestamp to indicate when the error was discovered, the implementation may chose to always return 0 and user must take it into account when utilizing the timestamp.

reserved1
    additional info to be intepreted based on the error_type

reserved2
    additional info to be intepreted based on the error_type

reserved3
    additional info to be intepreted based on the error_type

2.1.3 Asynchronous Notification Example

This is work-in-progress – the chapter needs to be written.

2.2 Open and close API

Since HSA core runtime is a user mode library, its state is a part of the application's process space. When the runtime is opened for the first time, a runtime instance for that application process is created. Closing a runtime destroys this instance. An application may open (or close) the HSA runtime multiple times within the same process and potentially within multiple threads – only a single instance of the runtime, per-process, will exist.

The core runtime defines a runtime context that acts as a reference counting mechanism and a scheme to differentiate multiple usages of the runtime within the same application process. The runtime context is generated when the runtime is opened or when a user calls the acquire API that is defined in this Section. As an example, consider an application that is using the runtime but also uses a library, this library also creates HSA queues and submits work to them. Both the library and the application may want to register callbacks, and to capture notifications/errors of their specific usage. The runtime context helps identify the different usages (within the same process) and channel errors and notifications to appropriate callbacks. It also acts as a reference counting mechanism; while correctly *acquired* , the runtime context ensures that the runtime instance will not be shutdown until the context is *released* (this, in effect, is the reference counting part of the context).

This section defines four new API, **hsa_open** to open the runtime instance, **hsa_close** to close it, **hsa_context_acquire** to create a new context (and increment the reference count), and, **hsa_context_release** to release the acquired context (and decrement the reference count).

Invocation of **hsa_open** initializes the HSA runtime if it is not already initialized. It is allowed for applications to invoke **hsa_open** multiple times and do multiple **hsa_close** API calls. The HSA open call returns a new context at every invocation. Reference counting is a mechanism that allows the runtime to keep a count of the number of different usages of the runtime API within the same application process. This ensures that the runtime stays active until a **hsa_close** is called by the user when the reference count represented by that runtime context is 1.

The definition of the **hsa_open** API is as follows:

```
hsa_status_t hsa_open(hsa_runtime_context_t **context);
```

context
    output, user allocated. A type for reference counting.

The open API returns `HSA_STATUS_SUCCESS` if the initialization was successful. Otherwise it returns one of the following errors:

  ▷ `HSA_STATUS_ERROR_OUT_OF_RESOURCES` if there is a failure in allocation of an internal structure required by the core runtime library. This error may also occur when the core runtime library needs to spawn threads or create internal OS-specific events.

  ▷ `HSA_STATUS_ERROR_COMPONENT_INITIALIZATION` if there is a non-specific failure in initializing one of the components.

  ▷ `HSA_STATUS_ERROR_CONTEXT_NULL` if the context pointer passed by the user is NULL. User is required to pass in a memory backed context pointer.

If the HSA runtime is already initialized, an asynchronous notification is generated by the runtime and `HSA_STATUS_SUCCESS` is returned. If the user chooses to capture this asynchronous notification, the user should define a callback and associate it with the context returned by the **hsa_open** call. Each **hsa_open** call increments the reference count before returning success.

The runtime defines **hsa_close** as the corresponding API call to finalize the use of the runtime API. This API takes in a context as input. This API updates the reference count for every invocation. Once the reference count is 0, it proceeds to relinquish any resources allocated for the runtime and closes the runtime instance. It is possible in a multi-threaded scenario that one thread is doing a close while the other is trying to acquire the runtime context or do an open. The core runtime specification defines that an acquire with an input context that represents a closed runtime instance will fail. However, **hsa_open** can be called to create a new instance of the runtime after it is closed. The API for **hsa_close** is defined as follows:

```
hsa_status_t hsa_close(hsa_runtime_context_t *context);
```

---

context
     the context to close

The close API returns `HSA_STATUS_SUCCESS` if the close was successful. Otherwise, it returns one of the following errors:

  ▷ `HSA_STATUS_ERROR_NOT_INITIALIZED` if the close was called (a) either before the runtime was initialized, or (b) after it has already been successfully closed.

  ▷ `HSA_STATUS_ERROR_RESOURCE_FREE` if some of the resources consumed during initialization by the runtime could not be freed.

If the close is called when the reference count is not 1, it is still considered a succesful invocation of close, in that, the `HSA_STATUS_SUCCESS` is returned with status `HSA_STATUS_CLOSE_CONTEXT_ACTIVE`, is generated by the runtime on the context that is till active before the API returns.

The HSA core runtime API for an acquire on a context, **hsa_context_acquire**, is defined as follows:

```
hsa_status_t hsa_context_acquire(hsa_runtime_context_t *input_context);
```

input_context
     input, user allocated. the context that the user is expilicitly reference counting, increment
     reference count if not 0

The open API returns HSA_STATUS_SUCCESS if the acquire was successful and if *output_context*
holds the new context generated. Otherwise it returns one of the following errors:

   ▷ HSA_STATUS_ERROR_NOT_INITIALIZED if the **hsa_acquire_context** was called (a) either
     before the runtime was initialized, or (b) after it has already been closed.

   ▷ HSA_STATUS_CONTEXT_LIMIT_REACHED if the increment has taken the context to the
     UINT64_MAX value defined in stdint.h C header file.

The corresponding release API, **hsa_context_release** is defined as follows:

```
hsa_status_t hsa_context_release(hsa_runtime_context_t *input_context);
```

context
     input, user allocated. the context that the user is expilicitly reference counting, decrement
     reference count if not 1

The **hsa_context_release** API returns HSA_STATUS_SUCCESS if the release was successful.
Otherwise it returns the following error:

   ▷ HSA_STATUS_ERROR_NOT_INITIALIZED if the **hsa_release_context** was called before the
     runtime, after reference count has reached a value of 0.

     Runtime context cannot be incremented beyond a 64bit unsigned integer. The context does
     not wrap around.

2.2.1 Open/Close Example

```
/* The runtime is opened for the first time */
status = hsa_open(&runtime_context);
assert(status == HSA_STATUS_SUCCESS);

/* now the context has been acquired, reference count associated with
   the
```

```
 * context incremented internally by the runtime*/
hsa_context_acquire(&runtime_context);

/* this close is just going to generate an asynchronous error and
   return the
 * success status */
hsa_close(&runtime_context);
assert(status == HSA_STATUS_SUCCESS);

/* this queue create will still succeed provided component is valid
   and
 * runtime can allocate required resources*/
status = hsa_queue_create(component, 1024, 0, NONE, runtime_context,
   &queue, NULL);
if(status != HSA_STATUS_SUCCESS) {
  assert(HSA_STATUS_SUCCESS == hsa_close());
  exit(1);
}

/*
 * ....
 * ....
 * ....
 */

/* this API will succeed, the context is now as it was after the
   hsa_open */
hsa_context_release(&runtime_context);
assert(status == HSA_STATUS_SUCCESS);

/* this close will attempt to release all runtime allocated releases,
 * components, topology table, queues, etc. are no longer valid */
hsa_close(&runtime_context);
assert(status == HSA_STATUS_SUCCESS);

/* this queue create will fail*/
status = hsa_queue_create(component, 1024, 0, NONE, runtime_context,
   &queue1, NULL);
if(status != HSA_STATUS_SUCCESS) {
  exit(1);
}
```

2.3 HSA Topology and Component

HSA platform topology information is provided by the runtime by way of data structures so user can gather details about how a HSA system/platform exposed its architectural details such as components, memory, caches and connectivity (platform topology requirement is described in the

Figure 2.1: Structure of the Topology Table

SAR document). This information could be utilized by the user in different ways including decisions on where to execute a particular user task. Core runtime specification defines the topology table data structure and other data structures to represent topology hierarchy. After the core runtime is initialized with **hsa_open**, the user may create a local copy of the topology information using the API **hsa_topology_table_create**. The user can parse this table representing the HSA system to gather details such as the number of different HSA Components on the system with local access to a particular set of memory resources.

The **hsa_topology_table_create** API is defined as follows:

```
hsa_status_t hsa_topology_table_create(hsa_topology_header_t **header);
```

header
    output, runtime allocated. The topology header, this includes the base pointers to the rest of the topology table.

The API returns HSA_STATUS_SUCCESS if the table has been successfully created and returned by way of the *header*. Otherwise, it returns one of the following errors:

▷ HSA_STATUS_ERROR_INVALID_ARGUMENT if *header* is NULL.

▷ HSA_STATUS_ERROR_OUT_OF_RESOURCES if there is a failure in allocation of an internal structure required by the core runtime or in the creation of table header or the actual table.

The table structure is shown in Figure 2.1. The first entity in the table is a table header. This is the output of the **hsa_topology_table_create** API. The table header is defined by the following structure:

```
typedef struct hsa_topology_header_s{
    uint32_t        header_size_bytes;
    hsa_platform_t  platform;
    uint32_t        table_size;
    void            *topology_table_base;
} hsa_topology_header_t;
```

header_size_bytes
    size of the header

platform
    the hierarchical platform structure that abstracts the table

table_size
    size of the table

topology_table_base
    table base address, points to the table which starts with hsa_platform_t structure

The table header structure includes the platform structure (`hsa_platform_t`). The platform information in the platform structure includes size/offset-array pairs for HSA agents (`hsa_agent_t`), memory (`hsa_memory_descriptor_t`) and cache (`hsa_cache_descriptor_t`). HSA platform can have a hierarchical structure with multiple components/agents and physical memories. The `hsa_platform_t` structure also includes properties such as the clock frequency that are common across the platform and also links to various elements in the topology table (see Figure 2.1).

The platform structure is defined as follows:

```
typedef struct hsa_platform_s{
    uint32_t                hsa_system_timestamp_frequency_mhz;
    uint8_t                 number_of_nodes;
    uint32_t                *node_id;
    uint32_t                number_of_agents;
    uint32_t                *agent_offset_list_bytes;
    uint32_t                number_memory_descriptors;
    uint32_t                *memory_descriptor_offset_list_bytes;
    uint32_t                number_cache_descriptors;
    uint32_t                *cache_descriptors_offset_list_bytes;
} hsa_platform_t;
```

hsa_system_timestamp_frequency_mhz
    1-400MHz.

---

number_of_nodes
    number of different nodes in this platform configuration.

node_id
    ids of the nodes.

number_of_agents
    number of agents.

agent_offset_list_bytes
    agent list, refers to the offsets in platform table.

number_memory_descriptors
    number of the different types of memories available to this agent.

memory_descriptor_offset_list_bytes
    each element in the array carries an offset into the topology table to where memory
    descriptors are located. Number of elements in array equals number_memory_descriptors.

number_cache_descriptors
    number of caches available to this agent/component

cache_descriptors_offset_list_bytes
    array of offsets (into the topology table) to cache descriptors. Number of elements in array
    equals number_cache_descriptors.

When no information is available for a particular element, the corresponding *number_ <element >s*
field is set to zero by the runtime in the platform structure. Platform structure maps to the agents,
cache and physical memory, etc. in the topology table for all nodes in the platform.

The core runtime defines the following structure to represent cache:

```
typedef struct hsa_cache_descriptor_s{
    uint32_t hsa_node_id;
    uint32_t hsa_cache_id;
    uint8_t  levels;
    uint8_t  *associativity;
    uint64_t *cache_size;
    uint64_t *cache_line_size;
    bool     *is_inclusive;
} hsa_cache_descriptor_t;
```

---

hsa_node_id
    id of the node this memory belongs to.

---

hsa_cache_id
    unique identified for this cache with in the system.

levels
    number of levels of cache (for a mult-level cache)

associativity
    associativity of this cache, array with number of entries = number of levels

cache_size
    size at each level, this array is of size = levels

cache_line_size
    cache line size at each level, this array is of size = levels

is_inclusive
    is the cache inclusive with the level above? The size of this array is level-1

The structure holds associativity, cache size, cache line size for all levels of cache and the inclusive property for all but the last level. Each cache in the HSA system has a unique cache ID identifying it.

The memory descriptor structure represents a physical memory block or region and includes elements to provide bandwidth, interleave characteristics and latency for accessing memory. Implementations may choose not to provide memory bandwidth or latency information. The memory descriptor structure is defined as follows:

```
typedef struct hsa_memory_descriptor_s{
    uint32_t       hsa_node_id;
    uint32_t       hsa_memory_id;
    hsa_segment_t  supported_segment_type_mask;
    uint64_t       physical_address_base;
    uint64_t       virtual_address_base;
    uint64_t       size_in_bytes;
    uint64_t       peak_bandwidth_mbps;
} hsa_memory_descriptor_t;
```

hsa_node_id
    id of the node this memory belongs to.

hsa_memory_id
    unique identified for this memory with in the system.

supported_segment_type_mask
    information on segments that can use this memory.

---

physical_address_base
    base of the physical address for this memory

virtual_address_base
    base of the virtual address for this memory, if applicable

size_in_bytes
    size

peak_bandwidth_mbps
    theoretical peak bandwidth in mega-bits per second to access this memory from the
    agent/component

The structure:

```
typedef struct hsa_segment_s{
    uint8_t global:1;
    uint8_t private:1;
    uint8_t group:1;
    uint8_t kernarg:1;
    uint8_t readonly:1;
    uint8_t spill:1;
    uint8_t arg:1;
    uint8_t reserved:1;
} hsa_segment_t;
```

---

global:1
    if bit is set, the element/mask represents global segment

private:1
    if bit is set, the element/mask represents private segment

group:1
    if bit is set, the element/mask represents group segment

kernarg:1
    if bit is set, the element/mask represents kernarg segment

readonly:1
    if bit is set, the element/mask represents readonly segment

spill:1
    if bit is set, the element/mask represents spill segment

arg:1
    if bit is set, the element/mask represents arg segment

```
reserved:1
    reserved
```

can represent any combination of the 7 HSA segments, a single bit for each segment.

The HSA Agent data structure represents an HSA component when the *agent_type* field in the agent structure is set to a 1 (i.e. bit 0 is set to 1). The structure contains elements that describe its properties. Each component has access to coherent global memory (the HSA global segment, and as per the requirement defined in SAR, has access to other segments as well). The *agent_type* is utilized as a bit-field. Setting bit 2 indicates that the agent is a host, bit 3 indicates that agent can participate in agent dispatches. All three bits or a combination of them can be set by the HSA runtime.

The structure of the HSA agent/component is defined as follows:

```
typedef struct hsa_agent_s{
    bool                    is_pic_supported;
    uint32_t                hsa_node_id;
    uint32_t                agent_id;
    hsa_agent_type_t        agent_type;
    char                    vendor[16];
    char                    name[16];
    uint64_t                *property_table;
    uint32_t                number_memory_descriptors;
    uint32_t                *memory_descriptors;
    uint32_t                number_cache_descriptors;
    uint32_t                *cache_descriptors;
    uint32_t                number_of_subagents;
    uint32_t                *subagent_offset_list;
    uint32_t                wavefront_size;
    uint32_t                queue_size;
    uint32_t                group_memory_size_bytes;
    uint32_t                fbarrier_max_count;
} hsa_agent_t;
```

is_pic_supported
    does it support position independent code?. Only applicable when the agent is a
    component.

hsa_node_id
    id of the node this agent/component belongs to.

agent_id
    Unique identifier for an HSA agent.

`agent_type`
    an identifier for the type of this agent.

`vendor[16]`
    The vendor of the agent/component. ISO/IEC 646 character encoding must be used. If the name is less than 16 characters then remaining characters must be set to 0.

`name[16]`
    The name of this agent/component. ISO/IEC 646 character encoding must be used. If the name is less than 16 characters then remaining characters must be set to 0.

`property_table`
    table of properties of the agent, any property that is not available has a value of 0

`number_memory_descriptors`
    number of the different types of memories available to this agent.

`memory_descriptors`
    Array of memory descriptor offsets. Number of elements in array equals number_memory_descriptors.

`number_cache_descriptors`
    Number of caches available to this agent/component

`cache_descriptors`
    Array of cache descriptor offsets. Number of elements in array equals number_cache_descriptors.

`number_of_subagents`
    Number of subagents.

`subagent_offset_list`
    subagent list of offsets, points to the offsets in the topology table.

`wavefront_size`
    Wave front size, i.e. number of work-items in a wavefront.

`queue_size`
    Maximum size of the user queue in bytes allocatable via the runtime.

`group_memory_size_bytes`
    Size (in bytes) of group memory available to a single work-group.

`fbarrier_max_count`
    max number of fbarrier that can be used in any kernel and functions it invokes.

Within the agent, the agent type is an enumeration that is defined as follows:

```
typedef enum hsa_agent_type_t{
    HOST=1,
```

```
    COMPONENT=2,
    AGENT_DISPATCH=4
}hsa_agent_type_t;
```

HOST=1
indcates that the agent represents the host.

COMPONENT=2
indicates that agent represents an HSA component.

AGENT_DISPATCH=4
indicates that the agent is capable of agent dispatches, and can serve as a target for
them.

The user must destroy the topology table before closing the runtime. The
`hsa_topology_table_destroy` API is defined by the runtime for the user to destroy the topology
table. Once a table is created, some parts of it may become invalid if any HW is
hot-plugged/unplugged or encounters an error. If such a change occurs, the HSA runtime
generates an asynchronous error (see Section 2.1.2) with the `hsa_status_t` enumeration of
`HSA_ERROR_TOPOLOGY_CHANGE`. This is an indication to the user that any current usage of
topology table must be stopped and a new topology table obtained by using the
hsa_topology_table_create API call. The runtime guarantees that any call made to
hsa_topology_table_create API after the asynchronous error is observed will return the latest
version of the topology table at the time of the API invocation. However, if the same HW was
hot-swapped out and in with the same interval, or if the error encountered in a component was
recovered, the topology table may not look different from the users perception.

2.3.1 Topology Example

This is work-in-progress – the chapter needs to be written.

2.4 Memory based Signals and Synchronization in HSA

In a HSA system, memory is coherent and can serve as a means for message passing,
asynchronous communication or synchronization between various elements. A signal is an
alternative, possibly more power-efficient, communication mechanism between two entities in a
HSA system. A signal carries a value, which can be updated or conditionally waited upon via an
API call or, optionally, an HSAIL instruction. A signal structure is opaque and is always typedef'ed
to `uint64_t`. Implementations can use the most power-efficient send-propagation and wait
techniques available to them on the HSA system.

HSA Agent is a participant in a HSA memory based signalling and synchronization. This feature requires a runtime API for allocation of signals that may be used for synchronization and states that the signal is opaque and may contain implementation specific information.

The signal handle, which represents a signal, and the value the signal carries, are defined as follows:

```
typedef uint64_t hsa_signal_handle_t;


typedef union signal_value_s{
    /**
    * Pointer to the base of the HSAIL segment
    */
    int value32;

    /**
    * Pointer to the base of the HSAIL segment
    */
    int64_t value64;
} hsa_signal_value_t;
```

An API, **hsa_signal_create**, to support creation of signals, is defined as follows:

```
 hsa_status_t hsa_signal_create(hsa_signal_value_t intial_signal_value,
                         hsa_signal_handle_t *signal_handle,
                         hsa_runtime_context_t *context);
```

initial_signal_value
    input. Initial value at the signal, the signal is initialized with this value.

signal_handle
    User allocated, output. The handle of the signal that this API creates. This handle is opaque.

context
    input. The context in which this signal is being created. Any errors/notifications will be reported via callbacks registered in the same context.

The signal create API returns HSA_STATUS_SUCCESS if the signal object has been successfully created. Otherwise, it returns one of the following:

  ▷ HSA_STATUS_ERROR_OUT_OF_RESOURCES if there is a failure in allocation of an internal structure required by the core runtime library in the context of the message queue creation. This error may also occur when the core runtime library needs to spawn threads or create internal OS-specific events.

▷ HSA_STATUS_ERROR_INVALID_ARGUMENT if *signal_handle* is NULL or an invalid pointer of an invalid/NULL context is passed in as an argument.

Once a signal is created for a particular context, it may be bound to other contexts. This is useful when signal is used across different components of a users application. An API to bind the signal to a particular runtime context is defined as follows:

```
hsa_status_t hsa_signal_bind(hsa_signal_handle_t signal_handle,
    hsa_runtime_context_t *context);
```

signal_handle
    input. Opaque handle that was obtained from **hsa_signal_create** API.

context
    input. Additional context to which this signal should be bound to

This API returns HSA_STATUS_SUCCESS if the bind was successful. Otherwise, it returns one of the following errors:

▷ HSA_STATUS_ERROR_INVALID_ARGUMENT if *signal_handle* is NULL or invalid or if the *context* is NULL or invalid.

The corresponding signal destruction API is defined as follows:

```
hsa_status_t hsa_signal_destroy(hsa_signal_handle_t signal_handle);
```

signal_handle
    input. Opaque handle that was obtained from **hsa_signal_create** API.

The signal destroy API returns HSA_STATUS_SUCCESS if the signal object has been successfully destroyed. Otherwise, it returns one of the following:

▷ HSA_STATUS_ERROR_INVALID_ARGUMENT if *signal_handle* is invalid.

A signal can also be unbound from a particular context if the user no longer wants to receive notifications about this signal in the callback registered for that context. The API to unbind is defined as follows:

```
hsa_status_t hsa_signal_unbind(hsa_signal_handle_t signal_handle,
    hsa_runtime_context_t *context);
```

signal_handle
   input. Opaque handle that was obtained from **hsa_signal_create** API.

context
   input. Unbind the signal from this context.

The API returns HSA_STATUS_SUCCESS if the signal is successfully unbound from the context. Otherwise, it can return one of the following errors:

▷ HSA_STATUS_ERROR_SIGNAL_NOT_BOUND if the signal was not already bound to that context.

▷ HSA_STATUS_ERROR_INVALID_ARGUMENT if *signal_handle* is NULL or invalid or if the *context* is NULL or invalid.

As per the HSA SAR specification the signals may only be created and operated on by either instructions in HSAIL or the HSA runtime API. Sending a signal entails updating a particular value at the signal. Waiting on a signal returns the current value at the opaque signal object – the wait has a runtime defined timeout which indicates the maximum amount of time that an implementation can spend waiting for a particular value before returning.

The API to query the timeout is defined as:

```
uint64_t hsa_signal_get_timeout();
```

This getter API does not return a status. This API returns the timeout, which indicates the maximum amount of time an implementation can spend in a wait operation on the signal. The return value is in the units of the system-wide clock who's frequency is available via the hsa_platform_t structure (see Section 2.3). As per SAR, the HSA system has a system-wide timestamp that operates at a fixed frequency. The frequency can be queried via the hsa_platform_t structure defined in Section 2.3. The timeout is incremented at the same frequency. The user can use this information to translate the timeout to a different frequency domain.

The send signal API sets the signal handle with caller specified value. Any subsequent wait on the signal handle would be given a copy of this new signal value after the wait condition is met (and before the timeout expires). The signal infrastructure allows for multiple waiters on a single signal. A multi-threaded user application can have multiple threads sending and waiting on signals.

In addition to the update of signals using Send, the API for send signal must support other atomic operations as well. HSA defines *AND, OR, XOR, Exchange, Add, Subtract, Increment, Decrement, Maximum, Minimum* and *CAS*. Apart from the no synchronization case, which is referred to as *none* synchronization, there are three types of synchronization defined in the systems architecture requirements:

Acquire synchronization

   No memory operation listed after the acquire can be executed before the
   acquire-synchronized operation. Acquire synchronization can be applied to various
   operations including a load operation.

Release synchronization

   No memory operation listed before the release can be executed after the
   release-synchronized operation. Release synchronization can be applied to various
   operations including a store operation.

Acquire-Release synchronization

   This acts like a fence. No memory operation listed before the Acquire-Release synchronized
   operation can be moved after it nor can any memory operation listed after the
   Acquire-Release synchronized operation be executed before it.

Relaxed synchronization

   No synchronization is applied to the send or wait operation.

Each operation on a signal value has the type of synchronization explicitly included in its name. For
example, Send-Release is a Send on a signal value with Release synchronization.

Hence, the following table 2.1 represents the complete set of actions (with associated
synchronization) that can be performed on a signal value:

For efficiency, a unique signal API has been created for each of these actions. In the description of
the API, for convenience, *value@signal_handle* is used to represent the value at a signal.

```
hsa_status_t hsa_signal_send_relaxed(hsa_signal_handle_t signal_handle,
                hsa_signal_value_t signal_value);
```

---

signal_handle
   input. Opaque handle of the signal object that is to be signaled.

signal_value
   input. Value of the signal, with relaxed semantics, signal_value is assigned to the
   value@signal_handle.

```
hsa_status_t hsa_signal_add_release(hsa_signal_handle_t signal_handle,
                hsa_signal_value_t add_value);
```

---

signal_handle
   input. Opaque handle of the signal object that is to be signaled.

add_value
   input. Value to add to the value@signal_handle (the addition is atomic).

Actions with associated synchronization
Send with release
Send with relaxed
AND with release
AND with relaxed
OR with release
OR with relaxed
XOR with release
XOR with relaxed
Exchange with acquire-release
Exchange with relaxed
Add with release
Add with relaxed
Subtract with release
Subtract with relaxed
Increment with release
Increment with relaxed
Decrement with release
Decrement with relaxed
Maximum with acquire-release
Maximum with relaxed
Minimum with relaxed
CAS release

Table 2.1: Actions with Associated Synchronization

```
hsa_status_t hsa_signal_add_relaxed(hsa_signal_handle_t signal_handle,
                hsa_signal_value_t add_value);
```

signal_handle
> input. Opaque handle of the signal object that is to be signaled.

add_value
> input. Value to add to the value@signal_handle (the addition is atomic).

```
hsa_status_t hsa_signal_and_release(hsa_signal_handle_t signal_handle,
                hsa_signal_value_t and_value);
```

signal_handle
> input. Opaque handle of the signal object that is to be signaled.

and_value
> input. Value to do an And with, value@signal_handle &= and_value.

```
hsa_status_t hsa_signal_and_relaxed(hsa_signal_handle_t signal_handle,
                hsa_signal_value_t and_value);
```

signal_handle
> input. Opaque handle of the signal object that is to be signaled.

and_value
> input. Value to do an And with, value@signal_handle &= and_value.

```
hsa_status_t hsa_signal_or_release(hsa_signal_handle_t signal_handle,
                hsa_signal_value_t or_value);
```

signal_handle
> input. Opaque handle of the signal object that is to be signaled.

or_value
> input. value@signal_handle OR= or_value.

```
hsa_status_t hsa_signal_or_relaxed(hsa_signal_handle_t signal_handle,
                hsa_signal_value_t or_value);
```

signal handle
    input. Opaque handle of the signal object that is to be signaled.

or value
    input. value@signal handle OR= or value.

```
hsa_status_t hsa_signal_xor_release(hsa_signal_handle_t signal_handle,
                hsa_signal_value_t xor_value);
```

signal handle
    input. Opaque handle of the signal object that is to be signaled.

xor value
    input. Value to do an XOR with, value@signal handle XOR= xor value.

```
hsa_status_t hsa_signal_xor_relaxed(hsa_signal_handle_t signal_handle,
                hsa_signal_value_t xor_value);
```

signal handle
    input. Opaque handle of the signal object that is to be signaled.

xor value
    input. Value to do an XOR with, value@signal handle XOR= xor value.

```
hsa_status_t hsa_signal_exchange_release(hsa_signal_handle_t
    signal_handle,
                hsa_signal_value_t exchange_value,
                hsa_signal_value_t *value_at_signal);
```

signal handle
    input. Opaque handle of the signal object that is to be signaled.

```
exchange_value
```
    user allocated, input/output. Exchange value, the value to be placed at signal,
    value@signal_handle , after being stored in value_at_signal, is overwritten with
    exchange_value.

```
value_at_signal
```
    user allocated, output. value_at_signal = value@signal_handle; value@signal_handle =
    exchange_value.

```
hsa_status_t hsa_signal_exchange_relaxed(hsa_signal_handle_t
    signal_handle,
                hsa_signal_value_t exchange_value,
                hsa_signal_value_t *value_at_signal);
```

```
signal_handle
```
    input. Opaque handle of the signal object that is to be signaled.

```
exchange_value
```
    user allocated, input/output. Exchange value, the value to be placed at signal,
    value@signal_handle , after being stored in value_at_signal, is overwritten with
    exchange_value.

```
value_at_signal
```
    user allocated, output. value_at_signal = value@signal_handle; value@signal_handle =
    exchange_value.

```
hsa_status_t hsa_signal_decrement_release(hsa_signal_handle_t
    signal_handle,
                hsa_signal_value_t decrement_value);
```

```
signal_handle
```
    input. Opaque handle of the signal object that is to be signaled.

```
decrement_value
```
    input. Value the signal is to be decremented with, value@signal_handle -=
    decrement_value.

```
hsa_status_t hsa_signal_decrement_relaxed(hsa_signal_handle_t
    signal_handle,
                hsa_signal_value_t decrement_value);
```

signal_handle
    input. Opaque handle of the signal object that is to be signaled.

decrement_value
    input. Value the signal is to be decremented with, value@signal_handle -=
    decrement_value.

hsa_status_t **hsa_signal_cas_release**(*hsa_signal_handle_t signal_handle,*
                hsa_signal_value_t value_compare,
                hsa_signal_value_t value_replace,
                hsa_signal_value_t *value_at_signal);

signal_handle
    input. Opaque handle of the signal object that is to be signaled.

value_compare
    input. The value to compare value@signal_handle against (operator is equal to).

value_replace
    input. the value to replace the value@signal_handle with if signal.value was equal to
    value_compare.

value_at_signal
    user allocated, output. The value at the signal, prior to the atomic replace, if the
    comparision was successful.

hsa_status_t **hsa_signal_max**(*hsa_signal_handle_t signal_handle,*
                hsa_signal_value_t compare_value,
                hsa_signal_value_t *max_value);

signal_handle
    input. Opaque handle of the signal object that is to be signaled.

compare_value
    input. Compared with value@signal_handle to determine maximum.

max_value
    user allocated, output. MAX(compare_value, value@signal_handle). The signal value is
    updated with the max value.

```
hsa_status_t hsa_signal_min(hsa_signal_handle_t signal_handle,
                hsa_signal_value_t compare_value,
                hsa_signal_value_t *min_value);
```

signal_handle
    input. Opaque handle of the signal object that is to be signaled.

compare_value
    input. value@signal_handle is compared with this and the minimum of the two returned.

min_value
    user allocated, output. min(compare_value, value@signal_handle). The signal value is
    updated with the min value.

```
hsa_status_t hsa_signal_send_release(hsa_signal_handle_t signal_handle,
                hsa_signal_value_t signal_value);
```

signal_handle
    input. Opaque handle of the signal object that is to be signaled

signal_value
    input. Value of the signal

```
hsa_status_t hsa_signal_subtract_release(hsa_signal_handle_t
    signal_handle,
                hsa_signal_value_t subtract_value);
```

signal_handle
    input. Opaque handle of the signal object that is to be signaled

subtract_value
    input. Value to be subtracted from the value@signal_handle.

```
hsa_status_t hsa_signal_subtract_relaxed(hsa_signal_handle_t
    signal_handle,
                hsa_signal_value_t subtract_value);
```

> signal_handle
> > input. Opaque handle of the signal object that is to be signaled
>
> subtract_value
> > input. Value to be subtracted from the value@signal_handle.

hsa_status_t **hsa_signal_increment_release**(*hsa_signal_handle_t signal_handle*,
    hsa_signal_value_t increment_value);

> signal_handle
> > input. Opaque handle of the signal object that is to be signaled
>
> increment_value
> > input. Value to do the increment with

hsa_status_t **hsa_signal_increment_relaxed**(*hsa_signal_handle_t signal_handle*,
    hsa_signal_value_t increment_value);

> signal_handle
> > input. Opaque handle of the signal object that is to be signaled
>
> increment_value
> > input. Value to do the increment with

All of the **signal send** API return HSA_STATUS_SUCCESS if the send is successful. Any atomic operation that needed to be performed has been done successfully and any result value that needs to be returned has been copied into the user-given location. One of the following error values may be returned in case the send is not successful:

▷ HSA_STATUS_ERROR_INVALID_ARGUMENT if (a) the user is expecting an output but the pointer to the output signal value is invalid, (b) the *signal_value* doesn't represent a valid signal.

The user may wait on a signal, with a condition specifying the terms of wait. The wait can be done either in the HSA Component via an HSAIL wait instruction or via a runtime API defined here.

Waiting on a signal returns the current value at the signal. The wait may return before the condition is satisfied or even before a valid value is obtained from the signal. It is the users burden to check the return status of the wait API before consuming the returned value.

Wait *reads* the value, hence Acquire and Acquire-Release synchronizations may be applied to the read. The synchronization should only assume to have been applied if the status returned by the wait API indicates a success (i.e. return type is HSA_STATUS_SUCCESS). The two wait APIs to support both synchronizations are defined as follows:

```
hsa_status_t hsa_signal_wait_acquire(hsa_signal_handle_t signal_handle,
                           hsa_signal_condition_t cond,
                           hsa_signal_value_t compare_value,
                           hsa_signal_value_t *return_value);
```

signal_handle
    input. Opaque handle of the signal whose value is to be retrieved.

cond
    input. apply this condition to compare the wait_value with value@signal_handle and return the value@signal_handle only when the condition is met.

compare_value
    input. value to compare with.

return_value
    user allocated, output. Pointer to where the current value@signal_handle must be read into.

```
hsa_status_t hsa_signal_wait_acquire_release(hsa_signal_handle_t
    signal_handle,
                           hsa_signal_condition_t cond,
                           hsa_signal_value_t compare_value,
                           hsa_signal_value_t *return_value);
```

signal_handle
    input. Opaque handle of the signal whose value is to be retrieved.

cond
    input. apply this condition to compare the compare_value with value@signal_handle and return the value@signal_handle only when the condition is met.

compare_value
    input. value to compare with.

> return_value
>> user allocated, output. Pointer to where the current value@signal_handle must be read into.

The user must always check the return value of the wait before considering the *wait_value* as the wait may have returned due to a timeout. The wait API can return the following status:

 ▷ If an error is signaled on the signal the user is waiting on, the wait API returns HSA_STATUS_ERROR to indicate that an error has occurred. The API still returns the current value at the signal. The user may also inspect the value returned, when an error occurred (see Section 2.4.1).

 ▷ HSA_STATUS_ERROR_INVALID_ARGUMENT if (a) the user is expecting an output but the pointer to the output signal value is invalid, (b) the *signal_value* doesn't represent a valid signal.

 ▷ HSA_STATUS_INFO_SIGNAL_TIMEOUT the signal wait has timed out.

The hsa_wait_condition_t is defined as follows:

```
typedef enum hsa_signal_condition_t{
    HSA_EQUALS,
    HSA_NOTEQUALS,
    HSA_GREATER,
    HSA_GREATER_EQUALS,
    HSA_LESSER,
    HSA_LESSER_EQUALS
}hsa_signal_condition_t;
```

> HSA_EQUALS
>> the return from the wait API will be either when signal.value == wait_value or the max timeout has been reached.
>
> HSA_NOTEQUALS
>> the return from the wait API will be either when signal.value != wait_value or the max timeout has been reached.
>
> HSA_GREATER
>> the return from the wait API will be either when signal.value ¿ wait_value or the max timeout has been reached.
>
> HSA_GREATER_EQUALS
>> the return from the wait API will be either when signal.value ¿= wait_value or the max timeout has been reached.

HSA_LESSER
    the return from the wait API will be either when signal.value ¡ wait_value or the max
    timeout has been reached.

HSA_LESSER_EQUALS
    the return from the wait API will be either when signal.value ¡= wait_value or the max
    timeout has been reached.

The runtime also defines an API to query the current signal value. If the signal is being updated by
the component or other threads, there is no guarantee that the value returned by the query API is
the value of the signal even at the instance it has been returned. Queried value may be used to
check progress of a kernel, if the kernel were updating the signal at various stages of its execution.
Query is a non-blocking API and does not take `hsa_wait_condition_t` as input. It merely
obtains the current value at the signal.

The `hsa_signal_query_acquire` API is defined as follows:

```
hsa_status_t hsa_signal_query_acquire(hsa_signal_handle_t
    signal_handle,
                            hsa_signal_value_t *value);
```

---

signal_handle
    input. Opaque handle of the signal whose value is to be retrieved.

value
    user allocated, output. Pointer to where the current value@signal_handle must be read
    into.

The `hsa_signal_query_acquire` API returns `HSA_STATUS_SUCCESS` when the value at the
signal has been successfully returned. Otherwise, it returns one of the following errors:

 ▷ `HSA_STATUS_ERROR_INVALID_ARGUMENT` if *signal_handle* is invalid.

Signals may be utilized in many ways. For example, a running kernel, after it finishes producing a
part of its computation, may set the signal in the dependency packet of another kernel dispatch so
that the queue processor can resolve the dependency and launch the kernel.

Signals cannot be used for Inter-Process Communication (IPC).

2.4.1   Indicating Errors with Signals

To put the signal in error state, the two most significant bits in the signal value are set and all other
bits cleared. It is the users burden to check if an error has occurred by looking at the return code of
the `hsa_signal_wait<acquire_release/Acquire>` API. Any negative value at the signal
triggers the `HSA_STATUS_ERROR` return code from the wait API. A signal that is already in error
may further be decremented to a larger negative value.

2.4.2 Usage Example

This is work-in-progress – the chapter needs to be written.

2.5 Architected Queue in HSA

HSA hardware supports kernel dispatch through user mode queues. A queue in HSA is associated with a specific HSA component. There are two kinds of queues that are supported, an AQL queue which can consume any kind of packets discussed in Section 2.6. A service queue is defined a queue that consumes AGENT_DISPATCH packets. AGENT_DISPATCH packets can be used to specify runtime-defined or user registered functions that will be executed on the agent (typically, the host CPU).

An HSA component can have multiple AQL and service queues associated with it. Conceptually, user mode queues are ring buffers that expose separate memory locations defining the current read and write state of the queue. The HSA runtime allows the user to create a user mode queue via. `hsa_queue_create` API. The same API also allows to user to create a service queue. The user may chose to manage their own service queue.

In a HSA system, agents write AQL packets to the user mode queue queue to enqueue work on to the HSA components. The queue memory is processed by HSA packet processor(s) as though it is a ring buffer. The details on how commands can be written to the queue via AQL packets and the structure of the AQL packet are discussed in Section 2.6.

A queue in HSA is defined with the following structure:

```
typedef struct hsa_queue_s{
  uint32_t queue_type;
  uint32_t queue_features;
  uint64_t base_address;
  hsa_signal_handle_t doorbell_signal;
  uint32_t size;
  uint32_t queue_id;
  uint32_t queue_active_group_count_global;
  uint64_t service_queue;
} hsa_queue_t;
```

queue_type
>   used for dynamic queue protocol determination. Currently, 0, the default queue type, is the only type supported.

queue_features
>   bitfield to indicate specific features supported by queue. On a queue creation, if user observes that some unknown bits are set, then the user should ignore them.

`base_address`
> A 64-bit pointer to the base of the virtual memory which holds the AQL packets for the queue. At the time of queue creation, the address passed in by the user as queue memory is copied here. This address must be 64-byte aligned.

`doorbell_signal`
> After writing a packet to the queue, user must signal this signal object with the most recent write_offset. The packet may already have been processed by the packet processor by the time this doorbell is signaled, however, it may not be processed at all if the doorbellSignal is not signaled.

`size`
> A 32-bit unsigned integer which specifies the maximum size of the queue in the number of packets. The size of the queue is always aligned with a power of two number of AQL packets.

`queue_id`
> A 32-bit ID for a queue which is unique-per-process.

`queue_active_group_count_global`
> maximum number of concurrent workgroups that can run out of this queue

`service_queue`
> A pointer to another User Mode Queue that can be used by the HSAIL kernel to request system services. The serviceQueue property is provided by the application to the runtime API call when the queue is created, and may be NULL, the system provided serviceQueue or an application managed queue.

`base_address` is the starting address of the buffer where the packets will be written. `size` is simply the size of the queue in bytes. The `queue_id` member is the unique (per-process) identifier for a queue and helps identify a queue when more than one queue is present in the system.

Internally, the queue structure contains read index and write index. These are not exposed to the user directly. The user can access them by using the `hsa_queue_get/cas/add_write_index` and `hsa_queue_get_read_index` API. All of these API calls have different versions for different memory scopes.

The API is defined as follows:

```
uint64_t hsa_queue_get_read_index_relaxed(hsa_queue_t *queue);
```

`queue`
> input. The HSA queue structure.

```
uint64_t hsa_queue_get_read_index_acquire(hsa_queue_t *queue);
```

> queue
>> input. The HSA queue structure.

```
uint64_t hsa_queue_get_write_index_relaxed(hsa_queue_t *queue);
```

> queue
>> input. The HSA queue structure.

```
uint64_t hsa_queue_get_write_index_acquire(hsa_queue_t *queue);
```

> queue
>> input. The HSA queue structure.

```
uint64_t hsa_queue_set_write_index_relaxed(hsa_queue_t *q,
    uint64_t val);
```

> queue
>> input. The HSA queue structure.
>
> val input. The new value of the write index.

```
uint64_t hsa_queue_set_write_index_release(hsa_queue_t *q,
    uint64_t val);
```

> queue
>> input. The HSA queue structure.
>
> val input. The new value of the write index.

```
uint64_t hsa_queue_cas_write_index(hsa_queue_t *q,
    uint64_t old_val,
    uint64_t new_val);
```

> queue
>   input. The HSA queue structure.
>
> old_val
>   input. The value to compare with
>
> new_val
>   input. If a match is determined, the write index is updated with new_val

```
uint64_t hsa_queue_cas_write_index_release(hsa_queue_t *q,
    uint64_t old_val,
    uint64_t new_val);
```

> queue
>   input. The HSA queue structure.
>
> old_val
>   input. The value to compare with
>
> new_val
>   input. If a match is determined, the write index is updated with new_val

```
uint64_t hsa_queue_cas_write_index_acquire(hsa_queue_t *q,
    uint64_t old_val,
    uint64_t new_val);
```

> queue
>   input. The HSA queue structure.
>
> old_val
>   input. The value to compare with
>
> new_val
>   input. If a match is determined, the write index is updated with new_val

```
uint64_t hsa_queue_cas_write_index_relaxed(hsa_queue_t *q,
    uint64_t old_val,
    uint64_t new_val);
```

queue
    input. The HSA queue structure.

old_val
    input. The value to compare with

new_val
    input. If a match is determined, the write index is updated with new_val

```
uint64_t hsa_queue_cas_write_index_acquire_release(hsa_queue_t *q,
    uint64_t old_val,
    uint64_t new_val);
```

queue
    input. The HSA queue structure.

old_val
    input. The value to compare with

new_val
    input. If a match is determined, the write index is updated with new_val

```
uint64_t hsa_queue_add_write_index_relaxed(hsa_queue_t *q,
    uint64_t val);
```

queue
    input. The HSA queue structure

val input. The value to add to the write index

```
uint64_t hsa_queue_add_write_index_acquire(hsa_queue_t *q,
    uint64_t val);
```

```
queue
     input. The HSA queue structure

val input. The value to add to the write index
```

uint64_t **hsa_queue_add_write_index_release**(hsa_queue_t *q,
    uint64_t *val*);

```
queue
     input. The HSA queue structure

val input. The value to add to the write index
```

uint64_t **hsa_queue_add_write_index_acquire_release**(hsa_queue_t *q,
    uint64_t *val*);

```
queue
     input. The HSA queue structure

val input. The value to add to the write index
```

These API are all setter/getter APIs and hence do not return hsa_status_t. If the queue structure passed to the API is invalid, the behavior of the API is undefined. All the API return the value of the corresponding index. The CAS, ADD and WRITE API on the write index return the value of the write index prior to the update.

The write index is a unique identifier for AQL packets in the queue. The read index indicates the next AQL packet that will be consumed by the HSA packet processor. The write index memory is updated by the agents via the runtime defined API, while the read index memory location is updated by the HSA Component and can be read by the agent, a runtime specified API call, or the kernel via HSAIL operation.

The read index is automatically advanced when a packet is read by the HSA packet processor. When the agent observes that read index matches write index, at that instance, the queue can be considered empty (it does not mean that the kernels have finished execution, just that all packets have been consumed). The write index and the read index never wrap when the write index

reaches its maximum value. An asynchronous error is generated by the packet processor and queue is put in error state.

The *queue_active_group_count* is the count of maximum number of work-groups that can be executed in parallel for dispatches executed on this queue.

The `doorbell_signal` is a signal from the agent writing the AQL packet to the HSA packet processor indicating that it has work to do. The value which the `doorbell_signal` must be signaled with shall be the latest write index at which an AQL packet has been written into. The purpose of this signal is to *inform* the HSA packet processor that it has packets that need to be processed. However, packets may be processed by the HSA packet processor even before the `doorbell_signal` has been signaled by the agent writing the AQL packet. This is because when write index is advanced by the agent there are two scenarios that could arise:

- the HSA packet processor is in some low-powered state awaiting work and requires the `doorbell_signal` signal to *wake* it to continue reading packets.

- the HSA packet processor is already actively processing a packet and observes the write index being updated by the agent and continues to process the new packets written – even before the agent has signalled the `doorbell_signal`.

Hence, despite the fact that the AQL packet for which the agent is signalling the doorbell may already have been processed, the agent must ring the doorbell for every batch of AQL packets written.

The `hsa_queue` structure is the output of `hsa_queue_create` function, which is defined as follows:

```
hsa_status_t hsa_queue_create(const hsa_agent_t *component,
                              size_t size,
                              uint32_t queue_type,
                              hsa_service_queue_type_t
                                 service_queue_type,
                              hsa_runtime_context_t *context,
                              hsa_queue_t **queue,
                              hsa_queue_mailbox_t *mailbox);
```

`component`
    input. The component on which this queue is to be created.

`size`
    input. Size of the queue memory in number of packets in is expected to hold.

`queue_type`
    input. type of the queue (only type 0, which is default in-order issue queue, is supported at this time).

```
service_queue_type
```
   input. The user can choose between NONE (no service queue), COMMON (runtime
   provided service queue that is shared), NEW (require the runtime to create a new queue).

```
context
```
   input. The context in which this queue is being created. Any errors/notifications will be
   reported via callbacks registered in the same context.

```
queue
```
   runtime allocated, output. The queue structure, filled up and returned by the runtime.

```
mailbox
```
   mailbox to gather execution information to be used for debug trap. User may pass NULL
   here if the user doesn't want the mailbox created.

The `hsa_queue_create` API allocates the memory for the queue. Space for `size` number of
packets is allocated by the implementation. The size is required to be aligned with a power of two
number of AQL packets. The *hsa_queue_mailbox_t* structure returned by the queue create call
contains *mailbox_ptr* and *mailbox_signal*. Their purpose is getting execution information when a
debugtrap_u32 HSAIL instruction is used in the user kernel. The user can wait on the
mailbox_signal and process the information in the *mailbox_ptr* as discussed in Section 2.9. The
pointer to the beginning of the memory allocated can be obtained from the queue structure in the
field `base_address`. No memory shall be allocated by an implementation if the queue creation
fails. An implementation may or may not initialize the `hsa_queue` structure if queue creation fails.
Hence the user should rely on the error code to determine if the `hsa_queue` structure is valid.

This service queue is configured when a user mode queue is created. The service queue is visible
to HSA agents through the queue structure `service_queue` field and is serviced by an
appropriate HSA agent. The application may chose to not use a service queue, select the runtime
managed service queue, or a queue managed by the application via. the *hsa_service_queue_type_t*
enumeration input parameter. Address of the service queue associated with the user mode queue
is returned in the queue structure. If there is no associated service queue then the NULL address
will be returned. The API allows different user mode queues to have a different associated service
queue. It also allows for the service queue to be user managed. The API allow allows the user to
specify that runtime return a default shared service queue which is created when the runtime is
initialized.

The `hsa_queue_create` API returns `HSA_STATUS_SUCCESS` when the queue is successfully
created. Otherwise, it can return one of the following status messages:

▷ `HSA_STATUS_ERROR_INVALID_ARGUMENT` error code is returned when the queue size is not a
   power of two, when the error message queue handle is invalid, or the component is not valid.
   This error code is also returned when *queue* is NULL.

▷ `HSA_STATUS_ERROR_OUT_OF_RESOURCES` if there is a failure in allocation of an internal
   structure required by the core runtime library in the context of the queue creation. This error
   may also occur when the core runtime library needs to spawn threads or create internal

OS-specific events. This error is also returned when a service queue or a user mode queue cannot be allocated.

The first ratified version of the SAR specification does not define the `queue_type` and `queue_feature` – they have been marked as fields for future expansion.

The API to destroy a queue is defined as follows:

```
hsa_status_t hsa_queue_destroy(hsa_queue_t *queue);
```

queue
> input. The queue structure that points to the queue that needs to be destroyed, after destruction base_address or the rest of the queue structure, even if cached by the user, are no longer valid.

After queue destruction, it is considered undefined to access the memory pointed to be the `base_address` or the *service_queue*.

2.5.1   Inactivating a Queue

The queue can forcefully be inactivated by the user. This will kill any pending executions and prevent any new packets from being processed. Any more packets written to the queue once it is inactivated will be ignored by the packet processor.

The API for inactivating the queue is defined as follows:

```
hsa_status_t hsa_queue_inactivate(hsa_queue_t *queue);
```

queue
> input. The queue that needs to be inactivated.

2.5.2   Queue Error Reporting, Inactivation and Queue State

The HSA queue structure includes an error message queue, *message_queue_handle*, that the user must initialize and pass as an argument at queue creation. The error message queue may be created by the user using the `hsa_error_message_queue_create` API. The user may also use the default error message queue generated by the `hsa_initialize` API.

There are two primary kinds of errors that impact queue processing and render a queue inactive:

- Errors due to packet processing, such as invalid format, field-value, invalid signal, etc.

- Errors occurring during subsequent resource/dispatch setup or system errors during dispatch.

A queue in HSA, once created, can be in one of the following states: *active, error pending inactive, error inactive* or *destroyed*.

Active Once a queue is successfully created using the `hsa_queue_create` API, it enters an active state, packets can be put on the queue and when the write-index is updated and the doorbell is updated, the packet processor processes the packets. The actual initiation of dispatch may depend on the resources available for the dispatch. Only in the *active* state, writing packets to the queue, updating the write index or ringing the doorbell has any effect. The queue is no longer being monitored by a queue packet processor for new packets in any other state.

Error pending inactive When packet processing or dispatch setup encounters one of the errors described above, the queue packet processor stops packet processing. At this point, there might be in-flight kernels and resources (such as segment allocation) that have been setup for a dispatch but have not yet been freed. So the queue is not entirely inactive, but once the asynchronous activity concludes, it will become inactive. A queue in *error pending inactive* state is not to be considered as destroyed, it still needs to be destroyed so the runtime can reclaim the memory allocated for this queue. If the user provides a callback at queue creation time, the callback is invoked after the queue is marked inactive.

Inactive If all the asynchronous activity concludes, the queue enters the inactive state. A queue can also enter this state when the user explicitly invokes the `hsa_queue_force_inactivate` API (note that the callback implementation for the queue error callback can invoke this API). In an inactive state, the queue structure and its packets may be inspected. Only the packets that are between the read index and the write index in the queue structure are considered to be valid for inspection by the user. The packet processor guarantees that all the packets that have been consumed by the packet processor (see Section 2.6.1) will be signalled with either the completion information or an error. Invocation of `hsa_queue_force_inactivate` API when the queue already is in the inactive state has no effect.

Destroyed The queue has been destroyed by the user. The resources allocated to the queue and the memory for the queue are no longer valid. The queue structure is no longer valid.

A state diagram showing the various states and transitions is shown in Figure 2.2.

The queue will report packet processing or parsing error, system error, dependency resolution error, and signalling error (signal destroyed by the time it needed to be signalled by packet processor).

The queue error reporting infrastructure supports and reports a single error per queue and attempts to inactivate the queue on the first error it encounters.
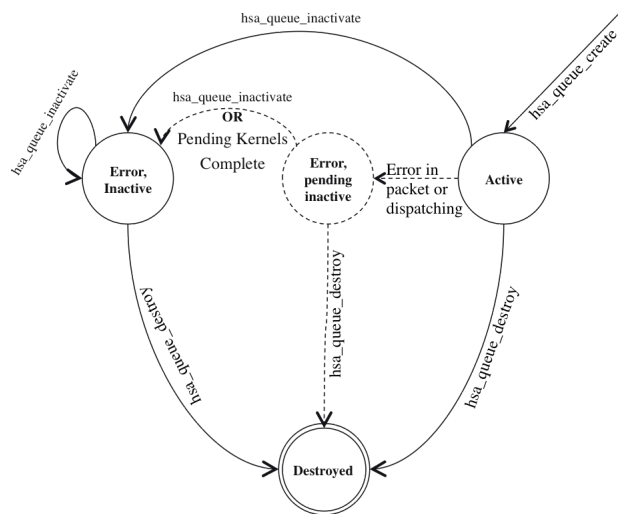
Figure 2.2: Once the queue is created and is active, any error in packet processing takes the queue into pending inactive state where the queue is performing tasks to get to inactive state. Failure during the attempt to inactivate results in queue reaching an error state. A queue that is in active, error or inactive state may be destroyed by using the `hsa_queue_destroy` API provided by the HSA runtime.

2.5.3 Multi-Threaded Queue Access

HSA Core API does not provide explicit API for synchronized access to the queues – the architected queue data structure and read/write index update API are sufficient to allow users to implement thread-safe packet insertion into the queue. Users can use several techniques to support multiple concurrent writers writing AQL packets to the queue. The following example code illustrates one such technique – several other techniques that allow concurrent writes to the queue can be utilized in a similar way.

The sample code below demonstrates a simple reader and writer logic to do a multi-threaded queue access using the queue structure above.

```c
// Read the current queue write offset via intrinsic
tmp_write_index = hsa_queue_get_write_index(q);

// wait until the queue is no longer full.
while(tmp_write_index == read_index + size) {}

// Atomically bump the WriteOffset via intrinsic
if (hsa_queue_cas_write_index(q, tmp_write_index , tmp_write_index +
    1) ==
tmp_write_index)
{
  // calculate index
  uint32_t index = tmp_write_index & (size -1);

  // copy over the packet, the format field is INVALID
  memcpy(q->base_address+index, pkt);

  // Update format field with release semantics
```

```
    q->base_address[index].hdr.format.store(DISPATCH,
        std::memory_order_release);

    // ring doorbell, with release semantics (could also amortize over
    // multiple packets)
    ring_doorbell(write_index+1);
}
```

2.6 Core Runtime Support for AQL

AQL is a command-interface for describing a dispatch or a dependency in a standard format for the
queue packet processor. To match with and support the AQL packet definitions in the HSA SAR,
HSA core base runtime includes structures for different types of AQL packets. SAR defines four
different kinds of AQL packets: invalid, component dispatch, agent dispatch and barrier. There is a
common packet header across these three packet types and is defined by the following structure:

```
typedef struct hsa_aql_packet_header_s {
  uint16_t format:8;
  uint16_t barrier:1;
  uint16_t acquire_fence_scope:2;
  uint16_t release_fence_scope:2;
  uint16_t reserved:3;
} hsa_aql_packet_header_t;
```

format:8
>   8 bits for describing the packet type, 0 for INVALID, 1 for COMPONENT DISPATCH, 2 for
>   BARRIER and 4 for AGENT DISPATCH. All other values are reserverd.

barrier:1
>   If set then processing of packet will only begin when all preceding packets are complete.

acquire_fence_scope:2
>   Determines the scope and type of the memory fence operation applied before the packet
>   enters the active phase. The different values are described in Table 2.2.

release_fence_scope:2
>   Determines the scope and type of the memory fence operation applied after kernel
>   completion but before the packet is completed. The different values are discussed in
>   Table 2.3.

reserved:3
>   must be 0

Acquire Fence Scope	Description
0	None – no fence is applied.
1	The acquire fence makes memory operations made by this HSA Agent prior to launch of this packet, visible to this packet operation.
2	The acquire fence makes memory operations made by HSA Agents prior to launch of this packet, visible to this packet operation.
3	reserved.

Table 2.2: Acquire Fence Scope Values and Actions

Release Fence Scope	Description
0	None – no fence is applied.
1	The release fence is applied to the HSA Agent only.
2	The release fence is applied globally to the HSA System.
3	reserved.

Table 2.3: Release Fence Scope Values and Actions

The `acquire_fence_scope` is used to control the ordering of memory operations before the packet enters the active state. There are three possible values for acquire fence scope. Each of the values defines a particular action by HSA agents and components. The details are described in Table 2.2.

Similarly, the release fence scope, which is also 2 bits, can be used to define the desired memory fence and cache actions at the end of kernel execution, but prior to the packet being marked as complete. Table 2.3 describes the different controls.

The `format` field in the header is used to specify the packet type. Beyond the four packet types defined, all the other packet types are reserved for implementation use. In addition to this, the last 15 bits in the packet header are also reserved for future or implementation specific use. The format field indicates the type of the packet. Of the three packet types, the dispatch and the barrier packet have individual packet-state diagrams that are discussed along with their description.

Invalid AQL packet   Indicates that the packet is not ready to be processed by the packet processor.

2.6.1   Dispatch AQL Packet

Dispatch packet type is used for dispatching a kernel on to a HSA component. The dispatch AQL packet can have five different states: *on queue*, *processing*, *error*, *active* or *complete*. Figure 2.3 shows the different states of a packet and transitions leading to those states.

Figure 2.3: Dispatch Packet State Diagram

On queue state A packet is considered to be in the on queue state once the format of the packet is changed from invalid (a value of 0) to a value of 1, 2 or 3. Any other value for format puts the packet and the queue in error state.

Processing state If this dispatch packet has the barrier bit set, then the processing of this packet occurs only after all prior kernels have completed execution. Otherwise, once the packets prior to this packet are processed, the packet processor begins to process this packet and the packet enters the processing state. From the launch state, two states are possible: error or active.

Error state the packet processor encountered an error processing this packet. This results in a queue error (see Figure 2.2) and the packet enters the error state (the completion object is signalled with error by the packet processor). The following errors are indicated via an error signalled to the completion object: processing parsing error, dependency resolution error, system error and premature termination due to queue inactivation. When the user invokes the **hsa queue inactivate** API or the **hsa queue destroy** API while the packet is in this state, the completion object will be signalled with an error.

Active state If the packet processing is successful and the kernel the packet represents is either executing or queued for execution, the packet enters the active state. From active state, either successful or failed execution both take the packet into the completed state. Alternatively, a user action (see 2.5.1) can also take the packet out of active state into complete state. When the user invokes the hsa queue inactivate API or the hsa queue destroy API while the packet is in this state, the completion object will be signalled with an error.

complete state A packet enters a complete state after its completion signal is signalled (either with success or error).

A dispatch packet is considered processed once the packet processor processes it and makes the queue slot occupied by this packet available. A processed dispatch packet may endure a period of time where it is awaiting its dispatch on to the HSA component. Even such packets awaiting execution are still considered as processed.

The structure for the dispatch AQL packet is shown below:

```
typedef struct hsa_aql_dispatch_packet_s{
  hsa_aql_packet_header_t header;
  uint16_t  workgroup_size_x;
  uint16_t  workgroup_size_y;
  uint16_t  workgroup_size_z;
  uint16_t  reserved2;
  uint32_t  grid_size_x;
  uint32_t  grid_size_y;
  uint32_t  grid_size_z;
  uint32_t  private_segment_size_bytes;
  uint32_t  group_segment_size_bytes;
  uint64_t  kernel_object_address;
  uint64_t  kernarg_address;
  uint64_t  reserved3;
  hsa_signal_handle_t completion_signal;
} hsa_aql_dispatch_packet_t;
```

header
     packet header packet header structure

workgroup_size_x
     x dimension of work-group (measured in work-items).

workgroup_size_y
     y dimension of work-group (measured in work-items).

workgroup_size_z
     z dimension of work-group (measured in work-items).

reserved2
     reserved

grid_size_x
     x dimension of grid (measured in work-items).

grid_size_y
     y dimension of grid (measured in work-items).

grid_size_z
     z dimension of grid (measured in work-items).

private_segment_size_bytes
     Total size in bytes of private memory allocation request (per work-item).

group_segment_size_bytes
     Total size in bytes of group memory allocation request (per work-group).

`kernel_object_address`
>   Address of an object in memory that includes an implementation-defined executable ISA image for the kernel.

`kernarg_address`
>   Address of memory containing kernel arguments.

`reserved3`
>   reserved

`completion_signal`
>   HSA signaling object used to indicate completion of the job.

Segment Sizes

If the kernel being dispatched uses private and group segments, the user is required to specify the sizes of these segments in the AQL dispatch packet. Manually calculating this information is not feasible and requires visual inspection of the user program, which itself may have been generated by a higher-level compiler. Hence the user must rely on the `finalizer` to get the corresponding segment sizes. Further details about determining segment sizes are described in Section A.

Of the other HSA segments, the kernarg segment is also a part of the AQL packet, but as a pointer. This is because the kernarg segment carries the arguments required to execute the kernel being dispatched and must be setup by the user (the layout of this segment is language/finalization specific and associated with the code object generated by finalization) prior to writing the AQL packet to the queue (unlike the group and private segments, whose lifespan spans only the active state of the AQL dispatch packet).

2.6.2 Agent Dispatch AQL Packet

Agent Dispatch AQL packets can be used to do dispatches on the agent queue. The HSA Queue API allows for creation of either agent queues or component queues in the core API (vendor-specific extensions may support queues that allow both agent and component dispatches, but it is not a core feature). The HSA core runtime structure for agent dispatches is defined as follows:

```
typedef struct hsa_aql_agent_dispatch_packet_s{
  hsa_aql_packet_header_t  header;
  uint16_t  type;
  uint32_t  reserved2;
  uint64_t  returnLocation;
  uint64_t  arg0;
  uint64_t  arg1;
  uint64_t  arg2;
  uint64_t  arg3;
```

```
    uint64_t   reserved3;
    uint64_t   completionSignal;
} hsa_aql_agent_dispatch_packet_t;
```

header
    packet header packet header structure

type
    The function to be performed by the destination HSA Agent. The type value is split into
    the following ranges: 0x0000:0x3FFF Vendor specific 0x4000:0x7FFF HSA runtime
    0x8000:0xFFFF User registered function

reserved2
    Must be 0.

returnLocation
    Pointer to location to store the function return value(s) in.

arg0
    64-bit direct or indirect arguments.

arg1
    64-bit direct or indirect arguments.

arg2
    64-bit direct or indirect arguments.

arg3
    64-bit direct or indirect arguments.

reserved3
    Must be 0.

completionSignal
    Address of HSA signaling object used to indicate completion of the job.

2.6.3   Barrier AQL packet

The barrier packet allows the user to specify up to 5 dependencies as hsa_signal objects and
requires the packet processor to resolve them before proceeding. The barrier packet is a blocking
packet, in that the processing of the barrier packet *completes* the packet and its completion object
is signalled. This is unlike a dispatch packet whose completion may occur at some future time after
the packet has finished processing. The HSA core runtime structure for the AQL barrier packet is
shown below:

```
typedef struct hsa_aql_barrier_packet_s{
```

```
    hsa_aql_packet_header_t header;
    uint32_t reserved2;
    uint64_t dep_signal0;
    uint64_t dep_signal1;
    uint64_t dep_signal2;
    uint64_t dep_signal3;
    uint64_t dep_signal4;
    uint64_t reserved3;
    uint64_t completion_signal;
} hsa_aql_barrier_packet_t;
```

header
> packet header packet header structure

reserved2
> reserved

dep_signal0
> The first dependency signal, a negative value means dependency not met and the completion signal for this packet will be set to

dep_signal1
> The first dependency signal, a negative value means dependency not met and the completion signal for this packet will be set to

dep_signal2
> The first dependency signal, a negative value means dependency not met and the completion signal for this packet will be set to

dep_signal3
> The first dependency signal, a negative value means dependency not met and the completion signal for this packet will be set to

dep_signal4
> The first dependency signal, a negative value means dependency not met and the completion signal for this packet will be set to

reserved3
> reserved

completion_signal
> HSA signaling object used to indicate completion of the dependency resolution, success of failure

If any of the dependent signals have been signalled with a negative value, the barrier packet is complete, and will indicate failure in its completion signal. The `completion signal` will be signalled with the error value as discussed in Section 2.4.1.

Figure 2.4: Barrier Packet State Diagram

If the queue is not already in an error state (e.g. the job generating the error was processed in a different queue) then the HSA Packet Processor should consider the error code on the dependent signal to indicate an error in the queue itself and subsequently signal the `error_signal` in the queue.

When all of the dependent signals have been signalled with the value 0, the `completion_signal` will be signalled with the value 0 to indicate a successful completion.

The barrier packet also has a barrier bit that indicates that this packet may only be processed when all previous packets have been marked as completed.

Alike the dispatch packet, the barrier packet can also be in one of the following states: *on queue*, *processing*, *completed, error* or *completed, success*.

On queue state A packet is considered to be in the on queue state once the format of the packet is changed from invalid (a value of 0) to a value of 1 or 2 or 3. Any other value for format puts the packet and the queue in error state.

Processing state If this barrier packet has the barrier bit set, then the processing of this packet occurs only after all prior dispatch packets have completed execution. Otherwise, once the packets prior to this packet are processed, the packet processor begins to process this packet and the packet enters the processing state. From the launch state, two states are possible: completion, error or completion, success.

completed-error The barrier packet reaches this state from the processing state if (a) one of the dependency signals had an error, and (b) if the packet was malformed (e.g. bad signal object or invalid usage of reserved bits). A barrier packet can also reach this state when the user invokes the `hsa_queue_inactivate` API or the `hsa_queue_destroy` API while the packet is in processing state (the completion object will be appropriately signalled with an error).

completed-success The barrier packet had all its dependencies met, its completion object has been signalled with a value of 0.

A state diagram in Figure 2.4 shows these transitions.

2.6.4   Packet Setup Example

This example shows how to setup a component dispatch packet, an agent dispatch packet and a barrier packet. This is work-in-progress – the section needs to be written.

2.7   Memory Registration and Deregistration

2.7.1   Overview

One of the key features of HSA is its ability to share global pointers between the host application and code executing on the component. This ability means that an application can directly pass a pointer to memory allocated on the host to a kernel function dispatched to a component without an intermediate copy, as illustrated by the example shown in Core API Documentation.

When a buffer will be accessed by a kernel running on a HSA device, programmers are encouraged to register the corresponding address range beforehand by using the appropriate HSA core API invocation. While kernels running on HSA devices can access any valid system memory pointer allocated by means of standard libraries (for example, malloc in the C language) without resorting to registration, there might be a performance benefit from registering the buffer with the HSA core component. When an HSA program no longer needs to access a registered buffer in a device, the user should deregister that virtual address range by using the appropriate HSA core API invocation.

The API for registering and deregistering is defined as follows:

```
hsa_status_t hsa_memory_register(void *address,
    size_t size);
```

address
      input. A pointer to the base of the memory region to be registered. If a null pointer is passed, no operation is performed.

size
      input. Requested registration size in bytes. If a size of zero is passed, no operation is performed.

```
hsa_status_t hsa_memory_deregister(void *address);
```

      input. address A pointer to the base of the memory region to be deregistered. If a NULL pointer is passed, no operation is performed.

.

This API returns HSA_STATUS_SUCCESS to indicate that registration/deregistration has been successfully performed. Otherwise, it returns one of the following for status:

▷ HSA_STATUS_ERROR_INVALID_ARGUMENT if a NULL value is passed for *address* or 0 for size.

▷ HSA_STATUS_INFO_NOT_REGISTERED this is applicable to **hsa_memory_deregister** API and indicates that a deregistration is attempted on an address of memory that has not been registered or a part of prior registered ranges.

▷ HSA_STATUS_INFO_OUT_OF_RESOURCES if there is a failure in allocating necessary resources to perform registration. Note that this is just info, since registration doesn't impact functionality, the user can still continue considering this an info.

2.7.2 Usage

A buffer is registered by indicating its starting address and a size. The size does not need to match that of the original allocation. For example:

```
void* ptr = malloc(16);
status = hsa_memory_register(ptr, 8);
if(status == HSA_STATUS_ERROR_INVALID_ARGUMENT)
  handle_error(status);
```

is a valid program. On the other hand:

```
void* ptr = malloc(16);
status = hsa_memory_register(ptr, 20);
if(status == HSA_STATUS_ERROR_INVALID_ARGUMENT)
  handle_error(status);
```

is not a valid program, because we are registering a range that spans several allocations, or might not be entirely allocated.

Registrations can overlap previously registered intervals. A special case of overlapped registrations is multiple registration. If the same interval is registered several times with different sizes, the HSA core component will select the maximum as the size of all the registrations. Therefore, the following program:

```
status = hsa_memory_register(ptr, 8);
if(status == HSA_STATUS_ERROR_INVALID_ARGUMENT)
    handle_error(status);
status = hsa_memory_register(ptr, 16);
if(status == HSA_STATUS_ERROR_INVALID_ARGUMENT)
    handle_error(status);
```

behaves identically to this program:

```
hsa_memory_register(ptr, 16);
if(status == HSA_STATUS_ERROR_INVALID_ARGUMENT)
    handle_error(status);
hsa_memory_register(ptr, 16);
if(status == HSA_STATUS_ERROR_INVALID_ARGUMENT)
    handle_error(status);
```

While the described behavior might seem counterintuitive, consider the following scenario: A pointer is registered twice with different sizes s1 and s2. When the pointer is deregistered, which interval should be deregistered: (p, s1) or (p, s2)? If all the registrations of the same pointer are considered identical by the core runtime, that problem is eliminated.

Deregistering a pointer that has not been previously registered results in an *info* status indicating the same.

The following code snippet revisits the introductory example. The code is almost identical to the original, except that we register the buffers that will be accessed from the device after allocating them, and we deregister all that memory before releasing it. In some platforms, we expect this version to perform better than the original one.

2.8   Core Memory Allocation and Copy API

While a HSA component is capable of accessing pageable system memory by definition, for scenarios where wants memory allocated that has already been registered (combine the allocation with memory registration described in Section 2.7.1), the HSA runtime provides an interface, hsa_memory_allocate to allocate memory that is internally registered by the runtime.

Note that since registration is a performance hint, allocating memory through this API corresponds to the *hint*. The API is defined as follows:

```
hsa_status_t hsa_memory_allocate(size_t size_bytes, void **address);
```

size_bytes
>   input, size of the segment that needs to be allocated in bytes.

address
>   input, address pointer allocated by the user. Dereferenced and assigned to the pointer to the memory allocated for this request.

If the memory is allocated successfully, the API returns HSA_STATUS_SUCCESS. Otherwise, it returns one of the folling error status values:

 ▷ HSA_STATUS_ERROR_OUT_OF_RESOURCES if there is a failure in allocation. This error may also occur when the core runtime library needs to spawn threads or create internal OS-specific events.

 ▷ HSA_STATUS_ERROR_INVALID_ARGUMENT if *address* or *size* are not valid Allocation of size 0 is allowed and returns a NULL pointer.

2.8.1   Allocation of Kernarg Memory

The kernarg memory that AQL packet points to (see Section 2.6) holds information about any arguments required to execute AQL dispatch on a HSA component. While any system memory may be used for kernarg memory, implementation/platform specific optimizations are possible if HSA core runtime provided API are utilized for allocating and copying to the allocated kernarg memory. To facilitate such optimizations, HSA core runtime defines the following API:

```
hsa_status_t hsa_memory_copy_kernarg_to_system(void *dst,
    const void *src,
    size_t size);
```

dst user allocated, output. A valid pointer to the destination array where the content is to be copied.

src input. A valid pointer to the source of data to be copied.

size
>   input. Number of bytes to copy.

```
hsa_status_t hsa_memory_copy_system_to_kernarg(void *dst,
    const void *src,
    size_t size);
```

`dst` user allocated, output. A valid pointer to the destination array where the content is to be copied.

`src` input. A valid pointer to the source of data to be copied.

`size`
  input. Number of bytes to copy.

`signal`
  input. The signal that will be incremented by the runtime when the copy is complete.

If the memory is allocated/copied successfully, the API returns `HSA_STATUS_SUCCESS`. Otherwise, one of the folling error status values are returned:

▷ `HSA_STATUS_ERROR_OUT_OF_RESOURCES` if there is a failure in allocation. This error may also occur when the core runtime library needs to spawn threads or create internal OS-specific events.

▷ `HSA_STATUS_ERROR_INVALID_ARGUMENT` if *address* or *size* are not valid. Allocation of size 0 is allowed and returns a NULL pointer.

@ `HSA_STATUS_ERROR_COPY_SOURCE_INVALID` if the source of the copy operation is invalid.

@ `HSA_STATUS_ERROR_COPY_DESTINATION_INVALID` if the destination of the copy operation is invalid. If both source and destination are invalid, only destination error is reported.

2.8.2   Component Local Memory

Component local memory is a memory type that is dedicated specifically for a particular HSA component. This memory could provide higher bandwidth for component access (than system memory) with the limitation that the host might not be able to access it directly.

HSA runtime provides host interface to allocate/deallocate and access component local memory. The result of the allocation is a pointer to an address in application processes address space, which can be accessed directly by the component during kernel execution.

The API is defined as follows:

```
hsa_status_t hsa_memory_allocate_component_local( const hsa_agent_t
   *component,
   size_t size,
   void **address);
```

`component`
> input. A valid pointer to the HSA device for which the specified amount of global memory
> is to be allocated.

`size`
> input. Requested allocation size in bytes. If size is 0, NULL is returned.

`address`
> output. A valid pointer to the location of where to return the pointer to the base of the
> allocated region of memory.

hsa_status_t **hsa_memory_free_component_local**(void *address);

`address`
> input. A pointer to the address to be deallocated. If the pointer is NULL, no operation is
> performed.

hsa_status_t **hsa_memory_copy_component_local_to_system**(void *dst,
    const void *src,
    size_t size,
    hsa_signal_handle_t signal);

`dst` user allocated, output. A valid pointer to the destination array where the content is to be
> copied.

`src` input. A valid pointer to the source of data to be copied.

`size`
> input. Number of bytes to copy.

`signal`
> input. The signal that will be incremented by the runtime when the copy is complete.

All three API return HSA_STATUS_SUCCESS when the allocate/free/copy is successful. Otherwise,
one of the following status values is returned:

▷ HSA_STATUS_ERROR_OUT_OF_RESOURCES if there is a failure in allocation of an internal
structure required by the core runtime library. This error may also occur when the core runtime

library needs to spawn threads or create internal OS-specific events. The
hsa memory allocate component local and
hsa memory copy component local to system API can return this error.

▷ HSA STATUS ERROR INVALID ARGUMENT if *component*, *address*, *src* or *dst* are not valid and
if the *address* pointer in **hsa memory allocate component local** is NULL. Allocation of
size 0 is allowed.

2.8.3 Usage

Component memory is allocated by indicating the size and the HSA device it corresponds to. For
example, the following code allocates 1024 bytes of device local memory:

```
void* component_ptr = NULL;
hsa_memory_allocate_component_local(1024, component, &component_ptr);
```

To access component memory from the host, the user can call
hsa memory copy component local to host in similar fashion as in memcpy. This interface
allows the user to perform component-to-host memory copy. For example:

```
const size_t DATA_SIZE = 1024;
void* src_ptr = malloc(DATA_SIZE);
void* dest_ptr = NULL;
hsa_memory_allocate_component_local(DATA_SIZE, device, &dest_ptr);
hsa_memory_copy_component_local_to_system(dest_ptr, src_ptr,
    DATA_SIZE);
```

copies 1024 bytes from system to component local memory.

The user should not register or deregister component local memory.

2.9 Execution Control At the Core Level

As per the systems architecture specification, the HSA system must support debugging of a HSAIL
kernel. The HSA Programmers Reference Manual (PRM) describes that the "block" section could
hold debug data and such a section can be placed within a function. This allows the high-level
compiler that generates HSAIL to embed debug specific information. This information makes its
way into the ".debug" section in the brig. This information can be used for associating a HSAIL
level instruction to the higher level functionality. In addition to this, the PRM also discusses the

debugtrap_u32 that halts the current wavefront and transfers control to the agent. The single operand to **debugtrap_u32**, "src" is passed to the agent and can be used to identify the trap.

To support this infrastructure in the runtime, the Core API defines a structure that can be used to exchange information between the kernel executing on the HSA component and the agent.

The core runtime defines a structure, mailbox, whose purpose is to exchange information as a part of execution control. Mailbox is a synchronous communication mechanism between the HSA component and any agents. The HSA component indicates a **debugtrap_u32** or syscall activity by sending a signal indicating it has written to some location in the mailbox.

The HSA PRM defines:

queueactivegroupcount_global_u32 *dest, address* Returns the maximum number of work-groups that can be executed in parallel for dispatches executed on the User Mode Queue with address.

activegroupid index that ranges from 0 through **queueactivegroupcount_global_u32**-1.

The mailbox is an array of structures of size **queueactivegroupcount_global_u32**. Since **activegroupid** is always unique within a queue for any concurrent execution of kernels in that queue, indexing into the mailbox by different work items happens without conflicts. When a workgroup encounters a syscall or a **debugtrap_u32**, the component indexes into its mailbox by accessing it via **activegroupid** from within the **queueptr**. Once the corresponding mailbox is accessed, pertinent information (see structure below) for each work group is populated. Subsequently the component sets the full flag, sends a signal to agent by accessing the *mailbox_signal* inside the queue structure (see Section 2.5), and waits for the full flag to be emptied. The mailbox structure is defined as follows.

```
typedef enum hsa_interrupt_condition_t{
    HSA_DEBUGTRAP=1,
    HSA_SYSCALL=4,
    HSA_OTHER_INTERRUPT=8
} hsa_interrupt_condition_t;
```

HSA_DEBUGTRAP=1
    caused by debugtrap_u32 instruction

HSA_SYSCALL=4
    caused by syscall

HSA_OTHER_INTERRUPT=8
    caused by other interrupt

```
typedef struct hsa_group_execution_info_s{
    hsa_signal_handle_t full_flag;
    uint16_t workgroup_size;
    hsa_interrupt_condition_t *condition;
    uint32_t *workitem_id;
    uint32_t *compute_unit_id;
    uint64_t *aql_packet_ptr;
    uint64_t *virtual_address;
    uint64_t *current_program_counter;
    uint64_t args;
    uint64_t **syscall_output;
} hsa_group_execution_info_t;
```

full_flag
    indicates the mailbox is full and needs to be consumed.

workgroup_size
    the size of the workgroup, all pointers below are arrays of that size.

condition
    what caused this execution to stop.

workitem_id
    the flattend workitem IDs, array[workgroup_size].

compute_unit_id
    the ID of the compute unit, array[workgroup_size].

aql_packet_ptr
    pointer to the AQL packet, array[workgroup_size].

virtual_address
    any pertinent virtual address, array[workgroup_size].

current_program_counter
    the current program counter, array[workgroup_size].

args
    location to where the arguments have been stored. The size and contents are written by the component and need to be decoded by the agent when reading this.

syscall_output
    if the condition is syscall, location to where the outputs need to be stored. This is array[workgroup_size].

The Agent waits on the signal, processes the mailbox, and clears the full flag.

If this kernel had a debugtrap_u32, a simple check for debugtrap can be written the following way:

```
    assert (HSA_STATUS_SUCCESS ==
        hsa_signal_wait(queue_ptr->mailbox_signal,
          HSA_GREATER_EQUALS, 1));
    for(i = 0; i < queue->active_group_global_count; i++) {
      assert(HSA_STATUS_SUCCESS ==
          hsa_signal_query_acquire(queue->mailbox[i].full_flag,value);
      if(value.u64_value == 1){
        for(j=0; j < queue->mailbox[i].workgroup_size; j++) {
          printf("\n workitemid=%d computeunitid=%d\n",
                queue->mailbox[i].workitem_id[j],
                queue->mailbox[i].compute_unit_id[j]);
                //here we can check PC, etc.
        }
        hsa_signal_send(queue_ptr->mailbox[i].full_flag, 0);
      }
    }
```

2.10   Agent Dispatch Support at the Core Level

The core runtime supports agent dispatches from an HSA component/Agent. The runtime defines a default service queue for every user mode queue created by the user. This default service queue is available to the HSAIL program HSAIL programs and the user applications may submit agent dispatch packets to the service queue or any user mode queue. The service queue shares the same structure as the regular HSA queue. The default service queues are monitored by the runtime.

```
 hsa_status_t hsa_register_agent_dispatch_callback(hsa_queue_t
    *agent_dispatch_queue,
    void (*agent_dispatch_callback)(uint64_t arg0, uint64_t arg1,
      uint64_t
      arg2, uint64_t arg3, uint64_t returnaddr),
    hsa_runtime_context_t *context);
```

notification_callback
      input. the callback that the user is registering, the callback is called with 5 64 bit args (see Section **??**) as a parameter.

2.11   Extensions to the Core Runtime API

When an implementor of the core runtime specification is not supporting any of the extension API, they will return HSA_STATUS_ERROR_EXTENSION_NOT_SUPPORTED as a return status for that API.

Individual vendors may define vendor extensions to HSA core runtime, or multiple vendors may collaborate to define an extension. The difference is in the naming scheme used for the symbols (defines, structures, functions, etc.) associated with the function:

- Symbols for single-vendor extensions that are defined in the global namespace must use the following naming convention:

  - *hsa_svext_<COMPANY_NAME >_*. For example, a company "ACME" defining a single-vendor extension would use the prefix *hsa_ext_acme_*. Company names must be registered with the HSA Foundation, must be unique, and may be abbreviated to improve the readability of the symbols.

- Symbols for multi-vendor extensions that are defined in the global namespace must use the following naming convention:

  - *hsa_ext_* For example, if another company embraces extension in the example above from Company "ACME", the resulting symbols would use the prefix *hsa_mvext_*.

Any constant definitions in the extension (#define/enumerations) use the same naming convention, except using all capital letters. So, using the single-vendor extension example from above, the associated defines and enumerations would have the prefix HSA_EXT_ACME_.

The symbols for all vendor extensions (both single-vendor and multi-vendor) are captured in the file **hsa/vendor_extensions.h**. This file is maintained by the HSA Foundation. This file includes the enumeration hsa_vendor_extension_t which defines a unique code for each vendor extension and multi-vendor extension. Vendors can reserve enumeration encodings through the HSA Foundation. Multi-vendor enumerations begin at the value of 1000000. For example, using the examples above, the hsa_vendor_extension_t enumeration might be:

```
typedef enum hsa_vendor_extension_s {
  HSA_SVEXT_START= 0,
  HSA_SVEXT_ACME_FOO = 1,
  HSA_SVEXT_ACME_ANOTHER_EXT = 2,
  HSA_MVEXT_START = 1000000,
  HSA_MVEXT_FOO   = 1000001
}hsa_vendor_extension_t;
```

HSA_SVEXT_START= 0
      start of the single vendor extension range

```
HSA_SVEXT_ACME_FOO = 1
      Company ACME, starts with FOO symbol

HSA_SVEXT_ACME_ANOTHER_EXT = 2
      Company ACME has another_ext symbol

HSA_MVEXT_START = 1000000
      multi vendor extension starts at 1000000

HSA_MVEXT_FOO = 1000001
      multivendor extension has a symbol foo
```

HSA defines the following query function for vendor extensions:

```
hsa_status_t hsa_vendor_extension_query(hsa_vendor_extension_t
    extension, void *extension_structure);
```

```
extension
      input. The vendor extention that is being queried
```

This API returns HSA_STATUS_SUCCESS if the extension is supported. Additionally, *extension_structure* is written with extension-specific information such as version information, function pointers, and data values. **hsa/vendor_extension.h** defines a unique structure for each extension. If the vendor extension is not supported, HSA_STATUS_ERROR_EXTENSION_UNSUPPORTED is returned, and extension_structure is not modified.

2.11.1 Example Definition And Usage of an Extension

An example that shows a hypothetical single-vendor extension "Foo" registered by company "ACME". The example includes four defines and two API functions. Note the use of the structure hsa_svext_acme_foo_t and how this interacts with the **hsa_query_vendor_extension** API call.

```
//---
// Sample hsa/vendor_extensions.h
// Company name is "ACME" and extension is "Foo"

#define HSA_EXT_ACME_MYDEFINE1 0x1000
#define HSA_EXT_ACME_MYDEFINE2 0x0100
#define HSA_EXT_ACME_MYDEFINE3 0x0010
#define HSA_EXT_ACME_MYDEFINE4 0x0001
```

```
// The structure which defines the version, functions, and data for
// the extension:
typedef struct hsa_ext_acme_foo_s {
    int major_version;  // major version number of the extension.
    int minor_version;  // minor version number of the extension.


    // Function pointers:
    int (*function1) ( int p1, int *p2, float p3, int p4);
    int (*function2) ( int* p1, int p2);

    // Data:
    unsigned foo_data1;
} hsa_ext_acme_foo_t;

main() {
    struct hsa_ext_acme_foo_t acmeFoo;

    hsa_status_t status = hsa_query_vendor_extension( HSA_EXT_ACME_FOO,
        &acmeFoo);
    if (status == HSA_STATUS_SUCCESS) {
        (*(acmeFoo.function2))(0, 0);
    }
}
```

Appendices

A

Compilation Unit, Finalizer and ISA Linking

A.1 Finalization, Compilation Unit, Code, Debug and Symbol Objects

Compilation support in the HSA core runtime comprises of a finalization step. The primary functionality of this step is to translate the HSAIL to a component specific instruction set and produce a compilation unit. It resolves user defined symbols that are not already bound via callbacks provided by the user.

HSA components may support HSAIL natively or may have a native ISA that the HSAIL needs to be translated into. However, compilation is a necessary step and is required despite a HSA components' native support of HSAIL. This is because of two primary reasons:

1. The HSA kernel code objects (accessible via the `hsa_compilationunit_code_t` structure, discussed in Section A.1.2) are an output of the finalization process and required as an input for the AQL dispatch packet; to obtain values from the finalizer for segment sizes etc. and also to act as a container for component specific execution information.

2. In order to support kernel dispatches from the HSA component, the kernel code object must reside in a memory layout specification since dispatches initiated from components are able to use memory operations to get the information necessary for a dispatch.

Before discussing BRIG, etc., a constant `HSA_CODE_VERSION` is defined by the runtime to represent the version of the code object formats being used in this instance of the runtime.

```
typedef uint32_t hsa_code_version32_t;
enum hsa_code_version_t {
  HSA_CODE_VERSION = 0
};
```

A.1.1   BRIG and .directive Section

The core runtime accepts HSAIL programs coded in the BRIG binary format, as defined in the HSA PRM [**?**], for its finalization process. BRIG is a binary format defined by the HSA PRM and includes 5 different sections, *.string*, *.directive*, *.code*, *.operand*, and *.debug*. More information on these sections is described in Section 19.1 of the HSA PRM. The core runtime structure that represents a BRIG is named `hsa_brig_t`. This structure is an in-memory representation of the BRIG. It is defined as follows:

```
typedef struct hsa_brig_s{
    uint8_t *string_section;
    uint8_t *directive_section;
    uint8_t *code_section;
    uint8_t *operand_section;
} hsa_brig_t;
```

string_section
>   From PRM: string section, containing all character strings and byte data used in the
>   compilation unit.

directive_section
>   The directives, which provide information for the finalizer. The directives do not generate
>   code.

code_section
>   All of the executable operations. Most operations contain offsets to the .operand section.

operand_section
>   The operands, such as immediate constants, registers, and address expressions, that
>   appear in the operations.

Of the different sections in BRIG, the .directive section provides information to the finalizer on functions, kernels, and global declarations, etc. The symbols in the .directive section have defined placement rules (see PRM for more information). For example, immediately after a function or kernel directive, BRIG requires the directives that describe the arguments to be in a certain order. Return arguments are first, followed by input arguments, followed by the directives that apply only to the function or kernel.

The `hsa_brig_t` structure has the base address of the directive section. A directive for any symbol is represented using an offset into the directive section. HSA core runtime defines a type `hsa_brig_directive_offset_t` to represent the .directive section offset. It is typedef to an unsigned 32 bit integer and is defined by all implementations as follows:

```
typedef uint32_t hsa_brig_directive_offset_t;
```

There are different types of directives specified in the PRM. Of these, the control directives are a means to allow implementations to pass information to the finalizer via HSAIL. HSA runtime defines a structure `hsa_control_directives_t` to represent the values of control directives both at finalization time and to record information in the kernel code object. Control directives may also be specified within the HSAIL code. When conflicting values are specified for a particular directive specified in HSAIL and at finalization time, the runtime will do one of the following: (a) perform a union (OR operation) when possible (e.g. when the value represents a bit field).

(b) when a union is not meaningful, the runtime will require that the value provided at finalization time via the `hsa_control_directives_t` structure match the value for this directive in the HSAIL kernel.

The `hsa_control_directives_t` structure is defined as follows:

```
typedef struct hsa_control_directives_s {
  hsa_control_directive_present64_t enabled_control_directives;
  hsa_exception_kind16_t enable_break_exceptions;
  hsa_exception_kind16_t enable_detect_exceptions;
  uint32_t max_dynamic_group_size;
  uint32_t max_flat_grid_size;
  uint32_t max_flat_workgroup_size;
  uint32_t requested_workgroups_per_cu;
  hsa_dim3_t required_grid_size;
  hsa_dim3_t required_workgroup_size;
  uint8_t required_dim;
  uint8_t reserved[75];
} hsa_control_directives_t;
```

`enabled_control_directives`
> If the value is 0 then there are no control directives specified and the rest of the fields can be ignored. The bits are accessed using the hsa_control_directives_present_mask_t. Any control directive that is not enabled in this bit set must have the value of all 0s.

`enable_break_exceptions`
> If enable break exceptions is not enabled in *enabled_control_directives*, then must be 0, otherwise must be non-0 and specifies the set of HSAIL exceptions that must have the BREAK policy enabled. If the HSAIL kernel being finalized has any enablebreakexceptions control directives, then the values specified by this argument are unioned with the values in these control directives. If any of the functions the kernel calls have an enablebreakexceptions control directive, then they must be equal or a subset of, this union.

`enable_detect_exceptions`
> If enable detect exceptions is not enabled in *enabled_control_directives*, then must be 0, otherwise must be non-0 and specifies the set of HSAIL exceptions that must have the DETECT policy enabled. If the kernel being finalized has any enabledetectexceptions control directives, then the values specified by this argument are unioned with the values

---

in these control directives. If any of the functions the kernel calls have an enabledetectexceptions control directive, then they must be equal or a subset of, this union.

max_dynamic_group_size

If max dynamic group size is not enabled in *enabled_control_directives* then this must be 0, and any amount of dynamic group segment can be allocated for a dispatch, otherwise the value specifies the maximum number of bytes of dynamic group segment that can be allocated for a dispatch. If the kernel being finalized has any maxdynamicsize control directives, then the values must be the same, and must be the same as this argument if it is enabled. This value can be used by the finalizer to determine the maximum number of bytes of group memory used by each work-group by adding this value to the group memory required for all group segment variables used by the kernel and all functions it calls, and group memory used to implement other HSAIL features such as fbarriers and the detect exception operations. This can allow the finalizer to determine the expected number of work-groups that can be executed by a compute unit and allow more resources to be allocated to the work-items if it is known that fewer work-groups can be executed due to group memory limitations.

max_flat_grid_size

If this is is not enabled in *enabled_control_directives* then must be 0, otherwise must be greater than 0. See HSA Programmer's Reference Manual description of maxflatgridsize control directive.

max_flat_workgroup_size

If this is is not enabled in *enabled_control_directives* then must be 0, otherwise must be greater than 0. See HSA Programmer's Reference Manual description of maxflatgridsize control directive.

requested_workgroups_per_cu

If this is is not enabled in *enabled_control_directives* then must be 0 and the finalizer may generate ISA that could result in any number of work-groups executing on a single compute unit. Otherwise, the finalizer will *attempt* to generate ISA that will allow the specified number of work-groups to execute on a single compute unit. This is only a hint and can be ignored by the finalizer. If the kernel being finalized, or any of the functions it calls, has the same control directive, then the values must be the same or the finalization can fail. This can be used to determine the number of resources that should be allocated to a single work-group and work-item.

required_grid_size

If not enabled then all elements for Dim3 must be 0, otherwise every element must be greater than 0. See HSA Programmer's Reference Manual description of requiredgridsize control directive.

required_workgroup_size

If not enabled then all elements for Dim3 must be 0, and the produced code can be dispatched with any legal work-group range consistent with the dispatch dimensions. Otherwise, the code produced must always be dispatched with the specified work-group range. No element of the specified range must be 0. It must be consistent with

required_dimensions and max_flat_workgroup_size. If the kernel being finalized, or any of the functions it calls, has a requiredworkgroupsize control directive, then the values must be the same. Specifying a value can allow the finalizer to optimize work-group id operations, and if the number of work-items in the work-group is less tha the WAVESIZE then barrier operations can be optimized to just a memory fence.

`required_dim`
> If required_dim is not enabled then must be 0 and the produced kernel code can be dispatched with 1, 2 or 3 dimensions. If enabled then the value is 1..3 and the code produced must only be dispatched with a dimension that matches. Other values are illegal. If the kernel being finalized, or any of the functions it calls, has a requireddimsize control directive, then the values must be the same. This can be used to optimize the code generated to compute the absolute and flat work-group and work-item id, and the dim HSAIL operations.

`reserved[75]`
> Reserved. Must be 0.

Where, the `hsa_control_directive_present64_t` is defined as a 64bit unsigned integer.

```
typedef uint64_t hsa_control_directive_present64_t;
```

The enumeration `hsa_control_directive_present64_t` is a bit set indicating which control directives have been specified. It is accessible via a mask, `hsa_control_directives_present_mask_t`, which is defined as follows:

```
typedef enum hsa_control_directive_present_mask_t {
  HSA_CONTROL_DIRECTIVE_ENABLE_BREAK_EXCEPTIONS = 0,
  HSA_CONTROL_DIRECTIVE_ENABLE_DETECT_EXCEPTIONS = 1,
  HSA_CONTROL_DIRECTIVE_MAX_DYNAMIC_GROUP_SIZE = 2,
  HSA_CONTROL_DIRECTIVE_MAX_FLAT_GRID_SIZE = 4,
  HSA_CONTROL_DIRECTIVE_MAX_FLAT_WORKGROUP_SIZE = 8,
  HSA_CONTROL_DIRECTIVE_REQUESTED_WORKGROUPS_PER_CU = 16,
  HSA_CONTROL_DIRECTIVE_REQUIRED_GRID_SIZE = 32,
  HSA_CONTROL_DIRECTIVE_REQUIRED_WORKGROUP_SIZE = 64,
  HSA_CONTROL_DIRECTIVE_REQUIRED_DIM = 128,
  HSA_CONTROL_DIRECTIVE_REQUIRE_NO_PARTIAL_WORKGROUPS = 256
} hsa_control_directive_present_mask_t;
```

`HSA_CONTROL_DIRECTIVE_ENABLE_BREAK_EXCEPTIONS = 0`
> mask that indicates break on exceptions is required by the user, the kernel pauses execution and the queue mailbox signal is signaled and mailbox updated.

`HSA_CONTROL_DIRECTIVE_ENABLE_DETECT_EXCEPTIONS = 1`
> says that exeptions are recorded

```
HSA_CONTROL_DIRECTIVE_MAX_DYNAMIC_GROUP_SIZE = 2
     says that max for dynamic group size is specified

HSA_CONTROL_DIRECTIVE_MAX_FLAT_GRID_SIZE = 4
     if enabled

HSA_CONTROL_DIRECTIVE_MAX_FLAT_WORKGROUP_SIZE = 8
     if enabled

HSA_CONTROL_DIRECTIVE_REQUESTED_WORKGROUPS_PER_CU = 16
     if enabled

HSA_CONTROL_DIRECTIVE_REQUIRED_GRID_SIZE = 32
     if enabled

HSA_CONTROL_DIRECTIVE_REQUIRED_WORKGROUP_SIZE = 64
     if enabled

HSA_CONTROL_DIRECTIVE_REQUIRED_DIM = 128
     if enabled

HSA_CONTROL_DIRECTIVE_REQUIRE_NO_PARTIAL_WORKGROUPS = 256
     if enabled
```

User can choose to either break on exceptions or just detect them. The PRM defines the policy to be exception-type specific, i.e. different IEEE exceptions supported by HSA (see the definition of the enumeration `hsa_exception_kind_mask_t` below) can be handled with different policies (BREAK vs. DETECT). The *enable_break_exceptions* field specifies the set of HSAIL exceptions that must have the BREAK policy enabled. It is possible that on some systems, enabling exceptions may result in lower code performance. If the kernel being finalized has any `enablebreakexceptions` control directives in HSAIL, then the runtime performs a union (OR operation) of values specified by this argument with the values in HSAIL control directives. If any of the functions the kernel calls have an enablebreakexceptions control directive, then they must be equal to or a subset of this union.

```
typedef enum hsa_exception_kind_mask_s {
  HSA_EXCEPTION_INVALID_OPERATION = 1,
  HSA_EXCEPTION_DIVIDE_BY_ZERO = 2,
  HSA_EXCEPTION_OVERFLOW = 4,
  HSA_EXCEPTION_UNDERFLOW = 8,
  HSA_EXCEPTION_INEXACT = 16
}hsa_exception_kind_mask_t ;
```

```
HSA_EXCEPTION_INVALID_OPERATION = 1
     IEEE 754 INVALID operation exception
```

```
HSA_EXCEPTION_DIVIDE_BY_ZERO = 2
     an operation on finite operands gives an exact infinite result

HSA_EXCEPTION_OVERFLOW = 4
     a result is too large to be represented correctly

HSA_EXCEPTION_UNDERFLOW = 8
     a result is very small (outside the normal range) and inexact

HSA_EXCEPTION_INEXACT = 16
     returns correctly rounded result by default
```

A.1.2 Code Objects

There are different code objects defined by the HSA core runtime specification in support of HSAIL: `hsa_compilationunit_code_t`, `hsa_kernel_code_t` and `hsa_function_code_t`. All of them are in memory and can be relocatable and/or position independent (which indicates that the object can be deep-copied to other memory locations for execution). The HSA runtime provides a query API to verify if a particular component supports position independent code objects (see Section 2.3).

The `hsa_compilationunit_code_t` is the header for the code object produced by the Finalizer and contains information that applies to all code entities in the compilation unit.

Since core runtime does not define a file format container, the core runtime provides API to work with HSAIL programs encoded in the BRIG binary format and supports generation of code objects that include kernels to be executed in the HSA component and binding of unresolved symbols associated with the code object generated.

Finalizer allocates a single contiguous area of memory to hold the generated code for all the code objects.

The structure of this contiguous area is as follows: Starting at offset 0, the "header" of this contiguous area is defined by the `hsa_compilationunit_code_t` structure. The `hsa_compilationunit_code_t` structure in turn contains an offset to an array of `hsa_code_entry_t`, one entry per code entity that the finalizer has produced code for. Each `hsa_code_entry_t` variable contains an offset to a `hsa_*code_t` object that describes that code entity (function/kernel/etc).

The kinds of code objects that can be contained in a `hsa_compilationunit_code_t` is defined by the following structure:

```
typedef enum hsa_code_kind_t {
  HSA_CODE_NONE = 0,
  HSA_CODE_KERNEL = 1,
  HSA_CODE_FUNCTION = 2,
  HSA_CODE_RUNTIME_FIRST = 0x40000000,
  HSA_CODE_VENDOR_FIRST = 0x80000000,
} hsa_code_kind_t;
```

```
HSA_CODE_NONE = 0
      Not a code object

HSA_CODE_KERNEL = 1
      HSAIL kernel that can be used with an AQL dispatch packet.

HSA_CODE_FUNCTION = 2
      HSAIL function.

HSA_CODE_RUNTIME_FIRST = 0x40000000
      HSA runtime code objects. For example, partially linked code objects.

HSA_CODE_VENDOR_FIRST = 0x80000000
      Vendor specific code objects.
```

The `hsa_code_entry_t` structure is defined as follows:

```
typedef struct hsa_code_entry_s {
  uint64_t code_id;
  int64_t code_byte_offset;
} hsa_code_entry_t;
```

```
code_id
      ID of the entity that generated the code. For HSAIL will be the BRIG directive offset of the
      kernel or function declaration. The array of hsa_code_entry_t are required to be ordered in
      ascending code_id to allow faster lookup.

code_byte_offset
      Byte offset from start of hsa_compilationunit_code_t to corresponding hsa_code_t. Every
      hsa_code_t starts with a common hsa_code_t, and hsa_code_t.code_type indicates what
      specific hsa_code_t it is.
```

The current version number of the HSA code object format is defined as follows:

```
typedef enum hsa_code_version_t {
  HSA_CODE_VERSION = 0
} hsa_code_version_t;
```

```
HSA_CODE_VERSION = 0
      Code version
```

Every `hsa_*_code_t` code objects start with a common header that also contains what kind of code object it is. The common header, `hsa_code_t`, is defined as follows:

```
typedef struct hsa_code_s {
  hsa_code_version32_t code_version;
  uint32_t struct_byte_size;
  int64_t compilationunit_byte_offset;
  hsa_code_kind32_t code_type;
} hsa_code_t;
```

code_version
> The code format version. The version of this defintion is specified by HSA_CODE_VERSION. Must match the value in the hsa_compilationunit_code_t that contains it.

struct_byte_size
> The byte size of the struct that contains this hsa_code_t. Must be set to sizeof(hsa_code_t). Used for backward compatibility.

compilationunit_byte_offset
> Offset from base of hsa_code_t to compilationunit_code_t that contains this hsa_code_t to the base of this hsa_code_t. Can be used to navigate back to the enclosing compilation unit. Since hsa_compilationunit_code_t is always at offset 0, this value must be negative.

code_type
> Type of code object.

A bit set of flags providing information about the code in a compilation unit. Unused flags must be 0. The `hsa_code_properties32_t` must be used as a type for this flag. The values/mask currently supported is defined as follows:

```
typedef enum hsa_code_properties_mask_t {
  HSA_CODE_PROPERTY_PIC = 1 << 0
} hsa_code_properties_mask_t;
```

HSA_CODE_PROPERTY_PIC = 1 << 0
> The code is position independent (can be executed at any address that meets the alignment requirement).

The `hsa_compilationunit_code_t` structure is defined as follows:

```
typedef struct hsa_compilationunit_code_s{
  hsa_code_version32_t code_version;
  uint32_t struct_byte_size;
  char component_vendor[16];
  char component_name[16];
  int64_t code_entry_byte_offset;
  uint32_t code_entry_count;
  hsa_powertwo8_t code_alignment;
  uint8_t reserved[3];
  uint64_t code_size_bytes;
  uint64_t code_base_address;
  hsa_code_properties32_t code_properties;
  uint32_t hsail_version_major;
  uint32_t hsail_version_minor;
  hsa_profile8_t hsail_profile;
  hsa_machine_model8_t hsail_machine_model;
  hsa_target_options16_t hsail_target_options;
} hsa_compilationunit_code_t;
```

code_version
>   The code format version. The version of this defintion is specified by
>   HSA_CODE_VERSION.

struct_byte_size
>   The byte size of this struct. Must be set to sizeof(hsa_compilationunit_code_t). Used for
>   backward compatibility.

component_vendor[16]
>   The vendor of the HSA Component on which this Kernel Code object can execute.
>   ISO/IEC 624 character encoding must be used. If the name is less than 16 characters
>   then remaining characters must be set to 0.

component_name[16]
>   The vendor's name of the HSA Component on which this Kernel Code object can execute.
>   ISO/IEC 646 character encoding must be used. If the name is less than 16 characters
>   then remaining characters must be set to 0.

code_entry_byte_offset
>   Byte offset from start of hsa_compilationunit_code_t to an array of code_entry_count
>   elements of type hsa_code_entry_t. Since hsa_compilationunit_code_t is always at offset 0,
>   this value must be positive.

code_entry_count
>   Number of code entries in this compilation unit.

code_alignment
>   The required alignment of this hsa_compilationunit_code_t expressed as a power of 2. The

Finalizer must set this to the value required by the HSA component it will execute on and the assumptions of the machine code it contains.

`reserved[3]`
Must be 0.

`code_size_bytes`
The size of the single contiguous block of memory which includes this hsa_compilationunit_code_t header and all following hsa_ _code_t and associated machine code.

`code_base_address`
The base address that this hsa_compilationunit_code_t must be allocated in order to execute the code it contains. The address must be a multiple of the alignment specified by the alignment field. If the code is position independent (can be executed at any address that meets the alignment requirement), then this field must be 0.

`code_properties`
A bit set of flags providing inforamtion about the code in this compilation unit. Unused flags must be 0.

`hsail_version_major`
The HSAIL major version. This information is from the HSAIL version directive. If this hsa_compilationunit_code_t is not generated from an HSAIL compilation unit then must be 0.

`hsail_version_minor`
The HSAIL minor version. This information is from the HSAIL version directive. If this hsa_compilationunit_code_t is not generated from an HSAIL compilation unit then must be 0.

`hsail_profile`
The HSAIL profile defines which features are used. This information is from the HSAIL version directive. If this hsa_compilationunit_code_t is not generated from an HSAIL compilation unit then must still indicate what profile is being used.

`hsail_machine_model`
The HSAIL machine model gives the address sizes used by the code. This information is from the HSAIL version directive. If not generated from an HSAIL compilation unit then must still indicate for what machine mode the code is generated.

`hsail_target_options`
The HSAIL target features. There are currently no target options so this field must be 0. If target options are added they will be specified by the HSAIL version directive. If this hsa_compilationunit_code_t is not generated from an HSAIL compilation unit then must be 0.

Both the `hsa_compilationunit_t` and `hsa_*_code_t` objects can have implementation define data towards the end of the structure. All elements in the contiguous location being accessed by

offsets enables such definition. This also allows the exact position of the various objects in the contiguous memory area to be implementation defined as long as memory alignment requirements are met.

Many of the `hsa_*_code_t` objects include a size field which is required to be set to the size of the structure. This allows forward compatibility and allows for structure definitions to change or include implementation specific information.

The `hsa_kernel_code_t` is required for all component dispatches and is referred to by the dispatch AQL packets placed in the HSA queue. This allows an implementation to rely on all such objects being the same size for more efficient navigation to the implementation specific data without need to first read the object's size field.

The `hsa_kernel_code_t` is the output of **hsa_finalize_brig** API and is defined as follows.

```
typedef struct hsa_kernel_code_s{
  hsa_code_t code;
  uint32_t workgroup_group_segment_byte_size;
  uint64_t kernarg_segment_byte_size;
  uint32_t workitem_private_segment_byte_size;
  uint32_t workgroup_fbarrier_count;
  hsa_control_directives_t control_directive;
  hsa_powertwo8_t wavefront_size;
  hsa_powertwo8_t kernarg_segment_alignment;
  uint8_t optimization_level;
  uint8_t reserved1[3];
} hsa_kernel_code_t;
```

code
>   Common header that all code objects start with. code.type must be
>   HSA_CODE_KERNEL.

workgroup_group_segment_byte_size
>   The amount of group segment memory required by a work-group in bytes. This does not
>   include any dynamically allocated group segment memory that may be added when the
>   kernel is dispatched.

kernarg_segment_byte_size
>   The size in bytes of the kernarg segment that holds the values of the arguments to the
>   kernel.

workitem_private_segment_byte_size
>   The amount of memory required for the combined private, spill and arg segments for a
>   work-item in bytes.

workgroup_fbarrier_count
>   Number of fbarrier's used in the kernel and all functions it calls. If the implemenation uses
>   group memory to allocate the fbarriers then that amount must already be included in the
>   *workgroup_group_segment_byte_size* total.

control_directive
>   The values are the actually values used by the finalizer in generating the code. This may be the union of values specified as finalizer arguments and explicit HSAIL control directives. If a finalizer implementation ignores a control directive, and not generate constrained code, then the control directive will not be marked as enabled even though it was present in the HSAIL or finalizer argument. The values are intended to reflect the constraints that the code actually requires to correctly execute, not the values that were actually specified at finalize time.

wavefront_size
>   Wavefront size expressed as a power of two. Must be a power of 2 in range 1..64 inclusive. Used to support runtime query that obtains wavefront size, which may be used by application to allocated dynamic group memory and set the dispatch work-group size.

kernarg_segment_alignment
>   The maximum byte alignment of variables used by the kernel in the specified memory segment. Expressed as a power of two. Must be at lease HSA_POWERTWO_16.

optimization_level
>   The optimization level specified when the kernel was finalized.

reserved1[3]
>   Reserved. Must be 0. Component specific fields can follow this field.

A.1.3   Finalize BRIG API

The **hsa_finalize_brig** API accepts hsa_brig_t as an input and produces relocatable code object in which global and group symbols are bound to actual, component recognizable, addresses. A symbol in the finalization step can be a variable, a function, image or a sampler. When finalization occurs, a hsa_kernel_code_t, representing the kernel that needs to be executed on the component is generated. However, all the symbols referenced in the kernel being finalized may not be resolved with the information provided at its finalization. User is allowed to define callbacks that can resolve the symbols declared in the global segment.

The callback, hsa_map_symbol_address_t, is defined as follows:

```
typedef hsa_status_t (*hsa_map_symbol_address_t)(
  hsa_finalize_compilationunit_caller_t caller,
  hsa_brig_directive_offset_t symbol_directive,
  uint64_t* address);
```

The finalization step, when successfully executed, can have two distinct outputs. The first is the compilation unit code object, hsa_compilationunit_code_t which is discussed in Section A.1.2. The second output is of type hsa_compilationunit_debug_t which is only generated when the code is compiled with a debug option. This debug information is currently implementation defined.

Since the contiguous memory that *compilationunit_code* represents and the memory for *compilationunit_debug* need to be allocated, the user is expected to provide callbacks for memory allocation.

The callback is required to have the following signature:

```
typedef hsa_status_t (*hsa_alloc_t)(
  hsa_finalize_brig_caller_t caller,
  size_t byte_size,
  size_t byte_alignment,
  void** address);
```

Where caller is the opaque pointer passed to the **hsa_finalize_brig** that is calling back this function. *byte_size* is the size in bytes of the memory to be allocated. *byte_alignment* is the required byte alignment of the memory allocated. Must be a power of 2. *address* is pointer to location that will be updated with address of allocated memory if successful, or NULL if not successful. *return value* is the HSA status of the allocation.

The callback for the allocation of *compilationunit_debug* is optional and is required when the user passes in a flag at finalization to generate debug information. This callback, when defined, must have the same signature as hsa_alloc_t above.

```
hsa_status_t hsa_finalize_brig( hsa_finalize_compilationunit_caller_t
   caller,
hsa_agent_t* component,
hsa_brig_t* brig,
size_t code_count,
hsa_brig_directive_offset_t* code_directive,
hsa_control_directives_t* control_directives,
hsa_map_symbol_address_t map_symbol_address,
hsa_alloc_t allocate_compilationunit_code,
hsa_alloc_t allocate_compilationunit_debug,
uint8_t optimization_level,
const char* options,
hsa_compilationunit_code_t** compilationunit_code,
hsa_compilationunit_debug_t** compilationunit_debug);
```

caller
    Opaque pointer and will be passed to all call back functions made by this call of the finalizer.

component
    input. A valid pointer to the hsa_agent_t.

brig
    input. A pointer to the in memory BRIG structure.

`code_count`
    The number of kernels plus functions to produce hsa_kernel_code_t and
    hsa_function_code_t objects for in the generated hsa_compilationunit_code_t.

`code_directive`
    A pointer to an array with code_count entries of hsa_brig_directive_offset_t. Each entry is
    the offset in the directive section of the passed brig of a kernel or function definition.
    These will be the kernels and functions that will have code objects generated in the
    produced hsa_compilationunit_code_t

`control_directives`
    The control directives that can be specified to influence how the finalizer generates code.
    If NULL then no control directives are used.

`map_symbol_address`
    Used by the finalizer to obtain the segment address for global segment symbols. The
    value of caller will be passed to every call.

`allocate_compilationunit_code`
    The callback function that the finalizer will use to allocate the contiguous block of memory
    that will be used for the hsa_compilationunit_code_t that is returned. It is the responsibility
    of the call of the finalizer to deallocate this memory, even if the finalizer does not report
    success.

`allocate_compilationunit_debug`
    The callback function that the finalizer will use to allocate the memory that will be used for
    the hsa_compilationunit_debug_t that is returned. It is the responsibility of the call of the
    finalizer to destroy this memory, even if the finalizer does not report success.

`optimization_level`
    an implementation defined value that control the level of optimization performed by the
    finalizer.

`options`
    implementation defined options that can be specified to the finalizer.
    compilationunit_code: if the return status is success then a pointer to the generated
    hsa_compilationunit_code_t for the HSA component must be written.

The **hsa_finalize_brig** API returns HSA_STATUS_SUCCESS when the finalization is successful.
Otherwise, it returns one of the following errors:

  ▷ HSA_STATUS_ERROR_INVALID_ARGUMENT If *brig* is NULL or invalid, or if *kernel_directive* is
    invalid.

  ▷ HSA_STATUS_INFO_UNRECOGNIZED_OPTIONS If the options are not recognized, no error is
    returned, just an info status is used to indicate invalid options.

  ▷ HSA_STATUS_ERROR_OUT_OF_RESOURCES If the finalize API cannot allocate memory for
    compilationunit_code/debug or the deserialize cannot allocate memory for *code_object*.

▷ HSA_STATUS_ERROR_DIRECTIVE_MISMATCH If the directive in the control directive structure
   and in the HSAIL kernel mismatch or if the same directive is used with a different value in one
   of the functions used by this kernel.

The `hsa_code_object_t` is defined as follows:

```
typedef struct hsa_control_directives_s {
  hsa_control_directive_present64_t enabled_control_directives;
  hsa_exception_kind16_t enable_break_exceptions;
  hsa_exception_kind16_t enable_detect_exceptions;
  uint32_t max_dynamic_group_size;
  uint32_t max_flat_grid_size;
  uint32_t max_flat_workgroup_size;
  uint32_t requested_workgroups_per_cu;
  hsa_dim3_t required_grid_size;
  hsa_dim3_t required_workgroup_size;
  uint8_t required_dim;
  uint8_t reserved[75];
} hsa_control_directives_t;
```

---

enabled_control_directives
   If the value is 0 then there are no control directives specified and the rest of the fields can
   be ignored. The bits are accessed using the hsa_control_directives_present_mask_t. Any
   control directive that is not enabled in this bit set must have the value of all 0s.

enable_break_exceptions
   If enable break exceptions is not enabled in *enabled_control_directives*, then must be 0,
   otherwise must be non-0 and specifies the set of HSAIL exceptions that must have the
   BREAK policy enabled. If the HSAIL kernel being finalized has any
   enablebreakexceptions control directives, then the values specified by this argument are
   unioned with the values in these control directives. If any of the functions the kernel calls
   have an enablebreakexceptions control directive, then they must be equal or a subset of,
   this union.

enable_detect_exceptions
   If enable detect exceptions is not enabled in *enabled_control_directives*, then must be 0,
   otherwise must be non-0 and specifies the set of HSAIL exceptions that must have the
   DETECT policy enabled. If the kernel being finalized has any enabledetectexceptions
   control directives, then the values specified by this argument are unioned with the values
   in these control directives. If any of the functions the kernel calls have an
   enabledetectexceptions control directive, then they must be equal or a subset of, this
   union.

max_dynamic_group_size
   If max dynamic group size is not enabled in *enabled_control_directives* then this must be
   0, and any amount of dynamic group segment can be allocated for a dispatch, otherwise
   the value specifies the maximum number of bytes of dynamic group segment that can be

---

allocated for a dispatch. If the kernel being finalized has any maxdynamicsize control directives, then the values must be the same, and must be the same as this argument if it is enabled. This value can be used by the finalizer to determine the maximum number of bytes of group memory used by each work-group by adding this value to the group memory required for all group segment variables used by the kernel and all functions it calls, and group memory used to implement other HSAIL features such as fbarriers and the detect exception operations. This can allow the finalizer to determine the expected number of work-groups that can be executed by a compute unit and allow more resources to be allocated to the work-items if it is known that fewer work-groups can be executed due to group memory limitations.

`max_flat_grid_size`
> If this is is not enabled in *enabled_control_directives* then must be 0, otherwise must be greater than 0. See HSA Programmer's Reference Manual description of maxflatgridsize control directive.

`max_flat_workgroup_size`
> If this is is not enabled in *enabled_control_directives* then must be 0, otherwise must be greater than 0. See HSA Programmer's Reference Manual description of maxflatgridsize control directive.

`requested_workgroups_per_cu`
> If this is is not enabled in *enabled_control_directives* then must be 0 and the finalizer may generate ISA that could result in any number of work-groups executing on a single compute unit. Otherwise, the finalizer will *attempt* to generate ISA that will allow the specified number of work-groups to execute on a single compute unit. This is only a hint and can be ignored by the finalizer. If the kernel being finalized, or any of the functions it calls, has the same control directive, then the values must be the same or the finalization can fail. This can be used to determine the number of resources that should be allocated to a single work-group and work-item.

`required_grid_size`
> If not enabled then all elements for Dim3 must be 0, otherwise every element must be greater than 0. See HSA Programmer's Reference Manual description of requiredgridsize control directive.

`required_workgroup_size`
> If not enabled then all elements for Dim3 must be 0, and the produced code can be dispatched with any legal work-group range consistent with the dispatch dimensions. Otherwise, the code produced must always be dispatched with the specified work-group range. No element of the specified range must be 0. It must be consistent with required_dimensions and max_flat_workgroup_size. If the kernel being finalized, or any of the functions it calls, has a requiredworkgroupsize control directive, then the values must be the same. Specifying a value can allow the finalizer to optimize work-group id operations, and if the number of work-items in the work-group is less tha the WAVESIZE then barrier operations can be optimized to just a memory fence.

`required_dim`
> If required_dim is not enabled then must be 0 and the produced kernel code can be

dispatched with 1, 2 or 3 dimensions. If enabled then the value is 1..3 and the code
produced must only be dispatched with a dimension that matches. Other values are
illegal. If the kernel being finalized, or any of the functions it calls, has a requireddimsize
control directive, then the values must be the same. This can be used to optimize the
code generated to compute the absolute and flat work-group and work-item id, and the
dim HSAIL operations.

reserved[75]
    Reserved. Must be 0.

```
typedef struct hsa_kernel_code_s{
  hsa_code_t code;
  uint32_t workgroup_group_segment_byte_size;
  uint64_t kernarg_segment_byte_size;
  uint32_t workitem_private_segment_byte_size;
  uint32_t workgroup_fbarrier_count;
  hsa_control_directives_t control_directive;
  hsa_powertwo8_t wavefront_size;
  hsa_powertwo8_t kernarg_segment_alignment;
  uint8_t optimization_level;
  uint8_t reserved1[3];
} hsa_kernel_code_t;
```

code
    Common header that all code objects start with. code.type must be
    HSA_CODE_KERNEL.

workgroup_group_segment_byte_size
    The amount of group segment memory required by a work-group in bytes. This does not
    include any dynamically allocated group segment memory that may be added when the
    kernel is dispatched.

kernarg_segment_byte_size
    The size in bytes of the kernarg segment that holds the values of the arguments to the
    kernel.

workitem_private_segment_byte_size
    The amount of memory required for the combined private, spill and arg segments for a
    work-item in bytes.

workgroup_fbarrier_count
    Number of fbarrier's used in the kernel and all functions it calls. If the implemenation uses
    group memory to allocate the fbarriers then that amount must already be included in the
    workgroup_group_segment_byte_size total.

control_directive
> The values are the actually values used by the finalizer in generating the code. This may be the union of values specified as finalizer arguments and explicit HSAIL control directives. If a finalizer implementation ignores a control directive, and not generate constrained code, then the control directive will not be marked as enabled even though it was present in the HSAIL or finalizer argument. The values are intended to reflect the constraints that the code actually requires to correctly execute, not the values that were actually specified at finalize time.

wavefront_size
> Wavefront size expressed as a power of two. Must be a power of 2 in range 1..64 inclusive. Used to support runtime query that obtains wavefront size, which may be used by application to allocated dynamic group memory and set the dispatch work-group size.

kernarg_segment_alignment
> The maximum byte alignment of variables used by the kernel in the specified memory segment. Expressed as a power of two. Must be at lease HSA_POWERTWO_16.

optimization_level
> The optimization level specified when the kernel was finalized.

reserved1[3]
> Reserved. Must be 0. Component specific fields can follow this field.

The user must ensure that the BRIG is valid, failing which the finalize API will return an error. The desired kernel to finalize is specified as a location in the BRIG. The location is described as an offset into the BRIG .directive section.

HSAIL does not define a set of options that a finalizer needs to specify. To ensure portability, **hsa_finalize_brig** must not return an error status if a given compilation option is not recognized.

A.1.4 Freeing The Compilation Unit Code/Debug Object

The core runtime also provides a corresponding destruction API that destroys the compilation unit code and debug objects. This will reclaim all memory used by the hsa_compilationunit_code_t (including the ISA it contains) and associated hsa_compilationunit_debug_t. It will also unregister the ISA memory if appropriate.

```
hsa_status_t
    hsa_compilationunit_code_destroy(hsa_compilationunit_code_t
    *code_object);
```

code_object
> input. a pointer to the compilation unit object that needs to be destroyed.

```
hsa_status_t
    hsa_compilationunit_debug_destroy(hsa_compilationunit_debug_t
    *debug_object);
```

> debug_object
>   input. a pointer to the compilation unit debug object that needs to be destroyed.

The **hsa_compilationunit_code_destroy** API destroys the compilation unit code object where as **hsa_compilationunit_debug_destroy** destroys the compilation unit debug object.

Both API return HSA_STATUS_SUCCESS if the destruction was successful. Otherwise, they can return one of the following errors:

  ▷ HSA_STATUS_ERROR_RESOURCE_FREE if some of the resources consumed during initialization by the runtime could not be freed.

  ▷ HSA_STATUS_ERROR_INVALID_ARGUMENT if *code_object* or *debug_object* is NULL or does not point to a valid code object.

Note that destroying does not impact any memory segments that may have been allocated/reserved for use in a kernel from this object. It merely releases resources used to build the object.

A.1.5 Serializing and Deserializing a Compilation Unit

Because of the opaque nature of what the compilation unit actually represents, and in order for the users of the core runtime to build file containers, serialization and deserialization API are defined by the core runtime. The definition is as follows:

```
hsa_status_t hsa_compilationunit_serialize(hsa_compilationunit_code_t
    *code_object,
    hsa_alloc_t allocate_compilationunit_code,
    void *serialized_object);
```

> code_object
>   input. The code object to serialize.
>
> size
>   output. Size of the serialized object that is generated.
>
> serialized_object
>   Output. Pointer to the serialized object.

```
hsa_status_t hsa_compilationunit_deserialize( void *serialized_object,
    hsa_compilationunit_code_t **code_object);
```

> `serialized_object`
>   input. Pointer to the serialized object.
>
> `size`
>   input. Size of the serialized object that is generated.
>
> `code_object`
>   runtime allocated, output. The code object generated as a part of serialization.

HSA_STATUS_SUCCESS is returned if serialize/deserialize API finish successfully. Success of serialize API means the *code_object* has been successfully serialized and then copied into the location in memory (*serialized_object*). A successful deserialization recreates the code object, allocates memory for it and returns it. One of the following errors may be returned:

  ▷ HSA_STATUS_ERROR_INVALID_ARGUMENT If *code_object* or is NULL or does not point to a valid code object in the serialize API. For the deserialize API, this means the *serialized_object* is either null or is not valid or the size is 0.

  ▷ HSA_STATUS_ERROR_OUT_OF_RESOURCES If the serialize API cannot allocate memory for *serialized_object* or the deserialize cannot allocate memory for *code_object*.

A.2 Group Memory Usage

Group memory can be allocated either statically when the finalizer is called or dynamically by the user at launch time.

Static allocation is supported as follows:

  • The **hsa_finalize_brig** routine writes the amount of group memory needed by the finalized ISA to the hsa_code_object_t.*workgroup_group_segment_size_byte* field. The group memory usage includes group memory which is statically allocated in the HSAIL kernel, as well as private group memory used by the finalizer. Different HSA implementations might allocate different amounts of group memory.

  • The user copies the requested group segment usage to the AQL dispatch packet's hsa_aql_dispatch_packet_t.*group_segment_size_bytes* field.

  • The packet processor reads the group memory usage field and reserves the required resources at dispatch time.

• Statically allocated group memory starts at a segment offset of 0.

Dynamically allocated group memory allows the user to specify the group memory size when the kernel is launched. This is useful to support dynamic group memory allocation features supported by languages such as OpenCL. Essentially, the user manually calculates the offset for each kernel argument (including the static allocation in the calculation) and passes these as arguments to the HSAIL kernel. Specifically:

• As above, the hsa_finalize_brig routine returns the requested static group allocation.

• HSAIL will use standard 32-bit arguments (that is, **kernarg_u32**) to specify group segment offsets. The user is responsible for computing the offset for each group memory argument location. The first argument must start just above the static allocation, so it always has the offset of `hsa_code_object_t`.*workgroup_group_segment_size_byte*

• After setting the offset for each group memory argument, the user must set the QL dispatch packet's `hsa_aql_dispatch_packet_t`.*group_segment_size_bytes* field to the total amount of group memory used (static and dynamic allocations).

See below for an example of setting up dynamic group memory arguments for a kernel.

```
// ... assume setup for component, queue

// User dynamically requests 3 group allocations of 256, 384, and
   500 bytes.
// These can be specified at launch-time.
int size1 = 256;
int size2 = 384;
int size3 = 500;

hsa_aql_dispatch_packet_t aql_packet;
uint64_t* kernel_arguments =
reinterpret_cast<uint64_t *>(malloc(argument_size));

memset(kernel_arguments, 0, argument_size);

// Copy parameters into the AQL packet, computing relative offsets:
kernel_arguments[0] =
    kernel_object_ptr->workgroup_group_segment_size_bytes;

kernel_arguments[1] =
    kernel_object_ptr->workgroup_group_segment_size_bytes+size1;

kernel_arguments[2] =
    kernel_object_ptr->workgroup_group_segment_size_bytes+size1+size2;
```

```
 // Set the total group memory size:
aql_packet.group_segment_size_bytes = kernel.group_memory_usage +
   size1 +
size2 + size3;
```

Here is the corresponding kernel and usage model:

```
kernel &myDynamicGroupMemKernel (
              kernarg_u32 %groupOffset0,
              kernarg_u32 %groupOffset1,
              kernarg_u32 %groupOffset2,
              kernarg_u32 %foo)
{
        ld_kernarg_u32 $s0, [%groupOffset0]
        workitemid         $s1
        add                $s2, $s1, $s0
        st_group        0, [$s2]
        barrier
        ...
```

B

Images and Samplers

B.1   Images

<div style="text-align: right;">

C

</div>

Component Initiated Dispatches

C.1   Component-Initiated Dispatch

Due to architected support for a queue and design of AQL, HSA supports component-initiated dispatch, which is the ability for a kernel to dispatch a new kernel by writing an AQL packet directly to a user queue. In simple use cases, the AQL packet can be created on the host and passed as a parameter to the kernel. This eliminates the need to do dynamic memory allocation on the component, but has the limitation that the problem fanout must be known at the time the first kernel is launched (so that the AQL packets can be preallocated). HSA also supports more advanced use cases where the AQL packet is dynamically allocated (including the memory space for kernel arguments and spill/arg/private space) on the component. This usage model obviously requires dynamic component-side memory allocation, for both host and component memory.

Some requirements to do component-initiated dispatch:

- Ability to dynamically choose a kernel to dispatch: Let us assume for example that there are three kernels (A, B and C). If the host launches A, then the user has the choice of launching B or C, or even A in case of recursion. So, the user should be able to get the ISA and segment size (HsaAqlKernel) from the corresponding BRIG dynamically. [caveat: The code sample here does not show how we can do this. It assumes that the HsaAqlKernel is being passed as an argument to the parent kernel (A in this case)]

- Ability to dynamically allocate memory from the shader: We need to allocate memory for AQLPacket, different kernel segments in the AQLPacket, kernel arguments, and so forth.

- Ability for a finalizer to identify a default HSA queue to write AQLPacket: The HSA queue information resides in the runtime layer of the stack. This needs to be exchanged with the compiler so it can be stored in the global space. This way, when the compiler sees the queue, it knows where to pick the HSA queue information to write the AQLPacket.

- Ability to notify the completion of all the component-initiated dispatches on the host:

  - The beginning of execution of the child kernel may or may not wait for the parent kernel's completion. This is determined by the user and could be algorithm dependent.
  - If the parent (initiated from host) kernel finishes successfully, it means all kernels it initiated also finished successfully.
  - To implement this, we need to track the list of kernels launched from the parent. Change the status of parent to complete, only if parent and all its child kernels have completed successfully.

Implementations that support component initiated dispatches will need to support these requirements. If the implementation supports the stated requirements, the following actions will allow a component to initiate a dispatch:

- The queue and `hsa_code_object_t` (describing the kernel to launch) can be passed as arguments to the parent (the one launched from the host) kernel. If the dispatch is to the same queue, it is accessible via an HSAIL instruction.

- If not, get the HsaAqlKernel from the BRIG for the kernel that is chosen to be dynamically dispatchd.

- When new work is to be created, the HSAIL code would:

  - Use the kernel dynamic memory allocator to allocate a new AQLPacket.
  - Use inline HSAIL to replicate the functionality of the HsaInitAQLPacket function. We could perhaps provide an HSAIL library to implement this functionality. Recall this function:
    * Copies some fields from the HsaAqlKernel structure (for example, the kernel ISA) to the AQLPacket
    * Uses a host allocator to allocate memory for the kernel arguments
    * Uses a component allocator to allocate memory for spill, private, and arg segments

- The HSAIL knows the signature of the called function and can fill in the AQL packet with regular HSAIL global store instructions.

- The HSA queue is architected, so the HSAIL can use memory store instructions to dispatch the kernel for dispatch. Depending how the user queues are configured, atomic accesses might be necessary to handle contention with other writers. Note that, if the queue information is not passed in as an argument, the default queue can be chosen by the finalizer as it was exchanged earlier from the runtime layer.

- We also need to handle deallocation of the kernel arguments and spill/private/arg space after the kernel completes.

- On the host, check if the parent has finished. If the parent has finished successfully, then it means that all the child kernels have finished successfully too. If the parent or any of the child kernels failed, an error code will be returned.

D

Error Status Structure and Defined Values

D.1   Error States in Core API and Extentions

Bibliography

[1] HSA Foundation. The Heterogeneous Systems Architecture Programmers Reference Manual. Technical report, HSA Foundation, 2013.