



HSA Runtime Programmer's Reference Manual
0.184

June 18, 2014

Draft

©2013-2014 HSA Foundation. All rights reserved.

The contents of this document are provided in connection with the HSA Foundation specifications. This specification is protected by copyright laws and contains material proprietary to the HSA Foundation. It or any components may not be reproduced, republished, distributed, transmitted, displayed, broadcast or otherwise exploited in any manner without the express prior written permission of HSA Foundation. You may use this specification for implementing the functionality therein, without altering or removing any trademark, copyright or other notice from the specification, but the receipt or possession of this specification does not convey any rights to reproduce, disclose, or distribute its contents, or to manufacture, use, or sell anything that it may describe, in whole or in part.

HSA Foundation grants express permission to any current Founder, Promoter, Supporter Contributor, Academic or Associate member of HSA Foundation to copy and redistribute UNMODIFIED versions of this specification in any fashion, provided that NO CHARGE is made for the specification and the latest available update of the specification for any version of the API is used whenever possible. Such distributed specification may be re-formatted AS LONG AS the contents of the specification are not changed in any way. The specification may be incorporated into a product that is sold as long as such product includes significant independent work developed by the seller. A link to the current version of this specification on the HSA Foundation web-site should be included whenever possible with specification distributions.

HSA Foundation makes no, and expressly disclaims any, representations or warranties, express or implied, regarding this specification, including, without limitation, any implied warranties of merchantability or fitness for a particular purpose or non-infringement of any intellectual property. HSA Foundation makes no, and expressly disclaims any, warranties, express or implied, regarding the correctness, accuracy, completeness, timeliness, and reliability of the specification. Under no circumstances will the HSA Foundation, or any of its Founders, Promoters, Supporters, Academic, Contributors, and Associates members or their respective partners, officers, directors, employees, agents or representatives be liable for any damages, whether direct, indirect, special or consequential damages for lost revenues, lost profits, or otherwise, arising from or in connection with these materials.

Contents

1	Introduction	1
1.1	Overview	1
1.2	Execution Model	3
1.2.1	Initial Setup	3
1.2.2	Kernel Execution	3
1.2.3	Kernel Completion	4
1.2.4	Example	4
2	HSA Core Programming Guide	7
2.1	Initialization and Shut Down	7
2.1.1	API	7
2.1.1.1	hsa_init	7
2.1.1.2	hsa_shut_down	8
2.2	Runtime Notifications	9
2.2.1	API	9
2.2.1.1	hsa_status_t	9
2.2.1.2	hsa_status_string	10
2.3	Agent Information	12
2.3.1	API	12
2.3.1.1	hsa_agent_t	12
2.3.1.2	hsa_node_t	12
2.3.1.3	hsa_component_feature_t	12
2.3.1.4	hsa_device_type_t	12
2.3.1.5	hsa_agent_info_t	13
2.3.1.6	hsa_agent_get_info	14
2.3.1.7	hsa_iterate_agents	14
2.4	Signals	16
2.4.1	API	16
2.4.1.1	hsa_signal_value_t	16

2.4.1.2	hsa_signal_t	17
2.4.1.3	hsa_signal_create	17
2.4.1.4	hsa_signal_destroy	17
2.4.1.5	hsa_signal_load	18
2.4.1.6	hsa_signal_store	18
2.4.1.7	hsa_signal_exchange	19
2.4.1.8	hsa_signal_cas	19
2.4.1.9	hsa_signal_add	20
2.4.1.10	hsa_signal_subtract	21
2.4.1.11	hsa_signal_and	22
2.4.1.12	hsa_signal_or	22
2.4.1.13	hsa_signal_xor	23
2.4.1.14	hsa_signal_condition_t	24
2.4.1.15	hsa_wait_latency_t	24
2.4.1.16	hsa_signal_wait	24
2.5	Queues	26
2.5.1	Queue States	27
2.5.2	API	27
2.5.2.1	hsa_queue_type_t	27
2.5.2.2	hsa_queue_feature_t	27
2.5.2.3	hsa_queue_t	28
2.5.2.4	hsa_queue_create	29
2.5.2.5	hsa_queue_destroy	29
2.5.2.6	hsa_queue_inactivate	30
2.5.2.7	hsa_queue_load_read_index	31
2.5.2.8	hsa_queue_load_write_index	31
2.5.2.9	hsa_queue_store_write_index	31
2.5.2.10	hsa_queue_cas_write_index	32
2.5.2.11	hsa_queue_add_write_index	32
2.5.2.12	hsa_queue_store_read_index	33
2.6	Architected Queuing Language Packets	34
2.6.1	Dispatch packet	34
2.6.1.1	Memory information	35
2.6.2	Agent Dispatch packet	35
2.6.3	Barrier packet	35
2.6.4	API	35
2.6.4.1	hsa_aql_packet_format_t	35

2.6.4.2	hsa_fence_scope_t	36
2.6.4.3	hsa_aql_packet_header_t	36
2.6.4.4	hsa_aql_dispatch_packet_t	37
2.6.4.5	hsa_aql_agent_dispatch_packet_t	38
2.6.4.6	hsa_aql_barrier_packet_t	39
2.7	Memory	40
2.7.1	API	40
2.7.1.1	hsa_region_t	40
2.7.1.2	hsa_segment_t	40
2.7.1.3	hsa_region_info_t	41
2.7.1.4	hsa_region_get_info	41
2.7.1.5	hsa_agent_iterate_regions	42
2.7.1.6	hsa_memory_allocate	42
2.7.1.7	hsa_memory_free	43
2.7.1.8	hsa_memory_copy	43
2.7.1.9	hsa_memory_register	44
2.7.1.10	hsa_memory_deregister	45
2.8	Extensions to the Core Runtime API	46
2.8.1	API	46
2.8.1.1	hsa_extension_t	46
2.8.1.2	hsa_vendor_extension_query	46
2.8.1.3	hsa_extension_query	47
2.8.2	Example	48
2.9	Common Definitions	49
2.9.1	API	49
2.9.1.1	hsa_powertwo8_t	49
2.9.1.2	hsa_powertwo_t	49
2.9.1.3	hsa_dim3_t	49
2.9.1.4	hsa_dim_t	50
2.9.1.5	hsa_runtime_caller_t	50
2.9.1.6	hsa_runtime_alloc_data_callback_t	50
3	HSA Extensions Programming Guide	51
3.1	HSAIL Finalization	51
3.1.1	API	51
3.1.1.1	hsa_ext_brig_profile8_t	51
3.1.1.2	hsa_ext_brig_profile_t	51

3.1.1.3	hsa_ext_brig_machine_model8_t	51
3.1.1.4	hsa_ext_brig_machine_model_t	52
3.1.1.5	hsa_ext_brig_section_id32_t	52
3.1.1.6	hsa_ext_brig_section_id_t	52
3.1.1.7	hsa_ext_brig_section_header_t	52
3.1.1.8	hsa_ext_brig_module_t	53
3.1.1.9	hsa_ext_brig_module_handle_t	53
3.1.1.10	hsa_ext_brig_code_section_offset32_t	54
3.1.1.11	hsa_ext_exception_kind16_t	54
3.1.1.12	hsa_ext_exception_kind_t	54
3.1.1.13	hsa_ext_control_directive_present64_t	55
3.1.1.14	hsa_ext_control_directive_present_t	55
3.1.1.15	hsa_ext_control_directives_t	57
3.1.1.16	hsa_ext_code_kind32_t	59
3.1.1.17	hsa_ext_code_kind_t	59
3.1.1.18	hsa_ext_program_call_convention_id32_t	60
3.1.1.19	hsa_ext_program_call_convention_id_t	60
3.1.1.20	hsa_ext_code_handle_t	60
3.1.1.21	hsa_ext_debug_information_handle_t	60
3.1.1.22	hsa_ext_code_descriptor_t	60
3.1.1.23	hsa_ext_finalization_request_t	63
3.1.1.24	hsa_ext_finalization_t	63
3.1.1.25	hsa_ext_symbol_definition_callback_t	64
3.1.1.26	hsa_ext_symbol_address_callback_t	64
3.1.1.27	hsa_ext_error_message_callback_t	64
3.1.1.28	hsa_ext_finalize	64
3.1.1.29	hsa_ext_destroy_finalization	66
3.1.1.30	hsa_ext_serialize_finalization	66
3.1.1.31	hsa_ext_deserialize_finalization	67
3.2	HSAIL Linking	69
3.2.1	API	69
3.2.1.1	hsa_ext_program_handle_t	69
3.2.1.2	hsa_ext_program_create	69
3.2.1.3	hsa_ext_program_destroy	70
3.2.1.4	hsa_ext_add_module	71
3.2.1.5	hsa_ext_finalize_program	71
3.2.1.6	hsa_ext_query_program_agent_id	73

3.2.1.7	hsa_ext_query_program_agent_count	73
3.2.1.8	hsa_ext_query_program_agents	74
3.2.1.9	hsa_ext_query_program_module_count	74
3.2.1.10	hsa_ext_query_program_modules	75
3.2.1.11	hsa_ext_query_program_brig_module	75
3.2.1.12	hsa_ext_query_call_convention	76
3.2.1.13	hsa_ext_query_symbol_definition	76
3.2.1.14	hsa_ext_define_program_allocation_global_variable_address	77
3.2.1.15	hsa_ext_query_program_allocation_global_variable_address	78
3.2.1.16	hsa_ext_define_agent_allocation_global_variable_address	78
3.2.1.17	hsa_ext_query_agent_global_variable_address	79
3.2.1.18	hsa_ext_define_readonly_variable_address	80
3.2.1.19	hsa_ext_query_readonly_variable_address	80
3.2.1.20	hsa_ext_query_kernel_descriptor_address	81
3.2.1.21	hsa_ext_query_indirect_function_descriptor_address	81
3.2.1.22	hsa_ext_validate_program	82
3.2.1.23	hsa_ext_validate_program_module	83
3.2.1.24	hsa_ext_serialize_program	83
3.2.1.25	hsa_ext_program_allocation_symbol_address_t	84
3.2.1.26	hsa_ext_agent_allocation_symbol_address_t	84
3.2.1.27	hsa_ext_deserialize_program	84
3.3	Images and Samplers	86
3.3.1	API	86
3.3.1.1	hsa_ext_image_handle_t	86
3.3.1.2	hsa_ext_image_format_capability_t	87
3.3.1.3	hsa_ext_image_info_t	87
3.3.1.4	hsa_ext_image_access_permission_t	88
3.3.1.5	hsa_ext_image_geometry_t	88
3.3.1.6	hsa_ext_image_channel_type_t	88
3.3.1.7	hsa_ext_image_channel_order_t	89
3.3.1.8	hsa_ext_image_format_t	90
3.3.1.9	hsa_ext_image_descriptor_t	90
3.3.1.10	hsa_ext_image_range_t	91
3.3.1.11	hsa_ext_image_region_t	91
3.3.1.12	hsa_ext_sampler_handle_t	91
3.3.1.13	hsa_ext_sampler_addressing_mode_t	92
3.3.1.14	hsa_ext_sampler_coordinate_mode_t	92

3.3.1.15 hsa_ext_sampler_filter_mode_t	92
3.3.1.16 hsa_ext_sampler_descriptor_t	93
3.3.1.17 hsa_ext_image_get_format_capability	93
3.3.1.18 hsa_ext_image_get_info	94
3.3.1.19 hsa_ext_image_create_handle	95
3.3.1.20 hsa_ext_image_import	96
3.3.1.21 hsa_ext_image_export	97
3.3.1.22 hsa_ext_image_copy	98
3.3.1.23 hsa_ext_image_clear	99
3.3.1.24 hsa_ext_image_destroy_handle	100
3.3.1.25 hsa_ext_sampler_create_handle	100
3.3.1.26 hsa_ext_sampler_destroy_handle	101
3.4 Component Initiated Dispatches	103
Index - Core APIs	105
Index - Extension APIs	107
Bibliography	109

Chapter 1

Introduction

1.1 Overview

Recent heterogeneous system designs have integrated CPU, GPU, and other accelerator devices into a single platform with a shared high-bandwidth memory system. Specialized accelerators now complement general purpose CPU chips and are used to provide both power and performance benefits. These heterogeneous designs are now widely used in many computing markets including cellphones, tablets, personal computers, and game consoles. The Heterogeneous System Architecture (HSA) builds on the close physical integration of accelerators that is already occurring in the marketplace, and takes the next step by defining standards for uniting the accelerators architecturally. The HSA specifications include requirements for virtual memory, memory coherency, architected dispatch mechanisms, and power-efficient signals. HSA refers to these accelerators as "components".

The system architecture defines a consistent base for building portable applications that access the power and performance benefits of the dedicated HSA components. Many of these components, including GPUs and DSPs, are capable and flexible processors that have been extended with special hardware for accelerating parallel code. Historically these devices have been difficult to program due to a need for specialized or proprietary programming languages. HSA aims to bring the benefits of these components to mainstream programming languages using similar or identical syntax to that which is provided for programming multi-core CPUs.

In addition to the system architecture, HSA defines a portable, low-level, compiler intermediate language called "HSAIL". A high-level compiler generates the HSAIL for the parallel regions of code. A low-level compiler called the "finalizer" translates the intermediate HSAIL to target machine code. Each HSA component provides its own implementation of the finalizer. For more information on HSAIL, refer to the HSA Programmer's Reference Manual [1].

The final piece of the puzzle is the HSA Runtime API. The runtime is a thin, user-mode API that provides the interfaces necessary for the host to launch compute kernels to the available components. This document describes the architecture and APIs for the HSA Runtime. Key sections of the runtime API include:

- Error Handling
- Runtime initialization and shutdown
- Agent information
- Signals and synchronization
- Architected dispatch
- Memory management

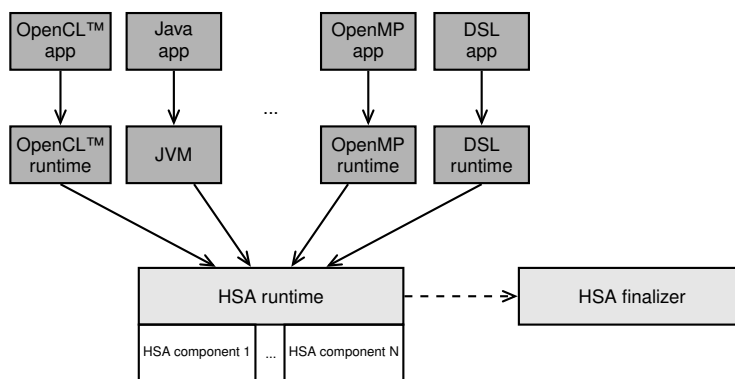


Figure 1.1: HSA Software Architecture

The remainder of this document describes the HSA software architecture and execution model, and includes functional descriptions for all of the HSA APIs and associated data structures.

Figure 1.1 shows how the HSA Runtime fits into a typical software architecture stack. At the top of the stack is a programming model such as OpenCL™, Java, OpenMP, or a domain-specific language (DSL). The programming model must include some way to indicate a parallel region that can be accelerated. For example, OpenCL has calls to `clEnqueueNDRangeKernel` with associated kernels and grid ranges. Java defines stream and lambda APIs, which provide support for both multi-core CPUs and HSA Components. OpenMP contains OMP pragmas that mark loops for parallel computing and that control other aspects of the parallel implementation. Other programming models can also build on this same infrastructure.

The language compiler is responsible for generating HSAIL code for the parallel regions of code. The code can be pre-compiled before runtime or compiled at runtime. A high-level compiler can generate the HSAIL before runtime, in which case when the application loads the finalizer converts the HSAIL to machine code for the target machine. Another option is to run the finalizer when the applications is built, in which case the resulting binary includes the machine code for the target architecture. The HSA Finalizer is an optional component of the HSA Runtime, which can reduce the footprint of the HSA software on systems where the finalization is done before runtime.

Each language also includes a "language runtime" that connects the language implementation to the HSA Runtime. When the language compiler generates code for a parallel region, it will include calls to the HSA Runtime to set up and dispatch the parallel region to the HSA Component. The language runtime is also responsible for initializing HSA, selecting target devices, creating execution queues, managing memory - and may use other HSA Runtime features as well. A runtime implementation may provide optional extensions. Applications can query the runtime to determine which extensions are available. This document describes the extensions for Finalization, Linking and Images.

The API for the HSA Runtime is standard across all HSA vendors, such that languages that use the HSA Runtime can execute on the different vendor's platforms that support the API. Each vendor is responsible for supplying their own HSA Runtime implementation that supports all of the HSA components in the vendor's platform. HSA does not provide a mechanism to combine runtimes from different vendors. The implementation of the HSA Runtime may include kernel-level components (required for some hardware components) or may be entirely user-space (for example, simulators or CPU implementations).

Figure 1.1 shows the "AQL" (Architected Queuing Language) path that application runtimes use to send commands directly to HSA components. For more information on AQL, refer to Section 2.6.

1.2 Execution Model

The HSA runtime exposes several details of the HSA hardware, including architected dispatches. The overall goal of the core runtime design is to provide a high-performance dispatch mechanism that is portable across multiple HSA vendor architectures. Two vendors with the same host ISA but different HSA-compliant GPUs will be able to run the same unmodified application binary file(s), because they support the HSA-architected AQL interface and supply a library that implements the architected core runtime API.

The HSA core runtime takes advantage of architected dispatch. Architected dispatch is the key feature in an HSA system that enables a user-level application to directly issue commands to the HSA Component hardware. Architected dispatch differentiates the HSA runtime from other higher-level runtime systems and programming models: other runtime systems provide software APIs for setting arguments and launching kernels, while HSA architects these at the hardware and specification level. The main advantage of architected dispatch is that the dispatch mechanism is architected at the HSA hardware level which means that an application can use regular memory operations and the wrapper API provided by the runtime to launch a kernel. The application creates user mode queues and AQL packets in memory, and then signals the HSA component to begin executing the packet using lightweight operations.

This section describes various features the HSA runtime provides to support architected dispatch, and the steps an application needs to perform in order to dispatch a kernel.

1.2.1 Initial Setup

One of the first steps an application must perform is agent (device) discovery. Agent discovery is performed after initialization of the runtime and information is made available to the user as opaque data structures. Section 2.3 describes how to find the available agents and query the runtime to obtain information about their attributes.

The next step in the setup is the creation of the component queues. Queues are an HSA architected mechanism that submits work to the HSA component hardware. For more information about the interfaces for queue creation, refer to Section 2.5. Different components may provide implementation-specific code under the HSA API for these functions.

1.2.2 Kernel Execution

The Systems Architecture Requirements document [2] specifies the structure of the *packets* (i.e. commands) that can be placed on the HSA user mode queues for execution by the component hardware. The Architected Queuing Language (AQL) describes the format of the packets. One type of AQL packet is a Dispatch packet. An application can create an AQL packet and initialize it with the code object that contains a component-specific ISA.

Optimized implementations can cache the result of this step and re-use the AQL packet for subsequent launches. Care must be taken to ensure that the AQL Dispatch packet (and the associated kernel and spill/arg/private memory) is not re-used before the launch completes. For simple cases such as a single-thread, synchronous launch, the AQL Dispatch packet(s) can be declared as a static variable and initialized at the same time the code is finalized. More complex cases might involve creating and tracking several Dispatch packets for a single kernel code object.

HSA hardware defines a packet process for processing these packets and a doorbell mechanism to inform the packet processing hardware that packets have been written into the queue. The HSA runtime defines a structure and update API to inform the hardware that the Dispatch packet has been written to the queue.

For more information on creating queues and the states of queues, refer to Section 2.5. For more information on packet formats and the states of packets, refer to Section 2.6.

After a packet is written and the hardware informed by way of the doorbell, execution can start. Execution of a kernel happens asynchronously. An application can write more packets to launch other kernels in the queue. This activity can overlap the actual execution of the kernel.

1.2.3 Kernel Completion

The architecture specification [2] defines signals as a mechanism for communication between different parts of an HSA system. The HSA core runtime defines signals as opaque objects. The API is defined to send a value to the signal and wait for a value at the signal. For more information on signals, refer to Section 2.4.

The AQL Dispatch packet provides a data field that allows an application to pass an opaque signal. When the HSA Component hardware observes a valid signal in an AQL packet, it sends a value to this signal when execution of the kernel completes (success or error). An application can wait on this signal to determine kernel completion. For more information on errors, refer to Section 2.2.

1.2.4 Example

```

/**
 * Dispatch a kernel in the command queue of a component.
 *
 * The source code has been simplified for readability. For instance, we do
 * not check the status code return by invocations of the HSA API, and we assume
 * that no asynchronous errors are generated by the runtime while executing the
 * kernel.
 */
#include "assert.h"
#include "stdio.h"
#include "string.h"

#include "hsa.h"

// Find component that can process Dispatch packets.
hsa_status_t get_dispatch_component(hsa_agent_t agent, void* data) {
    uint32_t features = 0;
    hsa_agent_get_info(agent, HSA_AGENT_INFO_COMPONENT_FEATURES, &features);
    if (features & HSA_COMPONENT_FEATURE_DISPATCH) {
        // Store component in user-provided buffer and return
        hsa_agent_t* ret = (hsa_agent_t*) data;
        *ret = agent;
        return HSA_STATUS_INFO_BREAK;
    }
    // Keep iterating
    return HSA_STATUS_SUCCESS;
}

int main() {
    // Initialize the runtime
    hsa_init();

    // Retrieve the component
    hsa_agent_t component;
    hsa_iterate_agents(get_dispatch_component, &component);

    // Create a queue in the selected component. The queue can hold up to four
    // packets, and has no callback or service queue associated with it.

```

```

hsa_queue_t *queue;
hsa_queue_create(component, 4, HSA_QUEUE_TYPE_SINGLE, NULL, NULL, &queue);

// Setup the packet encoding the task to execute
hsa_aql_dispatch_packet_t dispatch_packet;
const size_t packet_size = sizeof(dispatch_packet);
memset(&dispatch_packet, 0, packet_size); // reserved fields are zeroed
dispatch_packet.header.acquire_fence_scope = HSA_FENCE_SCOPE_COMPONENT;
dispatch_packet.header.release_fence_scope = HSA_FENCE_SCOPE_COMPONENT;
dispatch_packet.dimensions = 1;
dispatch_packet.workgroup_size_x = 256;
dispatch_packet.grid_size_x = 256;

// Indicate which ISA to run. The application is expected to have finalized a
// kernel (for example, using the finalization API). We will assume the object is
// located at address 0xDEADBEEF
dispatch_packet.kernel_object_address = 0xDEADBEEF;

// Assume our kernel receives no arguments, so no need to set the kernarg_address
// field in the packet.

// Create a signal with an initial value of one to monitor the task completion
hsa_signal_t signal;
hsa_signal_create(1, 0, NULL, &signal);
dispatch_packet.completion_signal = signal;

// Request a packet ID from the queue
// Since no packets have been enqueued yet, the expected ID is zero.
uint64_t write_index = hsa_queue_add_write_index_relaxed(queue, 1);

// Calculate the virtual address where to place the packet.
// No need to check if the queue is full or the index might wrap around...
void* dst = (void*)(queue->base_address + write_index * packet_size);
memcpy(dst, &dispatch_packet, packet_size);

// Notify the queue that the packet is ready to be processed
dispatch_packet.header.format = HSA_AQL_PACKET_FORMAT_DISPATCH;
hsa_signal_store_release(queue->doorbell_signal, write_index);

// Wait for the task to finish, which is the same as waiting for the value of the
// completion signal to be zero.
while (hsa_signal_wait_acquire(signal, HSA_EQ, 0, UINT64_MAX, HSA_WAIT_LATENCY_UNKNOWN) != 0);

// Done! The kernel has completed. Time to cleanup resources and leave.
hsa_signal_destroy(signal);
hsa_queue_destroy(queue);
hsa_shut_down();
return 0;
}

```


Chapter 2

HSA Core Programming Guide

This chapter describes the HSA Core runtime APIs, organized by functional area. For information on definitions that are not specific to any functionality, refer to Section 2.9.

Several operating systems allow functions to be executed when a DLL or a shared library is loaded (for example, DLL main in Windows and GCC *constructor/destructor* attributes that allow functions to be executed prior to main in several operating systems). Whether or not the HSA runtime functions are allowed to be invoked in such fashion may be implementation specific and is outside the scope of this specification.

Any header files distributed by the HSA foundation for this specification may contain calling-convention specific prefixes such as `__cdecl` or `__stdcall`, which are outside the scope of the API definition.

Unless otherwise stated, functions can be considered thread-safe.

2.1 Initialization and Shut Down

When an application initializes the runtime (**hsa_init**) for the first time in a given process, a runtime instance is created. The instance is internally reference counted such that multiple HSA clients within the same process do not interfere with each other. Invoking the initialization routine n times within a process does not create n runtime instances, but a unique runtime object with an associated reference counter of n . Shutting down the runtime (**hsa_shut_down**) is equivalent to decreasing its reference counter. When the reference counter is less than one, the runtime object ceases to exist and any reference to it (or to any resources created while it was active) results in undefined behavior.

2.1.1 API

2.1.1.1 hsa_init

```
hsa_status_t hsa_init();
```

Initialize the HSA runtime.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

If there is failure to allocate the resources required by the implementation.

HSA_STATUS_ERROR_REFCOUNT_OVERFLOW

If the runtime reference count reaches INT32_MAX.

Description

Initializes the HSA runtime if it is not already initialized, and increases the reference counter associated with the HSA runtime for the current process. Invocation of any HSA function other than **hsa_init** results in undefined behavior if the current HSA runtime reference counter is less than one.

2.1.1.2 hsa_shut_down

```
hsa_status_t hsa_shut_down();
```

Shut down the HSA runtime.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

Description

Decreases the reference count of the runtime instance. When the reference count reaches zero, the runtime is no longer considered valid but the application might call **hsa_init** to initialize the HSA runtime again.

Once the reference count of the runtime reaches zero, all the resources associated with it (queues, signals, agent information, etc.) are considered invalid and any attempt to reference them in subsequent API calls results in undefined behavior. When the reference count reaches zero, the HSA runtime might release resources associated with it.

2.2 Runtime Notifications

The runtime can report notifications (such as errors, and events) synchronously or asynchronously. The runtime uses the return value of functions in the HSA API to pass notifications synchronously. In this case, the notification is a status code of type `hsa_status_t` that indicates success or error.

The documentation of each function defines what constitutes a successful execution. When a HSA function does not execute successfully, the returned status code might help determining the source of the error. While a few error conditions can be generalized to a certain degree (e.g. failure in allocating resources) many errors can have implementation-specific explanations. For example, operations on signals (see Section 2.4) might fail if the implementation validates the signal object on which the signals operate. Because the representation of a signal is specific to the implementation, the reported error would simply indicate that the signal is invalid.

The `hsa_status_t` enumeration captures the result of any API function that has been executed, except for accessors and mutators. Success is represented by `HSA_STATUS_SUCCESS`, which has a value of zero. Error statuses are assigned positive integers and their identifiers start with the `HSA_STATUS_ERROR` prefix. The application might use `hsa_status_string` to obtain a string describing a status code.

The runtime passes *asynchronous* notifications in a different fashion. When the runtime detects an asynchronous event, it invokes a user-defined callback. For example, queues (see Section 2.5) are a common source of asynchronous events because the tasks queued by an application are asynchronously consumed by the packet processor. Callbacks are associated with queues when they are created. When the runtime detects an error in a queue, it invokes the callback associated with that queue and passes it an error flag (indicating what happened) and a pointer to the erroneous queue.

The application must use caution when using blocking functions within their callback implementation – a callback that does not return can render the runtime state to be undefined. The application cannot depend on thread local storage within the callbacks implementation and may safely kill the thread that registers the callback. The application is responsible for ensuring that the callback function is thread-safe. The runtime does not implement any default callbacks.

2.2.1 API

2.2.1.1 `hsa_status_t`

```
enum hsa_status_t
```

Values

`HSA_STATUS_SUCCESS = 0`

The function has been executed successfully.

`HSA_STATUS_INFO_BREAK`

A traversal over a list of elements has been interrupted by the application before completing.

`HSA_EXT_STATUS_INFO_ALREADY_INITIALIZED`

Indicates that initialization attempt failed due to prior initialization.

`HSA_EXT_STATUS_INFO_UNRECOGNIZED_OPTIONS`

TODO.

`HSA_STATUS_ERROR_INVALID_ARGUMENT`

One of the actual arguments does not meet a precondition stated in the documentation of the corresponding formal argument.

`HSA_STATUS_ERROR_INVALID_AGENT`

The agent is invalid.

HSA_STATUS_ERROR_INVALID_REGION

The memory region is invalid.

HSA_STATUS_ERROR_INVALID_SIGNAL

The signal is invalid.

HSA_STATUS_ERROR_INVALID_QUEUE

The queue is invalid.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

The runtime failed to allocate the necessary resources. This error may also occur when the core runtime library needs to spawn threads or create internal OS-specific events.

HSA_STATUS_ERROR_INVALID_PACKET_FORMAT

Indicates that the AQL packet is malformed.

HSA_STATUS_ERROR_RESOURCE_FREE

An error has been detected while releasing a resource.

HSA_STATUS_ERROR_NOT_INITIALIZED

An API other than **hsa_init** has been invoked while the reference count of the HSA runtime is zero.

HSA_STATUS_ERROR_REFCOUNT_OVERFLOW

The maximum reference count for the object has been reached.

HSA_EXT_STATUS_ERROR_DIRECTIVE_MISMATCH

TODO.

HSA_EXT_STATUS_ERROR_IMAGE_FORMAT_UNSUPPORTED

Image format is not supported.

HSA_EXT_STATUS_ERROR_IMAGE_SIZE_UNSUPPORTED

Image size is not supported.

2.2.1.2 **hsa_status_string**

```
hsa_status_t hsa_status_string(
    hsa_status_t status,
    const char ** status_string);
```

Query additional information about a status code.

Parameters

status

(in) Status code.

status_string

(out) A ISO/IEC 646 encoded English language string that potentially describes the error status. The string terminates in a NUL character.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *status_string* is NULL or *status* is not a valid status code.

2.3 Agent Information

The runtime exposes the list of agents that are available in the system. In the HSA API, agents are opaque handles of type `hsa_agent_t`. The application can traverse the list of available agents, and query attributes of a particular agent.

2.3.1 API

2.3.1.1 `hsa_agent_t`

```
typedef uint64_t hsa_agent_t
```

Opaque handle representing an agent, a device that participates in the HSA memory model. An agent can submit AQL packets for execution. Some agents are also components, i.e., they can be the target of AQL queries.

2.3.1.2 `hsa_node_t`

```
typedef uint64_t hsa_node_t
```

Opaque handle representing a NUMA node, a set of agents around memory which they can directly access.

2.3.1.3 `hsa_component_feature_t`

```
enum hsa_component_feature_t
```

Component features.

Values

`HSA_COMPONENT_FEATURE_DISPATCH = 1`

The component supports AQL packets of Dispatch type.

`HSA_COMPONENT_FEATURE_AGENT_DISPATCH = 2`

The component supports AQL packets of Agent Dispatch type.

2.3.1.4 `hsa_device_type_t`

```
enum hsa_device_type_t
```

Hardware device type.

Values

`HSA_DEVICE_TYPE_CPU = 0`

CPU device.

`HSA_DEVICE_TYPE_GPU = 1`

GPU device.

HSA_DEVICE_TYPE_DSP = 2
DSP device.

2.3.1.5 hsa_agent_info_t

```
enum hsa_agent_info_t
```

Agent attributes.

Values

HSA_AGENT_INFO_NAME

Agent name. The type of this attribute is char[64]. The string is NUL-terminated.

HSA_AGENT_INFO_NODE

NUMA node associated with the agent. The type of this attribute is hsa_node_t.

HSA_AGENT_INFO_COMPONENT_FEATURES

Component capabilities. Might be none if the agent is not a component. The type of this attribute is uint32_t, a bit-field of hsa_component_feature_t values.

HSA_AGENT_INFO_DEVICE

Device associated with the agent. The type of this attribute is hsa_device_type_t.

HSA_AGENT_INFO_VENDOR_NAME

Name of vendor. The type of this attribute is char[64]. The string is NUL-terminated.

HSA_AGENT_INFO_WAVEFRONT_SIZE

Number of work-items in a wavefront. The type of this attribute is uint32_t.

HSA_AGENT_INFO_CACHE_SIZE

Array of cache sizes (L1..L3). Each size is expressed in bytes. A size of 0 for a particular level indicates that there is no cache information for that level. The type of this attribute is uint32_t[4].

HSA_AGENT_INFO_MAX_GRID_DIM

Maximum number of work-items in a grid. The type of this attribute is hsa_dim3_t.

HSA_AGENT_INFO_MAX_WORKGROUP_DIM

Maximum number of work-items in a workgroup. The type of this attribute is hsa_dim3_t.

HSA_AGENT_INFO_QUEUE_MAX_PACKETS

Maximum capacity (in terms of packets) that a queue can hold. The type of this attribute is uint32_t.

HSA_AGENT_INFO_CLOCK

Current timestamp. The value of this attribute monotonically increases at a constant rate. The type of this attribute is uint64_t.

HSA_AGENT_INFO_CLOCK_FREQUENCY

Timestamp value increase rate, in MHz. The timestamp (clock) frequency is in the range 1-400MHz. The type of this attribute is uint16_t.

HSA_AGENT_INFO_MAX_SIGNAL_WAIT

Maximum duration of a signal wait operation. Expressed as a count based on the timestamp frequency. The type of this attribute is uint64_t.

2.3.1.6 hsa_agent_get_info

```
hsa_status_t hsa_agent_get_info(
    hsa_agent_t agent,
    hsa_agent_info_t attribute,
    void * value);
```

Get the current value of an attribute for a given agent.

Parameters

agent

(in) A valid agent.

attribute

(in) Attribute to query.

value

(out) Pointer to an application-allocated buffer where to store the value of the attribute. If the buffer passed by the application is not large enough to hold the value of *attribute*, the behavior is undefined.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_AGENT

If the agent is invalid.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *attribute* is not a valid agent attribute, or *value* is NULL.

2.3.1.7 hsa_iterate_agents

```
hsa_status_t hsa_iterate_agents(
    hsa_status_t(*) (hsa_agent_t agent, void *data) callback,
    void * data);
```

Iterate over the available agents, and invoke an application-defined callback on every iteration.

Parameters

callback

(in) Callback to be invoked once per agent.

data

(in) Application data that is passed to *callback* on every iteration. Might be NULL.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *callback* is NULL.

Description

If *callback* returns a status other than `HSA_STATUS_SUCCESS` for a particular iteration, the traversal stops and the function returns that status value.

2.4 Signals

HSA agents can communicate with each other by using coherent global memory, or by using signals. In an HSA system, a signal can be more power efficient than coherent global memory for communication.

A signal allows multiple producers and consumers. The runtime API uses opaque signal handlers of type `hsa_signal_t` to represent signals. A signal carries a value that can be updated or conditionally waited upon through an API call or HSAIL instruction. The value occupies four or eight bytes depending on the machine model being used.

Updating the value of a signal is equivalent to sending the signal. In addition to the regular update (store) of a signal value, an application can perform atomic operations such as add, subtract, or compare-and-swap. Each API operation also has specific memory order semantics associated with it. For example, store-release (**`hsa_signal_store_release`** function) is equivalent to storing a value on the signal handle with release memory ordering. The combinations of actions and memory orders available in the API match the corresponding HSAIL instructions [1].

The application may wait on a signal, with a condition specifying the terms of wait. The wait can be done either in the HSA Component by using an HSAIL **`wait`** instruction or by using a runtime API call. It is the responsibility of the application to detect that an error has occurred during the wait by checking the output status.

Waiting implies reading the current signal value (which is returned to the application) using an acquire (**`hsa_signal_wait_acquire`**) or a relaxed (**`hsa_signal_wait_relaxed`**) memory order. HSA implementations might establish a maximum waiting time.

As explained above, an application can modify or wait on signals by using the runtime API. However, signals also allow expressing complex tasks dependencies that are automatically handled by the packet processors. For example, if task *y* executing in one component consumes the result of task *x* executing in a different component, then *y depends* on *x*. In HSA, this dependency can be enforced across components by creating a signal that will be simultaneously used as a) the completion signal of a Dispatch packet *px* corresponding to *x* b) the dependency signal in a Barrier packet that precedes the Dispatch packet *py* corresponding to *y*. The packet processor enforces the task dependency by not launching *py* until *px* has completed.

For more information on HSA interfaces related to adding packets to a queue, refer to Section 2.5. For more information on AQL packets, refer to Section 2.6.

2.4.1 API

2.4.1.1 `hsa_signal_value_t`

```
#ifndef HSA_LARGE_MODEL
    typedef int64_t hsa_signal_value_t
#else
    typedef int32_t hsa_signal_value_t
#endif
```

Signal value. The value occupies 32 bits in small machine mode, and 64 bits in large machine mode.

2.4.1.2 hsa_signal_t

```
typedef uint64_t hsa_signal_t
```

Signal handle. The value 0 is reserved.

2.4.1.3 hsa_signal_create

```
hsa_status_t hsa_signal_create(
    hsa_signal_value_t initial_value,
    uint32_t num_consumers,
    hsa_agent_t * consumers,
    hsa_signal_t * signal);
```

Create a signal.

Parameters

initial_value

(in) Initial value of the signal.

num_consumers

(in) Size of *consumers*. If *consumers* is NULL, this argument is ignored.

consumers

(in) List of agents that might consume (wait on) the signal. Might be NULL.

signal

(out) Pointer to a memory location where the runtime will store the newly created signal handle.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

If there is failure to allocate the resources required by the implementation.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *signal* is NULL.

Description

The application might provide a list of agents that will wait on the signal. If an agent not listed in *consumers* waits on *signal*, the behavior is undefined.

2.4.1.4 hsa_signal_destroy

```
hsa_status_t hsa_signal_destroy(
    hsa_signal_t signal);
```

Destroy a signal previous created by **hsa_signal_create**.

Parameters

signal

(in) Signal.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_SIGNAL

If *signal* is invalid.

2.4.1.5 hsa_signal_load

```
hsa_signal_value_t hsa_signal_load_acquire(
    hsa_signal_t signal);
```

```
hsa_signal_value_t hsa_signal_load_relaxed(
    hsa_signal_t signal);
```

Atomically read the current signal value.

Parameters

signal

(in) Signal.

Returns

Value of the signal.

2.4.1.6 hsa_signal_store

```
void hsa_signal_store_relaxed(
    hsa_signal_t signal,
    hsa_signal_value_t value);
```

```
void hsa_signal_store_release(
    hsa_signal_t signal,
    hsa_signal_value_t value);
```

Atomically set the value of a signal.

Parameters

signal

(in) Signal.

value

(in) New signal value.

Description

Modifying the value of the signal implies sending the signal. All the agents waiting on *signal* for which *value*

satisfies their wait condition are awakened.

2.4.1.7 hsa_signal_exchange

```
hsa_signal_value_t hsa_signal_exchange_acq_rel(
    hsa_signal_t signal,
    hsa_signal_value_t value);

hsa_signal_value_t hsa_signal_exchange_acquire(
    hsa_signal_t signal,
    hsa_signal_value_t value);

hsa_signal_value_t hsa_signal_exchange_relaxed(
    hsa_signal_t signal,
    hsa_signal_value_t value);

hsa_signal_value_t hsa_signal_exchange_release(
    hsa_signal_t signal,
    hsa_signal_value_t value);
```

Atomically set the value of a signal and return its previous value.

Parameters

signal

(in) Signal. If *signal* is a queue doorbell signal, the behavior is undefined.

value

(in) New value.

Returns

Value of the signal prior to the exchange.

Description

Modifying the value of the signal implies sending the signal. All the agents waiting on *signal* for which *value* satisfies their wait condition are awakened.

2.4.1.8 hsa_signal_cas

```
hsa_signal_value_t hsa_signal_cas_acq_rel(
    hsa_signal_t signal,
    hsa_signal_value_t expected,
    hsa_signal_value_t value);
```

```

hsa_signal_value_t hsa_signal_cas_acquire(
    hsa_signal_t signal,
    hsa_signal_value_t expected,
    hsa_signal_value_t value);

hsa_signal_value_t hsa_signal_cas_relaxed(
    hsa_signal_t signal,
    hsa_signal_value_t expected,
    hsa_signal_value_t value);

hsa_signal_value_t hsa_signal_cas_release(
    hsa_signal_t signal,
    hsa_signal_value_t expected,
    hsa_signal_value_t value);

```

Atomically set the value of a signal if the observed value is equal to the expected value. The observed value is returned regardless of whether the replacement was done.

Parameters

signal

(in) Signal. If *signal* is a queue doorbell signal, the behavior is undefined.

expected

(in) Value to compare with.

value

(in) New value.

Returns

Observed value of the signal.

Description

If the value of the signal is modified, all the agents waiting on *signal* for which *value* satisfies their wait condition are awakened.

2.4.1.9 hsa_signal_add

```

void hsa_signal_add_acq_rel(
    hsa_signal_t signal,
    hsa_signal_value_t value);

```

```

void hsa_signal_add_acquire(
    hsa_signal_t signal,
    hsa_signal_value_t value);

void hsa_signal_add_relaxed(
    hsa_signal_t signal,
    hsa_signal_value_t value);

void hsa_signal_add_release(
    hsa_signal_t signal,
    hsa_signal_value_t value);

```

Atomically increment the value of a signal by a given amount.

Parameters

signal

(in) Signal. If *signal* is a queue doorbell signal, the behavior is undefined.

value

(in) Value to add to the value of the signal.

Description

Modifying the value of the signal implies sending the signal. All the agents waiting on *signal* for which the new value satisfies their wait condition are awakened.

2.4.1.10 hsa_signal_subtract

```

void hsa_signal_subtract_acq_rel(
    hsa_signal_t signal,
    hsa_signal_value_t value);

void hsa_signal_subtract_acquire(
    hsa_signal_t signal,
    hsa_signal_value_t value);

void hsa_signal_subtract_relaxed(
    hsa_signal_t signal,
    hsa_signal_value_t value);

void hsa_signal_subtract_release(
    hsa_signal_t signal,
    hsa_signal_value_t value);

```

Atomically decrement the value of a signal by a given amount.

Parameters

signal

(in) Signal. If *signal* is a queue doorbell signal, the behavior is undefined.

value

(in) Value to subtract from the value of the signal.

Description

Modifying the value of the signal implies sending the signal. All the agents waiting on *signal* for which the new value satisfies their wait condition are awakened.

2.4.1.11 hsa_signal_and

```
void hsa_signal_and_acq_rel(
    hsa_signal_t signal,
    hsa_signal_value_t value);

void hsa_signal_and_acquire(
    hsa_signal_t signal,
    hsa_signal_value_t value);

void hsa_signal_and_relaxed(
    hsa_signal_t signal,
    hsa_signal_value_t value);

void hsa_signal_and_release(
    hsa_signal_t signal,
    hsa_signal_value_t value);
```

Atomically perform a logical AND of the value of a signal and a given value.

Parameters

signal

(in) Signal. If *signal* is a queue doorbell signal, the behavior is undefined.

value

(in) Value to AND with the value of the signal.

Description

Modifying the value of the signal implies sending the signal. All the agents waiting on *signal* for which the new value satisfies their wait condition are awakened.

2.4.1.12 hsa_signal_or

```
void hsa_signal_or_acq_rel(
    hsa_signal_t signal,
    hsa_signal_value_t value);
```

```

void hsa_signal_or_acquire(
    hsa_signal_t signal,
    hsa_signal_value_t value);

void hsa_signal_or_relaxed(
    hsa_signal_t signal,
    hsa_signal_value_t value);

void hsa_signal_or_release(
    hsa_signal_t signal,
    hsa_signal_value_t value);

```

Atomically perform a logical OR of the value of a signal and a given value.

Parameters

signal

(in) Signal. If *signal* is a queue doorbell signal, the behavior is undefined.

value

(in) Value to OR with the value of the signal.

Description

Modifying the value of the signal implies sending the signal. All the agents waiting on *signal* for which the new value satisfies their wait condition are awakened.

2.4.1.13 hsa_signal_xor

```

void hsa_signal_xor_acq_rel(
    hsa_signal_t signal,
    hsa_signal_value_t value);

void hsa_signal_xor_acquire(
    hsa_signal_t signal,
    hsa_signal_value_t value);

void hsa_signal_xor_relaxed(
    hsa_signal_t signal,
    hsa_signal_value_t value);

void hsa_signal_xor_release(
    hsa_signal_t signal,
    hsa_signal_value_t value);

```

Atomically perform a logical XOR of the value of a signal and a given value.

Parameters

signal

(in) Signal. If *signal* is a queue doorbell signal, the behavior is undefined.

value

(in) Value to XOR with the value of the signal.

Description

Modifying the value of the signal implies sending the signal. All the agents waiting on *signal* for which the new value satisfies their wait condition are awakened.

2.4.1.14 hsa_signal_condition_t

```
enum hsa_signal_condition_t
```

Wait condition operator.

Values

HSA_EQ

The two operands are equal.

HSA_NE

The two operands are not equal.

HSA_LT

The first operand is less than the second operand.

HSA_GTE

The first operand is greater than or equal to the second operand.

2.4.1.15 hsa_wait_latency_t

```
enum hsa_wait_latency_t
```

Expected wait latency.

Values

HSA_WAIT_LATENCY_LOW

Low latency.

HSA_WAIT_LATENCY_HIGH

High latency.

HSA_WAIT_LATENCY_UNKNOWN

Unknown latency.

2.4.1.16 hsa_signal_wait

```
hsa_signal_value_t hsa_signal_wait_acquire(
    hsa_signal_t signal,
    hsa_signal_condition_t condition,
    hsa_signal_value_t compare_value,
    uint64_t timeout_hint,
    hsa_wait_latency_t latency_hint);
```



```

hsa_signal_value_t hsa_signal_wait_relaxed(
    hsa_signal_t signal,
    hsa_signal_condition_t condition,
    hsa_signal_value_t compare_value,
    uint64_t timeout_hint,
    hsa_wait_latency_t latency_hint);

```

Wait until the value of a signal satisfies a given condition.

Parameters

signal

(in) Signal.

condition

(in) Condition used to compare the signal value with *compare_value*.

compare_value

(in) Value to compare with.

timeout_hint

(in) Maximum duration hint. Specified in the same unit as the system timestamp. A value of UINT64_MAX indicates no maximum.

latency_hint

(in) Hint indicating whether the signal value is expected to meet the given condition in a short period of time or not. The HSA runtime may use this hint to optimize the wait implementation.

Returns

Observed value of the signal. The value might not satisfy the given condition.

Description

The wait may return before the condition is satisfied due to multiple reasons. It is the application's responsibility to check the returned signal value.

When the wait operation atomically loads the value of the passed signal, it uses the memory order indicated in the function name.

The application might indicate a preference about the maximum wait duration. The operation might block for a shorter or longer time even if the condition is not met.

2.5 Queues

HSA hardware supports packet execution through user-level queues. A queue is associated with a specific component, which might have several queues attached to it. Applications launch packets (explained in more detail in the next section) by placing the packets in the user mode queue of a component. The queue memory is then processed by the packet processor as though it were a ring buffer, with separate memory locations defining the current write and read addresses.

The HSA runtime provides an interface (**hsa_queue_create**) to create a user mode queue. When an application creates a queue, the runtime allocates memory for the data structure that represents the queue, as well as the buffer indicated by the *base_address* field.

In the HSA API, queues are of type `hsa_queue_t`. Queues are read-only data structures. Writing values directly to a queue struct results in undefined behavior.

In addition to the visible fields listed in `hsa_queue_t`, a queue also contains a read index and a write index. The write index is the number of packets allocated so far in that queue; the read index is the number of packets that have been processed and released by the queue's packet processor (i.e., the identifier of the next packet to be released). Both the write index and the read index are 64-bit unsigned integers that can exceed the maximum number of packets the queue can hold.

The write index and the read index are initialized to a value of 0. The runtime does not directly expose the index values to an application. An application can only access the values by using dedicated APIs. The available index functions differ on the index of interest (read or write), action to be performed (addition, compare and swap, etc.), and memory order (relaxed, release, etc.).

When the application observes that the read index matches the write index, the queue can be considered empty – this does not mean that the kernels have finished execution, just that all packets have been consumed. On the other hand, if the observed write index is greater or equal than the sum of the read index and the size of the queue, then the queue is full.

The *doorbell_signal* field contains a signal that the agent writing packets uses to indicate the packet processor that it has work to do. The value which the doorbell signal must be signaled with corresponds to the identifier of the packet that is ready to be launched. The new task might be consumed by the packet processor even before the doorbell signal has been signaled by the agent. This is because the packet processor might be already processing some other packet and observes that there is new work available, so it processes the new packets. In any case, the agent is required to signal the doorbell for every batch of packets it writes.

The runtime uses agent dispatch packets to specify runtime-defined or user-registered functions that will be executed on the agent (typically, the host CPU). Service queues consume agent dispatch packets. The

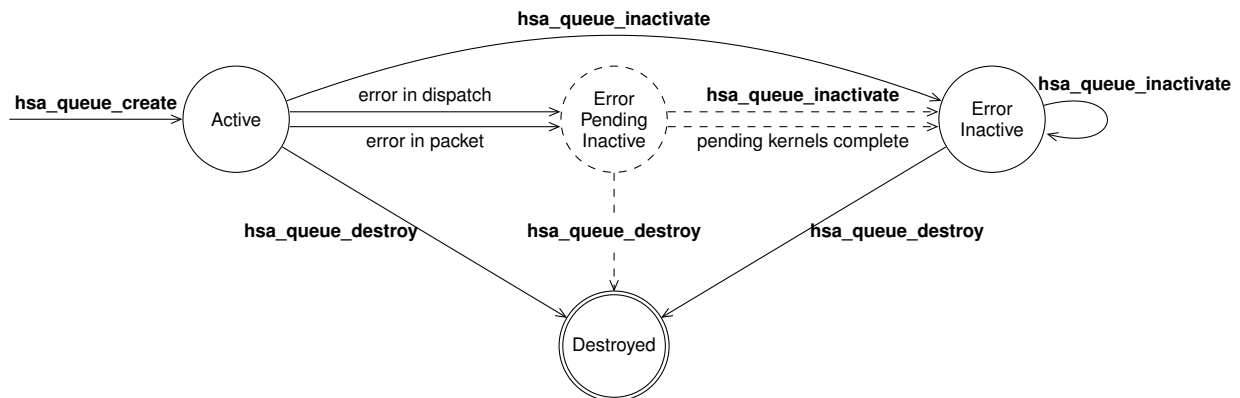


Figure 2.1: Queue state diagram.

service queue is visible to HSA agents through the queue structure *service_queue* field and is serviced by an appropriate HSA agent. The application may chose not to use a service queue.

2.5.1 Queue States

A queue in HSA, once created, can be in one of the following states: *active*, *error pending inactive*, *error inactive* or *destroyed*. A state diagram showing the various states and transitions is shown in Figure 2.1.

Active When a queue is successfully created using the **hsa_queue_create** API, it enters the active state. Packets can be added to the queue and are consumed by the packet processor. The actual initiation of dispatch may depend on the resources available for the dispatch. Writing packets to the queue, updating the write index or ringing the doorbell have an effect only when the queue is in the active state. The queue is not monitored by a packet processor in any other state.

Error pending inactive If an error is encountered during packet processing (for example, the packet format is invalid) or dispatch, the packet processor stops. At this point, there might be in-flight kernels and resources (such as segment allocation) that have been setup for a dispatch but have not yet been freed. So the queue is not entirely inactive, but when the asynchronous activity concludes, it will become inactive. A queue in *error pending inactive* state is not considered destroyed. It still needs to be destroyed so that the runtime can reclaim the memory allocated for this queue. If an application provides a callback at the time the application creates the queue, the callback is invoked after the queue is marked inactive.

Inactive If all the asynchronous activity concludes, the queue enters the inactive state. A queue can also enter this state when the user explicitly invokes the **hsa_queue_inactivate** API (note that the queue error callback can invoke this API). In an inactive state, the queue structure and its packets may be inspected. Only the packets that are between the read index and the write index in the queue structure are considered to be valid for inspection by the user. The packet processor guarantees that all the packets that have been consumed by the packet processor (see Section 2.6.1) will be notified. Inactivating a queue that is already in the inactive state has no effect.

Destroyed The queue has been destroyed by the user. The queue structure and the associated resources (ring buffer, indexes, etc.) are no longer valid.

2.5.2 API

2.5.2.1 hsa_queue_type_t

```
enum hsa_queue_type_t
```

Queue type. Intended to be used for dynamic queue protocol determination.

Values

HSA_QUEUE_TYPE_MULTI = 0

Multiple producers are supported.

HSA_QUEUE_TYPE_SINGLE = 1

Only a single producer is supported.

2.5.2.2 hsa_queue_feature_t

```
enum hsa_queue_feature_t
```

Queue features.

Values

HSA_QUEUE_FEATURE_DISPATCH = 1

Queue supports Dispatch packets.

HSA_QUEUE_FEATURE_AGENT_DISPATCH = 2

Queue supports Agent Dispatch packets.

2.5.2.3 hsa_queue_t

```
typedef struct hsa_queue_s {
    hsa_queue_type_t type;
    uint32_t features;
    uint64_t base_address;
    hsa_signal_t doorbell_signal;
    uint32_t size;
    uint32_t id;
    uint64_t service_queue;
} hsa_queue_t
```

User mode queue.

Data Fields

type

Queue type.

features

Queue features mask. This is a bit-field of `hsa_queue_feature_t` values. Applications should ignore any unknown set bits.

base_address

Starting address of the runtime-allocated buffer which is used to store the AQL packets. Aligned to the size of an AQL packet.

doorbell_signal

Signal object used by the application to indicate the ID of a packet that is ready to be processed.

The HSA runtime is responsible for the life cycle of the doorbell signal: replacing it with another signal or destroying it is not allowed and results in undefined behavior.

If *type* is `HSA_QUEUE_TYPE_SINGLE`, it is the application's responsibility to update the doorbell signal value with monotonically increasing indexes.

size

Maximum number of packets the queue can hold. Must be a power of two.

id

Queue identifier which is unique per process.

service_queue

A pointer to another user mode queue that can be used by the HSAIL kernel to request system services.

Description

Queues are read-only, but HSA agents can directly modify the contents of the buffer pointed by *base_*-

address, or use runtime APIs to access the doorbell signal or the service queue.

2.5.2.4 hsa_queue_create

```
hsa_status_t hsa_queue_create(
    hsa_agent_t component,
    size_t size,
    hsa_queue_type_t type,
    void(*) (hsa_status_t status, hsa_queue_t *queue) callback,
    hsa_queue_t * service_queue,
    hsa_queue_t ** queue);
```

Create a user mode queue.

Parameters

component

(in) Pointer to the component on which this queue is to be created.

size

(in) Number of packets the queue is expected to hold. Must be a power of two.

type

(in) Type of the queue.

callback

(in) Callback to be invoked for events related with this queue. Might be NULL.

service_queue

(in) Pointer to service queue to be associated with the newly created queue. Might be NULL, or another previously created queue that supports Agent Dispatch packets.

queue

(out) Memory location where the runtime stores a pointer to the newly created queue.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

If there is failure to allocate the resources required by the implementation.

HSA_STATUS_ERROR_INVALID_AGENT

If the component is invalid.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *component* is NULL, *size* is not a power of two, *type* is not a valid queue type, or *queue* is NULL.

Description

When a queue is created, the runtime also allocates the packet buffer and the completion signal. The application should only rely on the error code returned to determine if the queue is valid.

2.5.2.5 hsa_queue_destroy

```
hsa_status_t hsa_queue_destroy(
    hsa_queue_t * queue);
```

Destroy a user mode queue.

Parameters

queue

(in) Pointer to a queue.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_QUEUE

If the queue is invalid.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *queue* is NULL.

Description

A destroyed queue might not be accessed after being destroyed. When a queue is destroyed, the state of the AQL packets that have not been yet fully processed becomes undefined.

2.5.2.6 hsa_queue_inactivate

```
hsa_status_t hsa_queue_inactivate(
    hsa_queue_t * queue);
```

Inactivate a queue.

Parameters

queue

(in) Pointer to a queue.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_QUEUE

If the queue is invalid.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *queue* is NULL.

Description

Inactivating the queue aborts any pending executions and prevent any new packets from being processed. Any more packets written to the queue once it is inactivated will be ignored by the packet processor.

2.5.2.7 hsa_queue_load_read_index

```
uint64_t hsa_queue_load_read_index_acquire(
    hsa_queue_t * queue);

uint64_t hsa_queue_load_read_index_relaxed(
    hsa_queue_t * queue);
```

Atomically load the read index of a queue.

Parameters

queue

(in) Pointer to a queue.

Returns

Read index of the queue pointed by *queue*.

2.5.2.8 hsa_queue_load_write_index

```
uint64_t hsa_queue_load_write_index_acquire(
    hsa_queue_t * queue);

uint64_t hsa_queue_load_write_index_relaxed(
    hsa_queue_t * queue);
```

Atomically load the write index of a queue.

Parameters

queue

(in) Pointer to a queue.

Returns

Write index of the queue pointed by *queue*.

2.5.2.9 hsa_queue_store_write_index

```
void hsa_queue_store_write_index_relaxed(
    hsa_queue_t * queue,
    uint64_t value);

void hsa_queue_store_write_index_release(
    hsa_queue_t * queue,
    uint64_t value);
```

Atomically set the write index of a queue.

Parameters

queue

(in) Pointer to a queue.

value

(in) Value to assign to the write index.

2.5.2.10 hsa_queue_cas_write_index

```
uint64_t hsa_queue_cas_write_index_acq_rel(
    hsa_queue_t * queue,
    uint64_t expected,
    uint64_t value);

uint64_t hsa_queue_cas_write_index_acquire(
    hsa_queue_t * queue,
    uint64_t expected,
    uint64_t value);

uint64_t hsa_queue_cas_write_index_relaxed(
    hsa_queue_t * queue,
    uint64_t expected,
    uint64_t value);

uint64_t hsa_queue_cas_write_index_release(
    hsa_queue_t * queue,
    uint64_t expected,
    uint64_t value);
```

Atomically set the write index of a queue if the observed value is equal to the expected value. The application can inspect the returned value to determine if the replacement was done.

Parameters

queue

(in) Pointer to a queue.

expected

(in) Expected value.

value

(in) Value to assign to the write index if *expected* matches the observed write index.

Returns

Previous value of the write index.

2.5.2.11 hsa_queue_add_write_index

```
uint64_t hsa_queue_add_write_index_acq_rel(
    hsa_queue_t * queue,
    uint64_t value);
```



```

uint64_t hsa_queue_add_write_index_acquire(
    hsa_queue_t * queue,
    uint64_t value);

uint64_t hsa_queue_add_write_index_relaxed(
    hsa_queue_t * queue,
    uint64_t value);

uint64_t hsa_queue_add_write_index_release(
    hsa_queue_t * queue,
    uint64_t value);

```

Atomically increment the write index of a queue by an offset.

Parameters

queue

(in) Pointer to a queue.

value

(in) Value to add to the write index.

Returns

Previous value of the write index.

2.5.2.12 hsa_queue_store_read_index

```

void hsa_queue_store_read_index_relaxed(
    hsa_queue_t * queue,
    uint64_t value);

void hsa_queue_store_read_index_release(
    hsa_queue_t * queue,
    uint64_t value);

```

Atomically set the read index of a queue.

Parameters

queue

(in) Pointer to a queue.

value

(in) Value to assign to the read index.

Description

Modifications of the read index are not allowed if the queue is associated with an HSA component for which only the corresponding packet processor is permitted to update the read index. Invoking this operation in any other queue results in undefined behavior.

2.6 Architected Queuing Language Packets

The Architected Queuing Language (AQL) is a standard binary interface used to describe commands (such as a dispatch command or a barrier command). An AQL packet is a blob of data with a defined format encoding one command.

This section provides more information about AQL packet types and how the HSA API represents AQL packet types.

2.6.1 Dispatch packet

An application uses a Dispatch packet to submit tasks to an HSA component. The HSA API uses the `hsa_aql_dispatch_packet_t` type to represent a dispatch. A Dispatch packet can be in one of the following five states: *queued*, *launch*, *error*, *active* or *complete*. Figure 2.2 shows the state transition diagram.

Queued The queued state means that the format of the packet is not `HSA_AQL_PACKET_FORMAT_ALWAYS_RESERVED` nor `HSA_AQL_PACKET_FORMAT_INVALID`. If the *barrier* bit is set, the transition to launch state occurs after all the preceding kernels have completed execution. If the *barrier* bit is not set, then the transition to launch state occurs after the preceding kernels have completed their launch phase.

Launch The launch state indicates processing of the packet, but not execution of a task. This phase finalizes by applying an acquire memory fence. If an error is detected during launch, the queue is set into an error state and the event callback associated with the queue (if present) is invoked. The runtime passes an error status to the callback that indicates the source of the problem. The following error codes might be returned:

<code>HSA_STATUS_ERROR_INVALID_PACKET_FORMAT</code>	Malformed AQL packet
<code>HSA_STATUS_ERROR_OUT_OF_RESOURCES</code>	The packet processor is unable to allocate the resources required by the launch

Active The active state means that the kernel encoded by the packet has started execution. If an error is detected during this phase, a release fence is applied to the packet and its completion signal (if present) is assigned a negative value.

Error The error state means that an error was encountered during the launch or active phases.

Complete The complete state means that a memory release fence has been applied and the completion signal (if present) decremented.

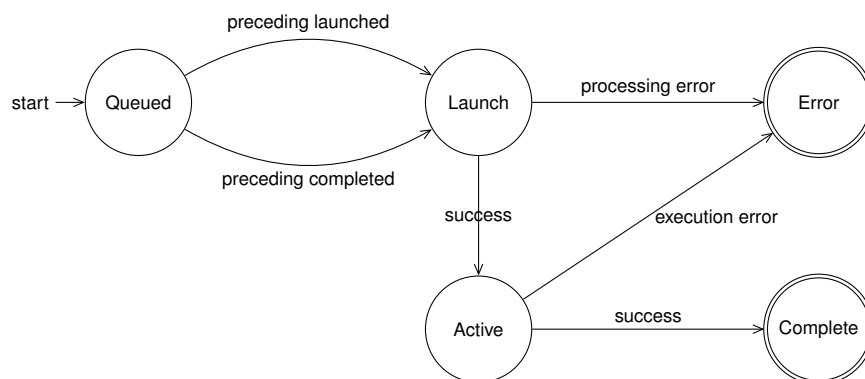


Figure 2.2: Packet State Diagram

2.6.1.1 Memory information

If the kernel being dispatched uses private or group segments, the HSA runtime requires the application to specify the sizes of the segments in the Dispatch packet. Manually calculating this information is not feasible and requires visual inspection of the input program, which itself may have been generated by a higher-level compiler. For this reason, the user must rely on the finalizer to get the corresponding segment sizes. For more information on determining segment sizes, refer to Section 3.1.

Another kind of HSA segment is the kernarg segment, used to pass arguments into a kernel. The application populates the *kernarg_address* field of the Dispatch packet with the address of a buffer of kernarg memory previously allocated. The buffer contains the actual parameters passed to the kernel being dispatched.

2.6.2 Agent Dispatch packet

Agent Dispatch packets dispatch jobs to HSA agents. The HSA API defines type `hsa_aql_agent_dispatch_packet_t` to represent Agent Dispatch packets. The states and state transitions for Agent Dispatch packets is identical to that of Dispatch packets.

2.6.3 Barrier packet

The Barrier packet `hsa_aql_barrier_packet_t` allows an application to specify up to five dependencies as `hsa_signal_t` objects and requires the HSA Packet Processor to resolve those dependencies before proceeding. The Barrier packet is a blocking packet in the sense that the processing of the Barrier packet completes the packet. This is unlike a Dispatch packet whose completion may occur at some future time after the packet has finished processing.

The HSA Packet Processor will not launch any further packets until the Barrier packet is complete. The execution phase for the Barrier packet completes when either of these two conditions have been met:

- Any of the dependent signals have been observed with a negative value after the Barrier packet launched, in which case the HSA Packet Processor assigns an error value to the *completion_signal*.
- All of the dependent signals have been observed with the value 0 after the Barrier packet launched (note that it is not required that all dependent signals are observed to be 0 at the same time).

As stated above, if any of the dependent signals have been assigned a negative value, the Barrier packet will indicate failure in its completion signal. The HSA Packet Processor assigns an error (negative) value to the *completion_signal*. For more information, refer to Section 2.4.

The state diagram of the Barrier packet is identical to that shown for Dispatch packets. However, each state might cause a different set of actions:

- No memory fence is applied in the launch phase.
- In the active phase, instead of executing a kernel, the packet waits for all the conditions to be met.
- No action is performed in the completion phase.

2.6.4 API

2.6.4.1 `hsa_aql_packet_format_t`

```
enum hsa_aql_packet_format_t
```

Packet type.

Values

HSA_AQL_PACKET_FORMAT_ALWAYS_RESERVED = 0

Initial format of a packets when the queue is created. Always reserved packet have never been assigned to the packet processor. From a functional view always reserved packets are equivalent to invalid packets. All queues support this packet format.

HSA_AQL_PACKET_FORMAT_INVALID = 1

The packet slot has been processed in the past, and has not been reassigned to the packet processor (is available). All queues support this packet format.

HSA_AQL_PACKET_FORMAT_DISPATCH = 2

Packet used by HSA agents for dispatching jobs to HSA components. Not all queues support packets of this type (see `hsa_queue_feature_t`).

HSA_AQL_PACKET_FORMAT_BARRIER = 3

Packet used by HSA agents to delay processing of subsequent packets, and to express complex dependencies between multiple packets. All queues support this packet format.

HSA_AQL_PACKET_FORMAT_AGENT_DISPATCH = 4

Packet used by HSA agents for dispatching jobs to HSA agents. Not all queues support packets of this type (see `hsa_queue_feature_t`).

2.6.4.2 hsa_fence_scope_t

```
enum hsa_fence_scope_t
```

Scope of the memory fence operation associated with a packet.

Values

HSA_FENCE_SCOPE_NONE = 0

No scope.

HSA_FENCE_SCOPE_COMPONENT = 1

The fence is applied with component scope for the global segment.

HSA_FENCE_SCOPE_SYSTEM = 2

The fence is applied with system scope for the global segment.

2.6.4.3 hsa_aql_packet_header_t

```
typedef struct hsa_aql_packet_header_s {
    hsa_aql_packet_format_t format : 8;
    uint16_t barrier : 1;
    hsa_fence_scope_t acquire_fence_scope : 2;
    hsa_fence_scope_t release_fence_scope : 2;
    uint16_t reserved : 3;
} hsa_aql_packet_header_t
```

AQL packet header.

Data Fields*format*

Packet type.

barrier

If set, the processing of the current packet only launches when all preceding packets (within the same queue) are complete.

acquire_fence_scope

Determines the scope and type of the memory fence operation applied before the packet enters the active phase. Must be HSA_FENCE_SCOPE_NONE for Barrier packets.

release_fence_scope

Determines the scope and type of the memory fence operation applied after kernel completion but before the packet is completed.

reserved

Must be 0.

2.6.4.4 hsa_aql_dispatch_packet_t

```
typedef struct hsa_aql_dispatch_packet_s {
    hsa_aql_packet_header_t header;
    uint16_t dimensions : 2;
    uint16_t reserved : 14;
    uint16_t workgroup_size_x;
    uint16_t workgroup_size_y;
    uint16_t workgroup_size_z;
    uint16_t reserved2;
    uint32_t grid_size_x;
    uint32_t grid_size_y;
    uint32_t grid_size_z;
    uint32_t private_segment_size;
    uint32_t group_segment_size;
    uint64_t kernel_object_address;
    uint64_t kernarg_address;
    uint64_t reserved3;
    hsa_signal_t completion_signal;
} hsa_aql_dispatch_packet_t
```

AQL Dispatch packet.

Data Fields*header*

Packet header.

dimensions

Number of dimensions specified in the grid size. Valid values are 1, 2, or 3.

reserved

Reserved, must be 0.

workgroup_size_x

X dimension of work-group (measured in work-items).

workgroup_size_y

Y dimension of work-group (measured in work-items).

workgroup_size_z

Z dimension of work-group (measured in work-items).

reserved2

Reserved. Must be 0.

grid_size_x

X dimension of grid (measured in work-items).

grid_size_y

Y dimension of grid (measured in work-items).

grid_size_z

Z dimension of grid (measured in work-items).

private_segment_size

Size (in bytes) of private memory allocation request (per work-item).

group_segment_size

Size (in bytes) of group memory allocation request (per work-group).

kernel_object_address

Address of an object in memory that includes an implementation-defined executable ISA image for the kernel.

kernarg_address

Pointer to a buffer containing the kernel arguments. Aligned to at least 16 bytes. The pointed buffer must remain allocated and not be modified until the kernel completes.

reserved3

Reserved. Must be 0.

completion_signal

Signal used to indicate completion of the job. Might be 0 (no signal).

2.6.4.5 hsa_aql_agent_dispatch_packet_t

```
typedef struct hsa_aql_agent_dispatch_packet_s {
    hsa_aql_packet_header_t header;
    uint16_t type;
    uint32_t reserved2;
    uint64_t return_location;
    uint64_t arg[4];
    uint64_t reserved3;
    hsa_signal_t completion_signal;
} hsa_aql_agent_dispatch_packet_t
```

Agent Dispatch packet.

Data Fields

header

Packet header.

type

The function to be performed by the destination HSA Agent. The type value is split into the following ranges: 0x0000:0x3FFF (vendor specific), 0x4000:0x7FFF (HSA runtime) 0x8000:0xFFFF (user registered function).

reserved2

Reserved. Must be 0.

return_location

Pointer to location to store the function return value(s) in.

arg

64-bit direct or indirect arguments.

reserved3

Reserved. Must be 0.

completion_signal

Signal used to indicate completion of the job. Might be 0 (no signal).

2.6.4.6 hsa_aql_barrier_packet_t

```
typedef struct hsa_aql_barrier_packet_s {
    hsa_aql_packet_header_t header;
    uint16_t reserved2;
    uint32_t reserved3;
    hsa_signal_t dep_signal[5];
    uint64_t reserved4;
    hsa_signal_t completion_signal;
} hsa_aql_barrier_packet_t
```

Barrier packet.

Data Fields

header

Packet header.

reserved2

Reserved. Must be 0.

reserved3

Reserved. Must be 0.

dep_signal

Array of dependent signal objects.

reserved4

Reserved. Must be 0.

completion_signal

Signal used to indicate completion of the job. Might be 0 (no signal).

2.7 Memory

One of the key features of HSA is its ability to share global pointers between the host application and code executing on the component. This ability means that an application can directly pass a pointer to memory allocated on the host to a kernel function dispatched to a component.

When a buffer created in the host is also accessed by a component, programmers are encouraged to *register* the corresponding address range beforehand by using the **hsa_memory_register** API. While kernels running on HSA devices can access any memory pointer allocated by means of standard libraries (for example, malloc in the C language), registering the buffer might result on improved access performance. When an HSA program no longer needs to access a registered buffer in a device, the application should deregister that virtual address range by using the appropriate HSA core API invocation.

The kernarg memory that AQL packet points to (see Section 2.6) holds information about any arguments required to execute AQL dispatch on a HSA component. While any system memory may be used for kernarg memory, implementation/platform specific optimizations are possible if HSA core runtime provided API are utilized for allocating and copying to the allocated kernarg memory.

2.7.1 API

2.7.1.1 hsa_region_t

```
typedef uint64_t hsa_region_t
```

A memory region represents a virtual memory interval that is visible to a particular agent, and contains properties about how memory is accessed from that agent. Two or more agents might have access to the same virtual memory interval, but the runtime uses multiple memory regions (one per agent) to represent the aperture.

2.7.1.2 hsa_segment_t

```
enum hsa_segment_t
```

Memory segments.

Values

HSA_SEGMENT_GLOBAL = 1
Global segment.

HSA_SEGMENT_PRIVATE = 2
Private segment.

HSA_SEGMENT_GROUP = 4
Group segment.

HSA_SEGMENT_KERNARG = 8
Kernarg segment.

HSA_SEGMENT_READONLY = 16
Readonly segment.

HSA_SEGMENT_IMAGE = 32
Image segment.

2.7.1.3 hsa_region_info_t

```
enum hsa_region_info_t
```

Attributes of a memory region.

Values

HSA_REGION_INFO_BASE_ADDRESS

Base (starting) address. The type of this attribute is void*.

HSA_REGION_INFO_SIZE

Size, in bytes. The type of this attribute is size_t.

HSA_REGION_INFO_NODE

NUMA node associated with this region. The type of this attribute is hsa_node_t

HSA_REGION_INFO_MAX_ALLOCATION_SIZE

Maximum allocation size in this region, in bytes. A value of 0 indicates that the host cannot allocate memory in the region. The type of this attribute is size_t.

HSA_REGION_INFO_SEGMENT

Segment where the memory in the region can be used. The type of this attribute is hsa_segment_t.

HSA_REGION_INFO_BANDWIDTH

Peak bandwidth for component access, in Mbs per second. The type of this attribute is uint32_t

HSA_REGION_INFO_CACHED

Cache levels that cache data in the region. The type of this attribute is bool[4].

2.7.1.4 hsa_region_get_info

```
hsa_status_t hsa_region_get_info(
    hsa_region_t region,
    hsa_region_info_t attribute,
    void * value);
```

Get the current value of an attribute of a region.

Parameters

region

(in) A valid region.

attribute

(in) Attribute to query.

value

(out) Pointer to a application-allocated buffer where to store the value of the attribute. If the buffer passed by the application is not large enough to hold the value of *attribute*, the behavior is undefined.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_REGION

If the region is invalid.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *attribute* is not a valid region attribute, or *value* is NULL.

2.7.1.5 hsa_agent_iterate_regions

```
hsa_status_t hsa_agent_iterate_regions(
    hsa_agent_t agent,
    hsa_status_t(*) (hsa_region_t region, void *data) callback,
    void * data);
```

Iterate over the memory regions that are visible to an agent, and invoke an application-defined callback on every iteration.

Parameters

agent

(in) A valid agent.

callback

(in) Callback to be invoked once per region that is visible from the agent.

data

(in) Application data that is passed to *callback* on every iteration. Might be NULL.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_AGENT

If the agent is invalid.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *callback* is NULL.

Description

If *callback* returns a status other than **HSA_STATUS_SUCCESS** for a particular iteration, the traversal stops and the function returns that status value.

2.7.1.6 hsa_memory_allocate

```
hsa_status_t hsa_memory_allocate(
    hsa_region_t region,
    size_t size,
    size_t alignment,
    void ** ptr);
```

Allocate a block of memory.

Parameters

region

(in) Region where to allocate memory from.

size

(in) Allocation size, in bytes. Allocation of size 0 is allowed and returns a NULL pointer.

alignment

(in) The alignment size (in bytes) for the base address of the resulting allocation. If the value is zero, no particular alignment will be applied. If the value is not zero, it must be a power of two that is not smaller than `sizeof(void*)`.

ptr

(out) A pointer to the location of where to return the pointer to the base of the allocated region of memory.

Return Values

`HSA_STATUS_SUCCESS`

The function has been executed successfully.

`HSA_STATUS_ERROR_NOT_INITIALIZED`

The runtime has not been initialized.

`HSA_STATUS_ERROR_OUT_OF_RESOURCES`

If no memory is available.

`HSA_STATUS_ERROR_INVALID_ARGUMENT`

If *alignment* is not a power of two, or *alignment* is smaller than `sizeof(void*)` but not 0, or *ptr* is NULL.

2.7.1.7 hsa_memory_free

```
hsa_status_t hsa_memory_free(
    void * ptr);
```

Deallocate a block of memory previously allocated using **hsa_memory_allocate**.

Parameters*ptr*

(in) Pointer to a memory block. If *ptr* is NULL, no action is performed.

Return Values

`HSA_STATUS_SUCCESS`

The function has been executed successfully.

`HSA_STATUS_ERROR_NOT_INITIALIZED`

The runtime has not been initialized.

2.7.1.8 hsa_memory_copy

```
hsa_status_t hsa_memory_copy(
    void * dst,
    const void * src,
    size_t size);
```

Copy block of memory.

Parameters*dst*

(out) A valid pointer to the destination array where the content is to be copied.

src

(in) A valid pointer to the source of data to be copied.

size(in) Number of bytes to copy. If *size* is 0, no copy is performed and the function returns success.**Return Values**

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If the source or destination pointers are NULL.

Description

Copying a number of bytes larger than the size of the memory regions pointed by *dst* or *src* results in undefined behavior.

2.7.1.9 hsa_memory_register

```
hsa_status_t hsa_memory_register(
    void * address,
    size_t size);
```

Register memory.

Parameters*address*

(in) A pointer to the base of the memory region to be registered. If a NULL pointer is passed, no operation is performed.

size(in) Requested registration size in bytes. A size of zero is only allowed if *address* is NULL.**Return Values**

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

If there is a failure in allocating the necessary resources.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *size* is 0 but *address* is not NULL.**Description**

Registering a buffer serves as an indication to the runtime that the passed buffer might be accessed from a component other than the host. Registration is a performance hint that allows the runtime implementation to know which buffers will be accessed by some of the components ahead of time.

Registrations should not overlap.

2.7.1.10 hsa_memory_deregister

```
hsa_status_t hsa_memory_deregister(  
    void * address,  
    size_t size);
```

Deregister memory previously registered using **hsa_memory_register**.

Parameters

address

(in) A pointer to the base of the memory region to be deregistered. If a NULL pointer is passed, no operation is performed.

size

(in) Size of the region to be deregistered.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

Description

If the memory interval being deregistered does not match a previous registration (start and end addresses), the behavior is undefined.

2.8 Extensions to the Core Runtime API

Extensions to the Core API can be optional (multi-vendor) or vendor specific. The difference is in the naming scheme used for the symbols (defines, structures, functions, etc.) associated with the function:

- Symbols for multi-vendor extensions defined in the global namespace must use the *hsa_ext_* prefix in their identifiers.
- Symbols for single vendor extensions defined in the global namespace must use the *hsa_svect_VENDOR_* prefix in their identifiers. Company names must be registered with the HSA Foundation, must be unique, and may be abbreviated to improve the readability of the symbols.

Any constant definitions in the extension (#define or enumeration values) use the same naming convention, except using all capital letters.

The symbols for all vendor extensions (both single-vendor and multi-vendor) are captured in the file **hsa/vendor_extensions.h**. This file is maintained by the HSA Foundation. This file includes the enumeration *hsa_extension_t* which defines a unique code for each vendor extension and multi-vendor extension. Vendors can reserve enumeration encodings through the HSA Foundation. Multi-vendor enumerations begin at the value of *HSA_EXT_START*, while single-vendor enumerations begin at *HSA_SVEXT_START*

2.8.1 API

2.8.1.1 *hsa_extension_t*

```
enum hsa_extension_t
```

HSA extensions.

Values

HSA_EXT_START = 0

Start of the multi vendor extension range.

HSA_EXT_FINALIZER = *HSA_EXT_START*

Finalizer extension. Finalizes the brig to compilation units that represent kernel and function code objects.

HSA_EXT_LINKER = 1

Linker extension.

HSA_EXT_IMAGES = 2

Images extension.

HSA_SVEXT_START = 10000

Start of the single vendor extension range.

2.8.1.2 *hsa_vendor_extension_query*

```
hsa_status_t hsa_vendor_extension_query(  
    hsa_extension_t extension,
```

```
void * extension_structure,
int * result);
```

Query vendor extensions.

Parameters

extension

(in) The vendor extension that is being queried.

extension_structure

(out) Extension structure.

result

(out) Pointer to memory location where to store the query result.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *extension* is not a valid value for a single vendor extension or *result* is NULL.

Description

If successful, the extension information is written with extension-specific information such as version information, function pointers, and data values. If the extension is not supported, the extension information is not modified.

2.8.1.3 hsa_extension_query

```
hsa_status_t hsa_extension_query(
    hsa_extension_t extension,
    int * result);
```

Query HSA extensions.

Parameters

extension

(in) The extension that is being queried.

result

(out) Pointer to memory location where to store the query result.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *extension* is not a valid value for a HSA extension or *result* is NULL.

2.8.2 Example

An example that shows a hypothetical single-vendor extension “Foo” registered by company “ACME”. The example includes four defines and two API functions. Note the use of the structure `hsa_svect_acme_foo_t` and how this interacts with the **`hsa_vendor_extension_query`** API call.

```
// Sample hsa/vendor_extensions.h
// Company name is "ACME" and extension is "Foo"
#define HSA_EXT_ACME_MYDEFINE1 0x1000
#define HSA_EXT_ACME_MYDEFINE2 0x0100

// The structure which defines the version, functions, and data for
// the extension:
typedef struct hsa_ext_acme_foo_s {
    int major_version; // major version number of the extension.
    int minor_version; // minor version number of the extension.
    // Function pointers:
    int (*function1) ( int p1, int *p2, float p3, int p4);
    int (*function2) ( int* p1, int p2);
    // Data:
    unsigned foo_data1;
} hsa_ext_acme_foo_t;

main() {
    struct hsa_ext_acme_foo_t acmeFoo;
    hsa_status_t status = hsa_vendor_extension_query(HSA_EXT_ACME_FOO, &acmeFoo);
    if (status == HSA_STATUS_SUCCESS) {
        (*(acmeFoo.function2))(0, 0);
    }
}
```


2.9 Common Definitions

2.9.1 API

2.9.1.1 hsa_powertwo8_t

```
typedef uint8_t hsa_powertwo8_t
```

Value expressed as a power of two.

2.9.1.2 hsa_powertwo_t

```
enum hsa_powertwo_t
```

Power of two between 1 and 256.

Values

HSA_POWER TWO_1 = 0

HSA_POWER TWO_2 = 1

HSA_POWER TWO_4 = 2

HSA_POWER TWO_8 = 3

HSA_POWER TWO_16 = 4

HSA_POWER TWO_32 = 5

HSA_POWER TWO_64 = 6

HSA_POWER TWO_128 = 7

HSA_POWER TWO_256 = 8

2.9.1.3 hsa_dim3_t

```
typedef struct hsa_dim3_s {
    uint32_t x;
    uint32_t y;
    uint32_t z;
} hsa_dim3_t
```

Three-dimensional coordinate.

Data Fields

x
X dimension.

y
Y dimension.

z
Z dimension.

2.9.1.4 hsa_dim_t

```
enum hsa_dim_t
```

Dimensions in a 3D space.

Values

HSA_DIM_X = 0
X dimension.

HSA_DIM_Y = 1
Y dimension.

HSA_DIM_Z = 2
Z dimension.

2.9.1.5 hsa_runtime_caller_t

```
typedef struct hsa_runtime_caller_s {
    uint64_t caller;
} hsa_runtime_caller_t
```

Opaque pointer that is passed to all runtime functions that use callbacks. The runtime passes this pointer as the first argument to all callbacks made by the function.

Data Fields

caller

Opaque pointer that is passed as the first argument to callback functions invoked by a runtime function.

2.9.1.6 hsa_runtime_alloc_data_callback_t

```
typedef hsa_status_t(* hsa_runtime_alloc_data_callback_t)(hsa_runtime_caller_t caller, size_t byte_size, void **address)
```

Call back function for allocating data.

Chapter 3

HSA Extensions Programming Guide

3.1 HSAIL Finalization

Note: The text in this section will be updated according to the API changes introduced in version 0.180 of the specification.

3.1.1 API

3.1.1.1 hsa_ext_brig_profile8_t

```
typedef uint8_t hsa_ext_brig_profile8_t
```

Profile is used to specify the kind of profile. This controls what features of HSAIL are supported. For more information see HSA Programmer's Reference Manual.

3.1.1.2 hsa_ext_brig_profile_t

```
enum hsa_ext_brig_profile_t
```

Profile kinds. For more information see HSA Programmer's Reference Manual.

Values

HSA_EXT_BRIG_PROFILE_BASE = 0

Base profile.

HSA_EXT_BRIG_PROFILE_FULL = 1

Full profile.

3.1.1.3 hsa_ext_brig_machine_model8_t

```
typedef uint8_t hsa_ext_brig_machine_model8_t
```

Machine model type. This controls the size of addresses used for segment and flat addresses. For more

information see HSA Programmer's Reference Manual.

3.1.1.4 hsa_ext_brig_machine_model_t

```
enum hsa_ext_brig_machine_model_t
```

Machine model kinds. For more information see HSA Programmer's Reference Manual.

Values

HSA_EXT_BRIG_MACHINE_SMALL = 0

Use 32 bit addresses for global segment and flat addresses.

HSA_EXT_BRIG_MACHINE_LARGE = 1

Use 64 bit addresses for global segment and flat addresses.

3.1.1.5 hsa_ext_brig_section_id32_t

```
typedef uint32_t hsa_ext_brig_section_id32_t
```

BRIG section id. The index into the array of sections in a BRIG module.

3.1.1.6 hsa_ext_brig_section_id_t

```
enum hsa_ext_brig_section_id_t
```

Predefined BRIG section kinds.

Values

HSA_EXT_BRIG_SECTION_DATA = 0

Textual character strings and byte data used in the module. Also contains variable length arrays of offsets into other sections that are used by entries in the hsa_code and hsa_operand sections. For more information see HSA Programmer's Reference Manual.

HSA_EXT_BRIG_SECTION_CODE = 1

All of the directives and instructions of the module. Most entries contain offsets to the hsa_operand or hsa_data sections. Directives provide information to the finalizer, and instructions correspond to HSAIL operations which the finalizer uses to generate executable ISA code. For more information see HSA Programmer's Reference Manual.

HSA_EXT_BRIG_SECTION_OPERAND = 2

The operands of directives and instructions in the code section. For example, immediate constants, registers and address expressions. For more information see HSA Programmer's Reference Manual.

3.1.1.7 hsa_ext_brig_section_header_t

```
typedef struct hsa_ext_brig_section_header_s {
    uint32_t byte_count;
    uint32_t header_byte_count;
    uint32_t name_length;
    uint8_t name[1];
} hsa_ext_brig_section_header_t
```

BRIG section header. Every section starts with a `hsa_ext_brig_section_header_t` which contains the section size, name and offset to the first entry. For more information see HSA Programmer's Reference Manual.

Data Fields

byte_count

Size in bytes of the section, including the size of the `hsa_ext_brig_section_header_t`. Must be a multiple of 4.

header_byte_count

Size of the header in bytes, which is also equal to the offset of the first entry in the section. Must be a multiple of 4.

name_length

Length of the section name in bytes.

name

Section name, *name_length* bytes long.

3.1.1.8 hsa_ext_brig_module_t

```
typedef struct hsa_ext_brig_module_s {
    uint32_t section_count;
    hsa_ext_brig_section_header_t * section[1];
} hsa_ext_brig_module_t
```

A module is the basic building block for HSAIL programs. When HSAIL is generated it is represented as a module.

Data Fields

section_count

Number of sections in the module. Must be at least 3.

section

A variable-sized array containing pointers to the BRIG sections. Must have *section_count* elements. Indexed by `hsa_ext_brig_section_id32_t`. The first three elements must be for the following predefined sections in the following order: `HSA_EXT_BRIG_SECTION_DATA`, `HSA_EXT_BRIG_SECTION_CODE`, `HSA_EXT_BRIG_SECTION_OPERAND`.

3.1.1.9 hsa_ext_brig_module_handle_t

```
typedef struct hsa_ext_brig_module_handle_s {
    uint64_t handle;
} hsa_ext_brig_module_handle_t
```

An opaque handle to the `hsa_ext_brig_module_t`.

Data Fields*handle*

HSA component specific handle to the brig module.

3.1.1.10 hsa_ext_brig_code_section_offset32_t

```
typedef uint32_t hsa_ext_brig_code_section_offset32_t
```

An entry offset into the code section of the BRIG module. The value is the byte offset relative to the start of the section to the beginning of the referenced entry. The value 0 is reserved to indicate that the offset does not reference any entry.

3.1.1.11 hsa_ext_exception_kind16_t

```
typedef uint16_t hsa_ext_exception_kind16_t
```

The set of exceptions supported by HSAIL. This is represented as a bit set.

3.1.1.12 hsa_ext_exception_kind_t

```
enum hsa_ext_exception_kind_t
```

HSAIL exception kinds. For more information see HSA Programmer's Reference Manual.

Values

HSA_EXT_EXCEPTION_INVALID_OPERATION = 1

Operations are performed on values for which the results are not defined. These are:

- Operations on signaling NaN (sNaN) floating-point values.
- Signalling comparisons: comparisons on quiet NaN (qNaN) floating-point values.
- Multiplication: mul(0.0, infinity) or mul(infinity, 0.0).
- Fused multiply add: fma(0.0, infinity, c) or fma(infinity, 0.0, c) unless c is a quiet NaN, in which case it is implementation-defined if an exception is generated.
- Addition, subtraction, or fused multiply add: magnitude subtraction of infinities, such as: add(positive infinity, negative infinity), sub(positive infinity, positive infinity).
- Division: div(0.0, 0.0) or div(infinity, infinity).
- Square root: sqrt(negative).
- Conversion: A cvt with a floating-point source type, an integer destination type, and a nonsaturating rounding mode, when the source value is a NaN, infinity, or the rounded value, after any flush to zero, cannot be represented precisely in the integer type of the destination.

HSA_EXT_EXCEPTION_DIVIDE_BY_ZERO = 2

A finite non-zero floating-point value is divided by zero. It is implementation defined if integer div or rem operations with a divisor of zero will generate a divide by zero exception.

HSA_EXT_EXCEPTION_OVERFLOW = 4

The floating-point exponent of a value is too large to be represented.

HSA_EXT_EXCEPTION_UNDERFLOW = 8

A non-zero tiny floating-point value is computed and either the ftz modifier is specified, or the ftz modifier was not specified and the value cannot be represented exactly.

HSA_EXT_EXCEPTION_INEXACT = 16

A computed floating-point value is not represented exactly in the destination. This can occur due to rounding. In addition, it is implementation defined if operations with the ftz modifier that cause a value to be flushed to zero generate the inexact exception.

3.1.1.13 hsa_ext_control_directive_present64_t

```
typedef uint64_t hsa_ext_control_directive_present64_t
```

Bit set of control directives supported in HSAIL. See HSA Programmer's Reference Manual description of control directives with the same name for more information. For control directives that have an associated value, the value is given by the field in `hsa_ext_control_directives_t`. For control directives that are only present or absent (such as `require_nopartial_workgroups`) they have no corresponding field as the presence of the bit in this mask is sufficient.

3.1.1.14 hsa_ext_control_directive_present_t

```
enum hsa_ext_control_directive_present_t
```

HSAIL control directive kinds. For more information see HSA Programmer's Reference Manual.

Values

HSA_EXT_CONTROL_DIRECTIVE_ENABLE_BREAK_EXCEPTIONS = 0

If not enabled then must be 0, otherwise must be non-0 and specifies the set of HSAIL exceptions that must have the BREAK policy enabled. If this set is not empty then the generated code may have lower performance than if the set is empty. If the kernel being finalized has any `enable_break_exceptions` control directives, then the values specified by this argument are unioned with the values in these control directives. If any of the functions the kernel calls have an `enable_break_exceptions` control directive, then they must be equal or a subset of, this union.

HSA_EXT_CONTROL_DIRECTIVE_ENABLE_DETECT_EXCEPTIONS = 1

If not enabled then must be 0, otherwise must be non-0 and specifies the set of HSAIL exceptions that must have the DETECT policy enabled. If this set is not empty then the generated code may have lower performance than if the set is empty. However, an implementation should endeavour to make the performance impact small. If the kernel being finalized has any `enable_detect_exceptions` control directives, then the values specified by this argument are unioned with the values in these control directives. If any of the functions the kernel calls have an `enable_detect_exceptions` control directive, then they must be equal or a subset of, this union.

HSA_EXT_CONTROL_DIRECTIVE_MAX_DYNAMIC_GROUP_SIZE = 2

If not enabled then must be 0, and any amount of dynamic group segment can be allocated for a dispatch, otherwise the value specifies the maximum number of bytes of dynamic group segment that can be allocated for a dispatch. If the kernel being finalized has any `maxdynamicssize` control directives, then the values must be the same, and must be the same as this argument if it is enabled. This value can be used by the finalizer to determine the maximum number of bytes of group memory used by each work-group by adding this value to the group memory required for all group segment variables used by the kernel and all functions it calls, and group memory used to implement other HSAIL features such as `fbarriers` and the `detect` exception operations. This can allow the finalizer to determine the expected number of work-groups that can be executed by a compute unit and allow more resources to be allocated to the work-items if it is known that fewer work-groups can be executed due to group memory limitations.

`HSA_EXT_CONTROL_DIRECTIVE_MAX_FLAT_GRID_SIZE = 4`

If not enabled then must be 0, otherwise must be greater than 0. Specifies the maximum number of work-items that will be in the grid when the kernel is dispatched. For more information see HSA Programmer's Reference Manual.

`HSA_EXT_CONTROL_DIRECTIVE_MAX_FLAT_WORKGROUP_SIZE = 8`

If not enabled then must be 0, otherwise must be greater than 0. Specifies the maximum number of work-items that will be in the work-group when the kernel is dispatched. For more information see HSA Programmer's Reference Manual.

`HSA_EXT_CONTROL_DIRECTIVE_REQUESTED_WORKGROUPS_PER_CU = 16`

If not enabled then must be 0, and the finalizer is free to generate ISA that may result in any number of work-groups executing on a single compute unit. Otherwise, the finalizer should attempt to generate ISA that will allow the specified number of work-groups to execute on a single compute unit. This is only a hint and can be ignored by the finalizer. If the kernel being finalized, or any of the functions it calls, has a `requested` control directive, then the values must be the same. This can be used to determine the number of resources that should be allocated to a single work-group and work-item. For example, a low value may allow more resources to be allocated, resulting in higher per work-item performance, as it is known there will never be more than the specified number of work-groups actually executing on the compute unit. Conversely, a high value may allocate fewer resources, resulting in lower per work-item performance, which is offset by the fact it allows more work-groups to actually execute on the compute unit.

`HSA_EXT_CONTROL_DIRECTIVE_REQUIRED_GRID_SIZE = 32`

If not enabled then all elements for `Dim3` must be 0, otherwise every element must be greater than 0. Specifies the grid size that will be used when the kernel is dispatched. For more information see HSA Programmer's Reference Manual.

`HSA_EXT_CONTROL_DIRECTIVE_REQUIRED_WORKGROUP_SIZE = 64`

If not enabled then all elements for `Dim3` must be 0, and the produced code can be dispatched with any legal work-group range consistent with the dispatch dimensions. Otherwise, the code produced must always be dispatched with the specified work-group range. No element of the specified range must be 0. It must be consistent with `required_dimensions` and `max_flat_workgroup_size`. If the kernel being finalized, or any of the functions it calls, has a `requiredworkgroupsize` control directive, then the values must be the same. Specifying a value can allow the finalizer to optimize work-group id operations, and if the number of work-items in the work-group is less than the `WAVESIZE` then barrier operations can be optimized to just a memory fence.

`HSA_EXT_CONTROL_DIRECTIVE_REQUIRED_DIM = 128`

If not enabled then must be 0 and the produced kernel code can be dispatched with 1, 2 or 3 dimensions. If enabled then the value is 1..3 and the code produced must only be dispatched with a dimension that matches. Other values are illegal. If the kernel being finalized, or any of the functions it calls, has a `requireddimsize` control directive, then the values must be the same. This can be used to optimize the code generated to compute the absolute and flat work-group and work-item id, and the `dim` HSAIL operations.

HSA_EXT_CONTROL_DIRECTIVE_REQUIRE_NO_PARTIAL_WORKGROUPS = 256

Specifies that the kernel must be dispatched with no partial work-groups. It can be placed in either a kernel or a function code block. This is only a hint and can be ignored by the finalizer.

It is undefined if the kernel is dispatched with any dimension of the grid size not being an exact multiple of the corresponding dimension of the work-group size.

A finalizer might be able to generate better code for `currentworkgroupsize` if it knows there are no partial work-groups, because the result becomes the same as the `workgroupsize` operation. An HSA component might be able to dispatch a kernel more efficiently if it knows there are no partial work-groups.

The control directive applies to the whole kernel and all functions it calls. It can appear multiple times in a kernel or function. If it appears in a function (including external functions), then it must also appear in all kernels that call that function (or have been specified when the finalizer was invoked), either directly or indirectly.

If `require_no_partial_work_groups` is specified when the finalizer is invoked, the kernel behaves as if the `require_no_partial_workgroups` control directive has been specified.

`require_no_partial_work_groups` does not have a field since having the bit set in `enabledControlDirectives` indicates that the `cpntrol` directive is present.

3.1.1.15 `hsa_ext_control_directives_t`

```
typedef struct hsa_ext_control_directives_s {
    hsa_ext_control_directive_present64_t enabled_control_directives;
    hsa_ext_exception_kind16_t enable_break_exceptions;
    hsa_ext_exception_kind16_t enable_detect_exceptions;
    uint32_t max_dynamic_group_size;
    uint32_t max_flat_grid_size;
    uint32_t max_flat_workgroup_size;
    uint32_t requested_workgroups_per_cu;
    hsa_dim3_t required_grid_size;
    hsa_dim3_t required_workgroup_size;
    uint8_t required_dim;
    uint8_t reserved[75];
} hsa_ext_control_directives_t
```

The `hsa_ext_control_directives_t` specifies the values for the HSAIL control directives. These control how the finalizer generates code. This struct is used both as an argument to `hsaFinalizeKernel` to specify values for the control directives, and is used in `HsaKernelCode` to record the values of the control directives that the finalizer used when generating the code which either came from the finalizer argument or explicit HSAIL control directives. See the definition of the control directives in HSA Programmer's Reference Manual which also defines how the values specified as finalizer arguments have to agree with the control directives in the HSAIL code.

Data Fields

enabled_control_directives

This is a bit set indicating which control directives have been specified. If the value is 0 then there are no control directives specified and the rest of the fields can be ignored. The bits are accessed using the `hsa_ext_control_directives_present_mask_t`. Any control directive that is not enabled in this bit set must have the value of all 0s.

enable_break_exceptions

If `enableBreakExceptions` is not enabled then must be 0, otherwise must be non-0 and specifies the set of HSAIL exceptions that must have the BREAK policy enabled. If this set is not empty then the generated code may have lower performance than if the set is empty. If the kernel being finalized has any `enablebreakexceptions` control directives, then the values specified by this argument are unioned with the values in these control directives. If any of the functions the kernel calls have an `enablebreakexceptions` control directive, then they must be equal or a subset of, this union.

enable_detect_exceptions

If `enableDetectExceptions` is not enabled then must be 0, otherwise must be non-0 and specifies the set of HSAIL exceptions that must have the DETECT policy enabled. If this set is not empty then the generated code may have lower performance than if the set is empty. However, an implementation should endeavour to make the performance impact small. If the kernel being finalized has any `enabledetectexceptions` control directives, then the values specified by this argument are unioned with the values in these control directives. If any of the functions the kernel calls have an `enabledetectexceptions` control directive, then they must be equal or a subset of, this union.

max_dynamic_group_size

If `maxDynamicGroupSize` is not enabled then must be 0, and any amount of dynamic group segment can be allocated for a dispatch, otherwise the value specifies the maximum number of bytes of dynamic group segment that can be allocated for a dispatch. If the kernel being finalized has any `maxdynamicsize` control directives, then the values must be the same, and must be the same as this argument if it is enabled. This value can be used by the finalizer to determine the maximum number of bytes of group memory used by each work-group by adding this value to the group memory required for all group segment variables used by the kernel and all functions it calls, and group memory used to implement other HSAIL features such as `fbarriers` and the detect exception operations. This can allow the finalizer to determine the expected number of work-groups that can be executed by a compute unit and allow more resources to be allocated to the work-items if it is known that fewer work-groups can be executed due to group memory limitations.

max_flat_grid_size

If `maxFlatGridSize` is not enabled then must be 0, otherwise must be greater than 0. See HSA Programmer's Reference Manual description of `maxflatgridsize` control directive.

max_flat_workgroup_size

If `maxFlatWorkgroupSize` is not enabled then must be 0, otherwise must be greater than 0. See HSA Programmer's Reference Manual description of `maxflatworkgroupsize` control directive.

requested_workgroups_per_cu

If `requestedWorkgroupsPerCu` is not enabled then must be 0, and the finalizer is free to generate ISA that may result in any number of work-groups executing on a single compute unit. Otherwise, the finalizer should attempt to generate ISA that will allow the specified number of work-groups to execute on a single compute unit. This is only a hint and can be ignored by the finalizer. If the kernel being finalized, or any of the functions it calls, has a `requested` control directive, then the values must be the same. This can be used to determine the number of resources that should be allocated to a single work-group and work-item. For example, a low value may allow more resources to be allocated, resulting in higher per work-item performance, as it is known there will never be more than the specified number of work-groups actually executing on the compute unit. Conversely, a high value may allocate fewer resources, resulting in lower per work-item performance, which is offset by the fact it allows more work-groups to actually execute on the compute unit.

required_grid_size

If not enabled then all elements for `Dim3` must be 0, otherwise every element must be greater than 0. See HSA Programmer's Reference Manual description of `requiredgridsize` control directive.

required_workgroup_size

If `requiredWorkgroupSize` is not enabled then all elements for `Dim3` must be 0, and the produced code can be dispatched with any legal work-group range consistent with the dispatch dimensions. Otherwise, the code produced must always be dispatched with the specified work-group range. No element of the specified range must be 0. It must be consistent with `required_dimensions` and `max_flat_workgroup_size`. If the kernel being finalized, or any of the functions it calls, has a `requiredworkgroupsize` control directive, then the values must be the same. Specifying a value can allow the finalizer to optimize work-group id operations, and if the number of work-items in the work-group is less than the `WAVESIZE` then barrier operations can be optimized to just a memory fence.

required_dim

If `requiredDim` is not enabled then must be 0 and the produced kernel code can be dispatched with 1, 2 or 3 dimensions. If enabled then the value is 1..3 and the code produced must only be dispatched with a dimension that matches. Other values are illegal. If the kernel being finalized, or any of the functions it calls, has a `requireddimsize` control directive, then the values must be the same. This can be used to optimize the code generated to compute the absolute and flat work-group and work-item id, and the dim HSAIL operations.

reserved

Reserved. Must be 0.

3.1.1.16 `hsa_ext_code_kind32_t`

```
typedef uint32_t hsa_ext_code_kind32_t
```

The kinds of code objects that can be contained in `hsa_ext_code_descriptor_t`.

3.1.1.17 `hsa_ext_code_kind_t`

```
enum hsa_ext_code_kind_t
```

Kinds of code object. For more information see HSA Programmer's Reference Manual.

Values

`HSA_EXT_CODE_NONE = 0`

Not a code object.

`HSA_EXT_CODE_KERNEL = 1`

HSAIL kernel that can be used with an AQL dispatch packet.

`HSA_EXT_CODE_INDIRECT_FUNCTION = 2`

HSAIL indirect function.

`HSA_EXT_CODE_RUNTIME_FIRST = 0x40000000`

HSA runtime code objects. For example, partially linked code objects.

`HSA_EXT_CODE_RUNTIME_LAST = 0x7fffffff`

`HSA_EXT_CODE_VENDOR_FIRST = 0x80000000`

Vendor specific code objects.

`HSA_EXT_CODE_VENDOR_LAST = 0xffffffff`

3.1.1.18 hsa_ext_program_call_convention_id32_t

```
typedef uint32_t hsa_ext_program_call_convention_id32_t
```

Each HSA component can support one or more call conventions. For example, an HSA component may have different call conventions that each use a different number of isa registers to allow different numbers of wavefronts to execute on a compute unit.

3.1.1.19 hsa_ext_program_call_convention_id_t

```
enum hsa_ext_program_call_convention_id_t
```

Kinds of program call convention IDs.

Values

HS_EXT_PROGRAM_CALL_CONVENTION_FINALIZER_DETERMINED = -1
Finalizer determined call convention ID.

3.1.1.20 hsa_ext_code_handle_t

```
typedef struct hsa_ext_code_handle_s {
    uint64_t handle;
} hsa_ext_code_handle_t
```

The 64-bit opaque code handle to the finalized code that includes the executable ISA for the HSA component. It can be used for the kernel dispatch packet kernel object address field.

Data Fields

handle

HSA component specific handle to the code.

3.1.1.21 hsa_ext_debug_information_handle_t

```
typedef struct hsa_ext_debug_information_handle_s {
    uint64_t handle;
} hsa_ext_debug_information_handle_t
```

An opaque handle to the debug information.

Data Fields

handle

HSA component specific handle to the debug information.

3.1.1.22 hsa_ext_code_descriptor_t

```
typedef struct hsa_ext_code_descriptor_s {
    hsa_ext_code_kind32_t code_type;
    uint32_t workgroup_group_segment_byte_size;
    uint64_t kernarg_segment_byte_size;
    uint32_t workitem_private_segment_byte_size;
    uint32_t workgroup_fbarrier_count;
    hsa_ext_code_handle_t code;
    hsa_powertwo8_t kernarg_segment_alignment;
    hsa_powertwo8_t group_segment_alignment;
    hsa_powertwo8_t private_segment_alignment;
    hsa_powertwo8_t wavefront_size;
    hsa_ext_program_call_convention_id32_t program_call_convention;
    hsa_ext_brig_module_handle_t module;
    hsa_ext_brig_code_section_offset32_t symbol;
    hsa_ext_brig_profile8_t hsail_profile;
    hsa_ext_brig_machine_model8_t hsail_machine_model;
    uint16_t reserved1;
    hsa_ext_debug_information_handle_t debug_information;
    char agent_vendor[24];
    char agent_name[24];
    uint32_t hsail_version_major;
    uint32_t hsail_version_minor;
    uint64_t reserved2;
    hsa_ext_control_directives_t control_directive;
} hsa_ext_code_descriptor_t
```

Provides the information about a finalization of the kernel or indirect function for a specific HSA component, and for indirect functions, a specific call convention of that HSA component.

Data Fields

code_type

Type of code object this code descriptor associated with.

workgroup_group_segment_byte_size

The amount of group segment memory required by a work-group in bytes. This does not include any dynamically allocated group segment memory that may be added when the kernel is dispatched.

kernarg_segment_byte_size

The size in bytes of the kernarg segment that holds the values of the arguments to the kernel.

workitem_private_segment_byte_size

The amount of memory required for the combined private, spill and arg segments for a work-item in bytes.

workgroup_fbarrier_count

Number of fbarrier's used in the kernel and all functions it calls. If the implementation uses group memory to allocate the fbarriers then that amount must already be included in the workgroupGroupSegment-ByteSize total.

code

The 64-bit opaque code handle to the finalized code that includes the executable ISA for the HSA component. It can be used for the kernel dispatch packet kernel object address field.

kernarg_segment_alignment

The maximum byte alignment of variables used by the kernel in the kernarg memory segment. Expressed as a power of two. Must be at least HSA_POWER TWO_16

group_segment_alignment

The maximum byte alignment of variables used by the kernel in the group memory segment. Expressed as a power of two. Must be at least `HSA_POWER TWO_16`

private_segment_alignment

The maximum byte alignment of variables used by the kernel in the private memory segment. Expressed as a power of two. Must be at least `HSA_POWER TWO_16`

wavefront_size

Wavefront size expressed as a power of two. Must be a power of 2 in range 1..64 inclusive. Used to support runtime query that obtains wavefront size, which may be used by application to allocated dynamic group memory and set the dispatch work-group size.

program_call_convention

Program call convention id this code descriptor holds.

module

BRIG module handle this code descriptor associated with.

symbol

BRIG directive offset this code descriptor associated with.

hsail_profile

The HSAIL profile defines which features are used. This information is from the HSAIL version directive. If this `hsa_ext_code_descriptor_t` is not generated from an **hsa_ext_finalize** then must still indicate what profile is being used.

hsail_machine_model

The HSAIL machine model gives the address sizes used by the code. This information is from the HSAIL version directive. If this `hsa_ext_code_descriptor_t` is not generated from an **hsa_ext_finalize** then must still indicate for what machine mode the code is generated.

reserved1

Reserved for BRIG target options if any are defined in the future. Must be 0.

debug_information

Opaque handle to debug information.

agent_vendor

The vendor of the HSA Component on which this Kernel Code object can execute. ISO/IEC 624 character encoding must be used. If the name is less than 24 characters then remaining characters must be set to 0.

agent_name

The vendor's name of the HSA Component on which this Kernel Code object can execute. ISO/IEC 624 character encoding must be used. If the name is less than 24 characters then remaining characters must be set to 0.

hsail_version_major

The HSAIL major version. This information is from the HSAIL version directive. If this `hsa_ext_code_descriptor_t` is not generated from an **hsa_ext_finalize** then must be 0.

hsail_version_minor

The HSAIL minor version. This information is from the HSAIL version directive. If this `hsa_ext_code_descriptor_t` is not generated from an **hsa_ext_finalize** then must be 0.

reserved2

Reserved. Must be 0.

control_directive

The values should be the actually values used by the finalizer in generating the code. This may be the union of values specified as finalizer arguments and explicit HSAIL control directives. If the finalizer chooses to ignore a control directive, and not generate constrained code, then the control directive should not be marked as enabled even though it was present in the HSAIL or finalizer argument. The values are intended to reflect the constraints that the code actually requires to correctly execute, not the values that were actually specified at finalize time.

3.1.1.23 `hsa_ext_finalization_request_t`

```
typedef struct hsa_ext_finalization_request_s {
    hsa_ext_brig_module_handle_t module;
    hsa_ext_brig_code_section_offset32_t symbol;
    hsa_ext_program_call_convention_id32_t program_call_convention;
} hsa_ext_finalization_request_t
```

Finalization request. Holds information about the module needed to be finalized. If the module contains an indirect function, also holds information about call convention id.

Data Fields

module

Handle to the `hsa_ext_brig_module_t`, which needs to be finalized.

symbol

Entry offset into the code section.

program_call_convention

If this finalization request is for indirect function, desired program call convention.

3.1.1.24 `hsa_ext_finalization_t`

```
typedef struct hsa_ext_finalization_s {
    uint32_t code_descriptor_count;
    uint32_t reserved1;
    hsa_ext_code_descriptor_t code_descriptors[1];
} hsa_ext_finalization_t
```

Finalization descriptor is the descriptor for the code object produced by the Finalizer and contains information that applies to the kernels/indirect function that were finalized.

Data Fields

code_descriptor_count

Number of code descriptors produced.

reserved1

Reserved. Must be 0.

code_descriptors

Dynamically sized array of code descriptors.

3.1.1.25 hsa_ext_symbol_definition_callback_t

```
typedef hsa_status_t(* hsa_ext_symbol_definition_callback_t)(hsa_runtime_caller_t caller, hsa_ext_brig_module_handle_t module, hsa_ext_brig_code_section_offset32_t symbol, hsa_ext_brig_module_handle_t *definition_module, hsa_ext_brig_module_t *definition_module_brig, hsa_ext_brig_code_section_offset32_t *definition_symbol)
```

Call back function to get the definition of a module scope variable/fbarrier or kernel/function.

3.1.1.26 hsa_ext_symbol_address_callback_t

```
typedef hsa_status_t(* hsa_ext_symbol_address_callback_t)(hsa_runtime_caller_t caller, hsa_ext_brig_module_handle_t module, hsa_ext_brig_code_section_offset32_t symbol, uint64_t *symbol_address)
```

Call back function to get the address of global segment variables, kernel table variable, indirect function table variable.

3.1.1.27 hsa_ext_error_message_callback_t

```
typedef hsa_status_t(* hsa_ext_error_message_callback_t)(hsa_runtime_caller_t caller, hsa_ext_brig_module_handle_t module, hsa_ext_brig_code_section_offset32_t statement, uint32_t indent_level, const char *message)
```

Call back function to get the string representation of the error message.

3.1.1.28 hsa_ext_finalize

```
hsa_status_t hsa_ext_finalize(
    hsa_runtime_caller_t caller,
    hsa_agent_t agent,
    uint32_t program_agent_id,
    uint32_t program_agent_count,
    size_t finalization_request_count,
    hsa_ext_finalization_request_t * finalization_request_list,
    hsa_ext_control_directives_t * control_directives,
    hsa_ext_symbol_definition_callback_t symbol_definition_callback,
    hsa_ext_symbol_address_callback_t symbol_address_callback,
    hsa_ext_error_message_callback_t error_message_callback,
    uint8_t optimization_level,
    const char * options,
    int debug_information,
    hsa_ext_finalization_t ** finalization);
```

Finalizes provided list of kernel and/or indirect functions.

Parameters

caller

(in) Opaque pointer and will be passed to all call back functions made by this call.

agent

(in) HSA agents for which code must be produced.

program_agent_id

(in) Program agent id.

program_agent_count

(in) Number of program agents.

finalization_request_count

(in) The number of kernels and/or indirect needed to be finalized.

finalization_request_list

(in) List of kernels and/or indirect functions needed to be finalized.

control_directives

(in) The control directives that can be specified to influence how the finalizer generates code. If NULL then no control directives are used. If this call is successful and control_directives is not NULL, then the resulting hsa_ext_code_descriptor_t object will have control directives which were used by the finalizer.

symbol_definition_callback

(in) Call back function to get the definition of a module scope variable/fbarrier or kernel/function.

symbol_address_callback

(in) Call back function to get the address of global segment variables, kernel table variables, indirect function table variable.

error_message_callback

(in) Call back function to get the string representation of the error message.

optimization_level

(in) An implementation defined value that control the level of optimization performed by the finalizer.

options

(in) Implementation defined options that can be specified to the finalizer.

debug_information

(in) The flag for including/excluding the debug information for *finalization*. 0 - exclude debug information, 1 - include debug information.

finalization

(out) the descriptor for the code object produced by the Finalizer and contains information that applies to all code entities in the program.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully.

HSA_EXT_STATUS_ERROR_DIRECTIVE_MISMATCH

If the directive in the control directive structure and in the HSA IL kernel mismatch or if the same directive is used with a different value in one of the functions used by this kernel.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *finalization_request_list* is NULL or invalid.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

If the finalize API cannot allocate memory for *finalization*.

HSA_EXT_STATUS_INFO_UNRECOGNIZED_OPTIONS

If the options are not recognized, no error is returned, just an info status is used to indicate invalid options.

Description

Invokes the finalizer on the provided list of kernels and indirect functions. A kernel can only be finalized once per program per agent. An indirect function can only be finalized once per program per agent per call convention. Only code for the HSA components specified when the program was created can be requested. The program must contain a definition for the requested kernels and indirect functions amongst the modules that have been added to the program.

3.1.1.29 hsa_ext_destroy_finalization

```
hsa_status_t hsa_ext_destroy_finalization(
    hsa_ext_finalization_t finalization);
```

Destroys the finalization.

Parameters

finalization

(in) Finalization to be destroyed.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *finalization* is NULL or does not point to a valid finalization structure.

HSA_STATUS_ERROR_RESOURCE_FREE

If some of the resources consumed during initialization by the runtime could not be freed.

3.1.1.30 hsa_ext_serialize_finalization

```
hsa_status_t hsa_ext_serialize_finalization(
    hsa_runtime_caller_t caller,
    hsa_agent_t agent,
    hsa_ext_finalization_t * finalization,
    hsa_runtime_alloc_data_callback_t alloc_serialize_data_callback,
    hsa_ext_error_message_callback_t error_message_callback,
    int debug_information,
    void * serialized_object);
```

Serializes the finalization.

Parameters

caller

(in) Opaque pointer and will be passed to all call back functions made by this call.

agent

(in) The HSA agent for which *finalization* must be serialized.

finalization

(in) Finalization descriptor to serialize.

alloc_serialize_data_callback

(in) Call back function for allocation.

error_message_callback

(in) Call back function to get the string representation of the error message.

debug_information

(in) The flag for including/excluding the debug information for *finalization*. 0 - exclude debug information, 1 - include debug information.

serialized_object

(out) Pointer to the serialized object.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *finalization* is either NULL or does not point to a valid finalization descriptor object.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

If no memory can be allocated for *serialized_object*.

Description

Serializes finalization descriptor for specified *agent*. The caller can set *debug_information* to 1 in order to include debug information of this finalization descriptor in the serialized object.

3.1.1.31 hsa_ext_deserialize_finalization

```
hsa_status_t hsa_ext_deserialize_finalization(
    hsa_runtime_caller_t caller,
    void * serialized_object,
    hsa_agent_t agent,
    uint32_t program_agent_id,
    uint32_t program_agent_count,
    hsa_ext_symbol_address_callback_t symbol_address_callback,
    hsa_ext_error_message_callback_t error_message_callback,
    int debug_information,
    hsa_ext_finalization_t ** finalization);
```

Deserializes the finalization.

Parameters

caller

(in) Opaque pointer and will be passed to all call back functions made by this call.

serialized_object

(in) Serialized object to be deserialized.

agent

(in) The HSA agent for which *finalization* must be serialized.

program_agent_id

(in) TODO.

program_agent_count

(in) TODO.

symbol_address_callback

(in) Call back function to get the address of global segment variables, kernel table variables, indirect function table variable.

error_message_callback

(in) Call back function to get the string representation of the error message.

debug_information

(in) The flag for including/excluding the debug information for *finalization*. 0 - exclude debug information, 1 - include debug information.

finalization

(out) Deserialized finalization descriptor.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *serialized_object* is either NULL, or is not valid, or the size is 0.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

If no memory can be allocated for *finalization*.

Description

Deserializes finalization descriptor for specified *agent*. The caller can set *debug_information* to 1 in order to include debug information of this finalization descriptor from the serialized object.

3.2 HSAIL Linking

Note: The text in this section will be updated according to the API changes introduced in version 0.180 of the specification.

3.2.1 API

3.2.1.1 hsa_ext_program_handle_t

```
typedef struct hsa_ext_program_handle_s {
    uint64_t handle;
} hsa_ext_program_handle_t
```

An opaque handle to the HSAIL program.

Data Fields

handle

HSA component specific handle to the program.

Description

An application can use the HSA runtime to create zero or more HSAIL programs, to which it can add zero or more HSAIL modules. HSAIL program manages linking of symbol declaration to symbol definitions between modules. In addition, the application can provide symbol definitions to an HSAIL program, and can obtain the address of symbols defined by the HSAIL program using the HSA runtime. An HSAIL program can be created with **hsa_ext_program_create**, which returns a handle to the created program **hsa_ext_program_handle_t**. A program handle has to be further used to add particular modules to the program using **hsa_ext_add_module**, perform various define, query and validation operations, finalize the program using **hsa_ext_finalize_program**. A program has to be destroyed once not needed any more using **hsa_ext_program_destroy**.

3.2.1.2 hsa_ext_program_create

```
hsa_status_t hsa_ext_program_create(
    hsa_agent_t * agents,
    uint32_t agent_count,
    hsa_ext_brig_machine_model8_t machine_model,
    hsa_ext_brig_profile8_t profile,
    hsa_ext_program_handle_t * program);
```

Creates an HSAIL program.

Parameters

agents

(in) One or more HSA components that are part of the HSA platform to create a program for.

agent_count

(in) Number of HSA components to create an HSAIL program for.

machine_model

(in) The kind of machine model this HSAIL program is created for. The machine model address size for the global segment must match the size used by the applications. All module added to the program must have the same machine model.

profile

(in) The kind of profile this HSAIL program is created for. All modules added to the program must have the same profile as the program.

program

(out) A valid pointer to a program handle for the HSAIL program created.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully, and an HSAIL program is created.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *agent* is NULL, or not valid. If *agent_count* is 0. If *machine_model* is not valid. If *profile* is not valid. In this case *program* will be NULL.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

If there is a failure to allocate resources required for program creation.

HSA_EXT_STATUS_INFO_ALREADY_INITIALIZED

If *program* is already a valid program. No error is returned, just an info status is used to indicate invalid options.

Description

Creates an HSAIL program. When an HSAIL program is created, one or more HSA components that are part of the HSA platform must be specified (*hsa_agent_t*), together with the machine model (*hsa_ext_brig_machine_model8_t*) and profile (*hsa_ext_brig_profile8_t*). The set of agents associated with the HSAIL being created cannot be changed after the program is created. The machine model address size for the global segment must match the size used by the application. All modules added to the program must have the same machine model and profile as the program. Once the program is created, the program handle (*hsa_ext_program_handle_t*) is returned. See *hsa_ext_program_handle_t* for more details.

3.2.1.3 hsa_ext_program_destroy

```
hsa_status_t hsa_ext_program_destroy(
    hsa_ext_program_handle_t program);
```

Destroys an HSAIL program.

Parameters*program*

(in) Program handle for the HSAIL program to be destroyed.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully, and an HSAIL program is destroyed.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *program* is not a valid *hsa_ext_program_handle_t* object.

HSA_STATUS_ERROR_RESOURCE_FREE

If *program* is already destroyed or has never been created.

Description

Destroys an HSAIL program pointed to by program handle *program*. When the program is destroyed, member code objects are destroyed as well.

3.2.1.4 hsa_ext_add_module

```
hsa_status_t hsa_ext_add_module(
    hsa_ext_program_handle_t program,
    hsa_ext_brig_module_t * brig_module,
    hsa_ext_brig_module_handle_t * module);
```

Adds an existing HSAIL module to an existing HSAIL program.

Parameters

program

(in) HSAIL program to add HSAIL module to.

brig_module

(in) HSAIL module to add to the HSAIL program.

module

(out) The handle for the *brig_module*.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully, and the module is added successfully.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If the *program* is not a valid HSAIL program. If the *brig_module* is not a valid HSAIL module. If the machine model and/or profile of the *brig_module* do not match the machine model and/or profile *program*.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

If there is a failure to allocate resources required for module addition.

HSA_EXT_STATUS_INFO_ALREADY_INITIALIZED

If the *brig_module* is already added to the HSAIL program. No error is returned, just an info status is used to indicate invalid options.

Description

Adds an existing HSAIL module to an existing HSAIL program. An HSAIL module is the unit of HSAIL generation, and can contain multiple symbol declarations and definitions. An HSAIL module can be added to zero or more HSAIL programs. Distinct instances of the symbols it defines are created within each program, and symbol declarations are only linked to the definitions provided by other modules in the same program. The same HSAIL module can be added to multiple HSAIL programs, which allows multiple instances of the same kernel and indirect functions that reference distinct allocations of global segment variables. The same HSAIL module cannot be added to the same HSAIL program more than once. The machine model and profile of the HSAIL module that is being added has to match the machine model and profile of the HSAIL program it is added to. HSAIL modules and their handles can be queried from the program using several query operations.

3.2.1.5 hsa_ext_finalize_program

```
hsa_status_t hsa_ext_finalize_program(
    hsa_ext_program_handle_t program,
    hsa_agent_t agent,
    size_t finalization_request_count,
```

```

hsa_ext_finalization_request_t * finalization_request_list,
hsa_ext_control_directives_t * control_directives,
hsa_ext_error_message_callback_t error_message_callback,
uint8_t optimization_level,
const char * options,
int debug_information);

```

Finalizes provided HSAIL program.

Parameters

program

(in) HSAIL program to be finalized.

agent

(in) Specific HSA component(s) to finalize program for.

finalization_request_count

(in) The number of kernels and indirect functions that are requested to be finalized.

finalization_request_list

(in) List of kernels and indirect functions that are requested to be finalized.

control_directives

(in) The control directives that can be specified to influence how the finalizer generates code. If NULL then no control directives are used. If this call is successful and *control_directives* is not NULL, then the resulting *hsa_ext_code_descriptor_t* object will have control directives which were used by the finalizer.

error_message_callback

(in) Call back function to get the string representation of the error message.

optimization_level

(in) An implementation defined value that control the level of optimization performed by the finalizer. For more information see HSA Programmer's Reference Manual.

options

(in) Implementation defined options that can be passed to the finalizer. For more information HSA Programmer's Reference Manual.

debug_information

(in) The flag for including/excluding the debug information for. 0 - exclude debug information, 1 - include debug information.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully, and the requested list of kernels/functions is finalized.

HSA_EXT_STATUS_ERROR_DIRECTIVE_MISMATCH

If the directive in the control directive structure and in the HSAIL kernel mismatch or if the same directive is used with a different value in one of the functions used by this kernel. The *error_message_callback* can be used to get the string representation of the error.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *agent* is NULL or invalid (if one of the specified HSA components is not the part of the HSAIL program). If the *program* is not a valid HSAIL program. If *finalization_request_list* is NULL or invalid. If *finalization_request_count* is 0. The *error_message_callback* can be used to get the string representation of the error.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

If there is a failure to allocate resources required for finalization. The *error_message_callback* can be used to get the string representation of the error.

HSA_EXT_STATUS_INFO_UNRECOGNIZED_OPTIONS

If the *options* or *optimization_level* are not recognized. No error is returned, just an info status is used to indicate invalid options.

Description

Finalizes provided HSA IL program. The HSA runtime finalizer can be used to generate code for kernels and indirect functions from a specific program for a specific HSA component. A kernel can only be finalized once per program per agent. An indirect function can only be finalized once per program per agent per call convention. Only code for HSA components specified when the program was created can be requested. The program must contain a definition for the requested kernels and indirect functions amongst the modules that have been added to the program. The modules of the program must collectively define all variables, fbarriers, kernels and functions referenced by operations in the code block. In addition, the caller of this function can specify control directives as an input argument, which will be passed to the finalizer. These control directives can be used for low-level performance tuning, for more information on control directives see HSA Programmer's Reference Manual.

3.2.1.6 hsa_ext_query_program_agent_id

```
hsa_status_t hsa_ext_query_program_agent_id(
    hsa_ext_program_handle_t program,
    hsa_agent_t agent,
    uint32_t * program_agent_id);
```

Queries HSA component's id for specified HSA component contained in specified HSA IL program.

Parameters

program

(in) HSA IL program to query HSA component's id from.

agent

(in) HSA component for which the id is queried.

program_agent_id

(out) HSA component's id contained in specified HSA IL program.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully, and HSA component's id is queried.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If provided *program* or *agent* is not valid. If *program_agent_id* is NULL.

3.2.1.7 hsa_ext_query_program_agent_count

```
hsa_status_t hsa_ext_query_program_agent_count(
    hsa_ext_program_handle_t program,
    uint32_t * program_agent_count);
```

Queries the number of HSA components contained in specified HSA IL program.

Parameters

program

(in) HSAIL program to query number of HSA components from.

program_agent_count

(out) Number of HSA components contained in specified HSAIL program.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully, and number of HSA components in the HSAIL program is queried.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If provided *program* is not valid. If *program_agent_count* is NULL.

3.2.1.8 hsa_ext_query_program_agents

```
hsa_status_t hsa_ext_query_program_agents(
    hsa_ext_program_handle_t program,
    uint32_t program_agent_count,
    hsa_agent_t * agents);
```

Queries specified number of HSA components contained in specified HSAIL program.

Parameters

program

(in) HSAIL program to query HSA agents from.

program_agent_count

(in) Number of HSA agents to query.

agents

(out) HSA agents contained in specified HSAIL program.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully, and HSA agents contained in the HSAIL program are queried.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If provided *program* is invalid.

3.2.1.9 hsa_ext_query_program_module_count

```
hsa_status_t hsa_ext_query_program_module_count(
    hsa_ext_program_handle_t program,
    uint32_t * program_module_count);
```

Queries the number of HSAIL modules contained in the HSAIL program.

Parameters

program

(in) HSAIL program to query number of HSAIL modules from.

program_module_count

(out) Number of HSAIL modules in specified HSAIL program.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully, and the number of HSAIL modules contained in the HSAIL program is queried.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If provided *program* is invalid. If *program_module_count* is NULL.

3.2.1.10 hsa_ext_query_program_modules

```
hsa_status_t hsa_ext_query_program_modules(
    hsa_ext_program_handle_t program,
    uint32_t program_module_count,
    hsa_ext_brig_module_handle_t * modules);
```

Queries specified number of HSAIL module handles contained in specified HSAIL program.

Parameters*program*

(in) HSAIL program to query HSAIL module handles from.

program_module_count

(in) Number of HSAIL module handles to query.

modules

(out) HSAIL module handles in specified HSAIL program.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully, and HSAIL module handles contained in the HSAIL program are queried.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If provided *program* is invalid. If *modules* is NULL.

3.2.1.11 hsa_ext_query_program_brig_module

```
hsa_status_t hsa_ext_query_program_brig_module(
    hsa_ext_program_handle_t program,
    hsa_ext_brig_module_handle_t module,
    hsa_ext_brig_module_t * brig_module);
```

Queries HSAIL module with specified handle that contained in specified HSAIL program.

Parameters*program*

(in) HSAIL program to query HSAIL modules from.

module

(in) HSAIL module handle for which to query the HSAIL module.

brig_module

(out) HSAIL module contained in specified HSAIL program.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully, and HSAIL module contained in the HSAIL program is queried.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If provided *program* is invalid, or *module* is invalid.

3.2.1.12 hsa_ext_query_call_convention

```
hsa_status_t hsa_ext_query_call_convention(
    hsa_ext_program_handle_t program,
    hsa_agent_t agent,
    hsa_ext_program_call_convention_id32_t * first_call_convention_id,
    uint32_t * call_convention_count);
```

Queries call convention IDs used for a specified HSA agent of a specified HSAIL program.

Parameters

program

(in) HSAIL program to query call convention IDs from.

agent

(in) HSA agent to query call convention IDs for.

first_call_convention_id

(out) Set of call convention IDs for specified HSA agent of a specified HSAIL program.

call_convention_count

(out) Number of call convention IDs available for specified HSA agent of a specified program.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully, and call convention IDs contained in the HSAIL program are queried.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If provided *program* is not valid. If *agent* is not valid or NULL. If *first_call_convention_id* is NULL. If *call_convention_count* is NULL.

3.2.1.13 hsa_ext_query_symbol_definition

```
hsa_status_t hsa_ext_query_symbol_definition(
    hsa_ext_program_handle_t program,
    hsa_ext_brig_module_handle_t module,
    hsa_ext_brig_code_section_offset32_t symbol,
    hsa_ext_brig_module_handle_t * definition_module,
    hsa_ext_brig_module_t * definition_module_brig,
    hsa_ext_brig_code_section_offset32_t * definition_symbol);
```

Queries the definition of a module scope variable/fbarrier or kernel/function for a specified HSAIL program.

Parameters

program

(in) HSAIL program to query symbol definition from.

module

(in) HSAIL module to query symbol definition from.

symbol

(in) Offset to query symbol definition from.

definition_module

(out) Queried HSAIL module handle.

definition_module_brig

(out) Queried HSAIL module.

definition_symbol

(out) Queried symbol.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully, and symbol definition contained in the HSAIL program is queried.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If provided *program* is not valid.

3.2.1.14 hsa_ext_define_program_allocation_global_variable_address

```
hsa_status_t hsa_ext_define_program_allocation_global_variable_address(
    hsa_ext_program_handle_t program,
    hsa_ext_brig_module_handle_t module,
    hsa_ext_brig_code_section_offset32_t symbol,
    hsa_ext_error_message_callback_t error_message_callback,
    void * address);
```

Defines global variable address in specified HSAIL program. Allows direct access to host variables from HSAIL.

Parameters*program*

(in) HSAIL program to define global variable address for.

module

(in) HSAIL module to define global variable address for.

symbol

(in) Offset in the HSAIL module to put the address on.

error_message_callback

(in) Call back function to get the string representation of the error message.

address

(in) Address to define in HSAIL program.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully, and specified global variable address is defined in specified HSAIL program.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If provided *program* is not valid. If *module* is not valid. If *address* is NULL.

3.2.1.15 hsa_ext_query_program_allocation_global_variable_address

```
hsa_status_t hsa_ext_query_program_allocation_global_variable_address(
    hsa_ext_program_handle_t program,
    hsa_ext_brig_module_handle_t module,
    hsa_ext_brig_code_section_offset32_t symbol,
    void ** address);
```

Queries global variable address from specified HSAIL program. Allows host program to directly access variables.

Parameters

program

(in) HSAIL program to query global variable address from.

module

(in) HSAIL module to query global variable address from.

symbol

(in) Offset in the HSAIL module to get the address from.

address

(out) Queried address.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully, and the global variable address is queried from specified HSAIL program.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If provided *program* is invalid, or *module* is invalid.

3.2.1.16 hsa_ext_define_agent_allocation_global_variable_address

```
hsa_status_t hsa_ext_define_agent_allocation_global_variable_address(
    hsa_ext_program_handle_t program,
    hsa_agent_t agent,
    hsa_ext_brig_module_handle_t module,
    hsa_ext_brig_code_section_offset32_t symbol,
    hsa_ext_error_message_callback_t error_message_callback,
    void * address);
```

Defines global variable address for specified HSA agent in specified HSAIL program. Allows direct access to host variables from HSAIL.

Parameters

program

(in) HSAIL program to define global variable address for.

agent

(in) HSA agent to define global variable address for.

module

(in) HSAIL module to define global variable address for.

symbol

(in) Offset in the HSAIL module to put the address on.

error_message_callback

(in) Call back function to get the string representation of the error message.

address

(in) Address to define for HSA agent in HSAIL program.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully, and specified global variable address is defined for specified HSA agent in specified HSAIL program.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If provided *program* is not valid. If *agent* is NULL or not valid If *module* is not valid. If *address* is NULL.

3.2.1.17 hsa_ext_query_agent_global_variable_address

```
hsa_status_t hsa_ext_query_agent_global_variable_address(
    hsa_ext_program_handle_t program,
    hsa_agent_t agent,
    hsa_ext_brig_module_handle_t module,
    hsa_ext_brig_code_section_offset32_t symbol,
    void ** address);
```

Queries global variable address for specified HSA agent from specified HSAIL program. Allows host program to directly access variables.

Parameters

program

(in) HSAIL program to query global variable address from.

agent

(in) HSA agent to query global variable address from.

module

(in) HSAIL module to query global variable address from.

symbol

(in) Offset in the HSAIL module to get the address from.

address

(out) Queried address.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully, and the global variable address is queried from specified HSAIL program.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If provided *program* is not valid. If *agent* is NULL or not valid If *module* is not valid. If *address* is NULL.

3.2.1.18 hsa_ext_define_readonly_variable_address

```
hsa_status_t hsa_ext_define_readonly_variable_address(
    hsa_ext_program_handle_t program,
    hsa_agent_t agent,
    hsa_ext_brig_module_handle_t module,
    hsa_ext_brig_code_section_offset32_t symbol,
    hsa_ext_error_message_callback_t error_message_callback,
    void * address);
```

Defines readonly variable address for specified HSA agent in specified HSAIL program. Allows direct access to host variables from HSAIL.

Parameters

program

(in) HSAIL program to define readonly variable address for.

agent

(in) HSA agent to define readonly variable address for.

module

(in) HSAIL module to define readonly variable address for.

symbol

(in) Offset in the HSAIL module to put the address on.

error_message_callback

(in) Call back function to get the string representation of the error message.

address

(in) Address to define for HSA agent in HSAIL program.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully, and specified readonly variable address is defined for specified HSA agent in specified HSAIL program.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If provided *program* is not valid. If *agent* is NULL or not valid If *module* is not valid. If *address* is NULL.

3.2.1.19 hsa_ext_query_readonly_variable_address

```
hsa_status_t hsa_ext_query_readonly_variable_address(
    hsa_ext_program_handle_t program,
    hsa_agent_t agent,
    hsa_ext_brig_module_handle_t module,
    hsa_ext_brig_code_section_offset32_t symbol,
    void ** address);
```

Queries readonly variable address for specified HSA agent from specified HSAIL program. Allows host program to directly access variables.

Parameters

program

(in) HSAIL program to query readonly variable address from.

agent

(in) HSA agent to query readonly variable address from.

module

(in) HSAIL module to query readonly variable address from.

symbol

(in) Offset in the HSAIL module to get the address from.

address

(out) Queried address.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully, and the readonly variable address is queried from specified HSAIL program.

HSA_STATUS_ERROR_INVALID_ARGUMENTIf provided *program* is not valid. If *agent* is NULL or not valid If *module* is not valid. If *address* is NULL.**3.2.1.20 hsa_ext_query_kernel_descriptor_address**

```

hsa_status_t hsa_ext_query_kernel_descriptor_address(
    hsa_ext_program_handle_t program,
    hsa_ext_brig_module_handle_t module,
    hsa_ext_brig_code_section_offset32_t symbol,
    void ** address);

```

Queries kernel descriptor address from specified HSAIL program. Needed to create a dispatch packet.

Parameters*program*

(in) HSAIL program to query kernel descriptor address from.

module

(in) HSAIL module to query kernel descriptor address from.

symbol

(in) Offset in the HSAIL module to get the address from.

address

(out) Queried address.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully, and the kernel descriptor address is queried from specified HSAIL program.

HSA_STATUS_ERROR_INVALID_ARGUMENTIf provided *program* is invalid, or *module* is invalid.**3.2.1.21 hsa_ext_query_indirect_function_descriptor_address**

```

hsa_status_t hsa_ext_query_indirect_function_descriptor_address(
    hsa_ext_program_handle_t program,
    hsa_ext_brig_module_handle_t module,
    hsa_ext_brig_code_section_offset32_t symbol,
    void ** address);

```

Queries indirect function descriptor address from specified HSAIL program. Allows host program to perform indirect function table variable initialization.

Parameters

program

(in) HSAIL program to query indirect function descriptor address from.

module

(in) HSAIL module to query indirect function descriptor address from.

symbol

(in) Offset in the HSAIL module to get the address from.

address

(out) Queried address.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully, and the indirect function descriptor address is queried from specified HSAIL program.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If provided *program* is invalid, or *module* is invalid.

3.2.1.22 hsa_ext_validate_program

```

hsa_status_t hsa_ext_validate_program(
    hsa_ext_program_handle_t program,
    hsa_ext_error_message_callback_t error_message_callback);

```

Validates HSAIL program with specified HSAIL program handle. Checks if all declarations and definitions match, if there is at most one definition. The caller decides when to call validation routines. For example, it can be done in the debug mode.

Parameters

program

(in) HSAIL program to validate.

error_message_callback

(in) Call back function to get the string representation of the error message.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully, and specified HSAIL program is a valid program.

If specified HSAIL program is not valid, refer to the error call back function to get string representation of the failure.

3.2.1.23 hsa_ext_validate_program_module

```
hsa_status_t hsa_ext_validate_program_module(
    hsa_ext_program_handle_t program,
    hsa_ext_brig_module_handle_t module,
    hsa_ext_error_message_callback_t error_message_callback);
```

Validates specified HSAIL module with specified HSAIL module handle for specified HSAIL program. Checks if BRIG for specified module is legal: operation operand type rules, etc. For more information about BRIG see HSA Programmer's Reference Manual. The caller decides when to call validation routines. For example, it can be done in the debug mode.

Parameters

program

(in) HSAIL program to validate HSAIL module in.

module

(in) HSAIL module handle to validate.

error_message_callback

(in) Call back function to get the string representation of the error message.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully, and specified HSAIL module is a valid module..

If the module is not valid, refer to the error call back function to get string representation of the failure.

3.2.1.24 hsa_ext_serialize_program

```
hsa_status_t hsa_ext_serialize_program(
    hsa_runtime_caller_t caller,
    hsa_ext_program_handle_t program,
    hsa_runtime_alloc_data_callback_t alloc_serialize_data_callback,
    hsa_ext_error_message_callback_t error_message_callback,
    int debug_information,
    void * serialized_object);
```

Serializes specified HSAIL program. Used for offline compilation.

Parameters

caller

(in) Opaque pointer to the caller of this function.

program

(in) HSAIL program to be serialized.

alloc_serialize_data_callback

(in) Call back function for memory allocation.

error_message_callback

(in) Call back function to get the string representation of the error message (if any).

debug_information

(in) The flag for including/excluding the debug information for *finalization*. 0 - exclude debug information, 1

- include debug information.

serialized_object

(out) Pointer to the serialized object.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully, and specified program is serialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *program* is not a valid program.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

If there is a failure to allocate resources required for serialization. The *error_message_callback* can be used to get the string representation of the error.

3.2.1.25 hsa_ext_program_allocation_symbol_address_t

```
typedef hsa_status_t(* hsa_ext_program_allocation_symbol_address_t)(hsa_runtime_caller_t caller, const char *name, uint64_t *symbol_address)
```

Call back function to get program's address of global segment variables, kernel table variable, indirect function table variable based on the symbolic name.

3.2.1.26 hsa_ext_agent_allocation_symbol_address_t

```
typedef hsa_status_t(* hsa_ext_agent_allocation_symbol_address_t)(hsa_runtime_caller_t caller, hsa_agent_t agent, const char *name, uint64_t *symbol_address)
```

Call back function to get agents's address of global segment variables, kernel table variable, indirect function table variable based on the symbolic name.

3.2.1.27 hsa_ext_deserialize_program

```
hsa_status_t hsa_ext_deserialize_program(
    hsa_runtime_caller_t caller,
    void * serialized_object,
    hsa_ext_program_allocation_symbol_address_t program_allocation_symbol_address,
    hsa_ext_agent_allocation_symbol_address_t agent_allocation_symbol_address,
    hsa_ext_error_message_callback_t error_message_callback,
    int debug_information,
    hsa_ext_program_handle_t ** program);
```

Deserializes the HSAIL program from a given serialized object. Used for offline compilation. Includes call back functions, where call back functions take symbolic name, this allows symbols defined by application to be relocated.

Parameters*caller*

(in) Opaque pointer to the caller of this function.

serialized_object

(in) Serialized HSAIL program.

program_allocation_symbol_address

(in) Call back function to get program's address of global segment variables, kernel table variable, indirect function table variable based on the symbolic name. Allows symbols defined by application to be relocated.

agent_allocation_symbol_address

(in) Call back function to get agent's address of global segment variables, kernel table variable, indirect function table variable based on the symbolic name. Allows symbols defined by application to be relocated.

error_message_callback

(in) Call back function to get the string representation of the error message.

debug_information(in) The flag for including/excluding the debug information for *finalization*. 0 - exclude debug information, 1 - include debug information.*program*

(out) Deserialized HSAIL program.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully, and HSAIL program is deserialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *serialized_object* is either NULL, or is not valid, or the size is 0.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

If there is a failure to allocate resources required for deserialization. The *error_message_callback* can be used to get the string representation of the error.

3.3 Images and Samplers

An HSA runtime uses an image handle `hsa_ext_image_handle_t` to access images. The image handle references the image data in memory and records information about resource layout and other properties. HSA decouples the storage of the image data and the description of how the device interprets that data. This allows the application developer to control the location of the image data storage and manage memory more efficiently.

The HSA image format is specified using a format descriptor (`hsa_ext_image_format_t`) that contains information about the image channel type and the channel order. The image channel type describes how the data is to be interpreted along with the bit size, and image channel order describes the number and the order. Not all image channel types and channel order combinations are valid on a HSA agent. All HSA agents have to support a required minimum set of image formats. For more information, refer to the HSA Programmer's Reference Manual[1]. An application can use **`hsa_ext_image_get_format_capability`** to query a runtime to obtain image format capabilities.

An implementation-independent image format descriptor (`hsa_ext_image_descriptor_t`) is composed of geometry along with the image format. The image descriptor is used to inquire the runtime for the HSA component-specific image data size and alignment details by calling **`hsa_ext_image_get_info`** for the purpose of determining the implementation's storage requirements.

The memory requirements (`hsa_ext_image_info_t`) include the size of the memory needed as well as any alignment constraints. An application can either allocate new memory for the image data, or sub-allocate a memory block from an existing memory if the memory size allows. Before the image data is used, an HSA agent-specific image handle must be created using it and if necessary, cleared and prepared according to the intended use.

A HSA agent-specific image handle (`hsa_ext_image_handle_t`) is used by the HSAIL language for reading or writing using HSAIL **`rd image`**, **`ldimage`** and **`stimage`** operations. **`hsa_ext_image_create_handle`** creates an image handle from a implementation-independent image format descriptor and independently allocated image data that conforms to the requirements provided by **`hsa_ext_image_get_info`**.

It must be noted that while the image data technically accessible from its pointer in the raw form, the data layout and organization is agent-specific and should be treated as opaque. The internal implementation of an optimal image data organization could vary depending on the attributes of the image format descriptor. As a result, there are no guarantees on the data layout when accessed from another HSA agent. The only reliable way to import or export image data from optimally organized images is to copy their data to and from a linearly organized data layout in memory, as specified by the image's format attributes.

The HSA Runtime provides interfaces to allow operations on images. Image data transfer to and from memory with a linear layout can be performed using **`hsa_ext_image_export`** and **`hsa_ext_image_import`** respectively. A portion of an image could be copied to another image using **`hsa_ext_image_copy`**. An image can be cleared using **`hsa_ext_image_clear`**. It is the application's responsibility to ensure proper synchronization and preparation of images on accesses from other image operations. See HSA System Architecture spec 2.13 for the HSA Image memory model.

A HSA agent-specific sampler handle (`hsa_ext_sampler_handle_t`) is used by the HSAIL language to describe how images are processed by the **`rdimage`** HSAIL operation. **`hsa_ext_sampler_create_handle`** creates a sampler handle from an agent independent sampler descriptor (`hsa_ext_sampler_descriptor_t`).

3.3.1 API

3.3.1.1 `hsa_ext_image_handle_t`

```
typedef struct hsa_ext_image_handle_s {
    uint64_t handle;
} hsa_ext_image_handle_t
```

Image handle, populated by **hsa_ext_image_create_handle**. Images handles are only unique within an agent, not across agents.

Data Fields

handle

HSA component specific handle to the image.

3.3.1.2 hsa_ext_image_format_capability_t

```
enum hsa_ext_image_format_capability_t
```

Image format capability returned by **hsa_ext_image_get_format_capability**.

Values

HSA_EXT_IMAGE_FORMAT_NOT_SUPPORTED = 0x0

Images of this format are not supported.

HSA_EXT_IMAGE_FORMAT_READ_ONLY = 0x1

Images of this format can be accessed for read operations.

HSA_EXT_IMAGE_FORMAT_WRITE_ONLY = 0x2

Images of this format can be accessed for write operations.

HSA_EXT_IMAGE_FORMAT_READ_WRITE = 0x4

Images of this format can be accessed for read and write operations.

HSA_EXT_IMAGE_FORMAT_READ_MODIFY_WRITE = 0x8

Images of this format can be accessed for read-modify-write operations.

HSA_EXT_IMAGE_FORMAT_ACCESS_INVARIANT_IMAGE_DATA = 0x10

Images of this format are guaranteed to have consistent data layout regardless of the how it is accessed by the HSA agent.

3.3.1.3 hsa_ext_image_info_t

```
typedef struct hsa_ext_image_info_s {
    size_t image_size;
    size_t image_alignment;
} hsa_ext_image_info_t
```

Agent-specific image size and alignment requirements. This structure stores the agent-dependent image data sizes and alignment, and populated by **hsa_ext_image_get_info**.

Data Fields

image_size

Component specific image data size in bytes.

image_alignment

Component specific image data alignment in bytes.

3.3.1.4 hsa_ext_image_access_permission_t

```
enum hsa_ext_image_access_permission_t
```

Defines how the HSA device expects to access the image. The access pattern used by the HSA agent specified in **hsa_ext_image_create_handle**.

Values

HSA_EXT_IMAGE_ACCESS_PERMISSION_READ_ONLY

Image handle is to be used by the HSA agent as read-only using an HSAIL roimg type.

HSA_EXT_IMAGE_ACCESS_PERMISSION_WRITE_ONLY

Image handle is to be used by the HSA agent as write-only using an HSAIL woimg type.

HSA_EXT_IMAGE_ACCESS_PERMISSION_READ_WRITE

Image handle is to be used by the HSA agent as read and/or write using an HSAIL rwimg type.

3.3.1.5 hsa_ext_image_geometry_t

```
enum hsa_ext_image_geometry_t
```

Geometry associated with the HSA image (image dimensions allowed in HSA). The enumeration values match the HSAIL BRIG type BrigImageGeometry.

Values

HSA_EXT_IMAGE_GEOMETRY_1D = 0

One-dimensional image addressed by width coordinate.

HSA_EXT_IMAGE_GEOMETRY_2D = 1

Two-dimensional image addressed by width and height coordinates.

HSA_EXT_IMAGE_GEOMETRY_3D = 2

Three-dimensional image addressed by width, height, and depth coordinates.

HSA_EXT_IMAGE_GEOMETRY_1DA = 3

Array of one-dimensional images with the same size and format. 1D arrays are addressed by index and width coordinate.

HSA_EXT_IMAGE_GEOMETRY_2DA = 4

Array of two-dimensional images with the same size and format. 2D arrays are addressed by index and width and height coordinates.

HSA_EXT_IMAGE_GEOMETRY_1DB = 5

One-dimensional image interpreted as a buffer with specific restrictions.

HSA_EXT_IMAGE_GEOMETRY_2DDEPTH = 6

Two-dimensional depth image addressed by width and height coordinates.

HSA_EXT_IMAGE_GEOMETRY_2DADEPTH = 7

Array of two-dimensional depth images with the same size and format. 2D arrays are addressed by index and width and height coordinates.

3.3.1.6 hsa_ext_image_channel_type_t


```
enum hsa_ext_image_channel_type_t
```

Component type associated with the image. See Image section in HSA Programming Reference Manual for definitions on each component type. The enumeration values match the HSAIL BRIG type BrigImageChannelType.

Values

```
HSA_EXT_IMAGE_CHANNEL_TYPE_SNORM_INT8 = 0
HSA_EXT_IMAGE_CHANNEL_TYPE_SNORM_INT16 = 1
HSA_EXT_IMAGE_CHANNEL_TYPE_UNORM_INT8 = 2
HSA_EXT_IMAGE_CHANNEL_TYPE_UNORM_INT16 = 3
HSA_EXT_IMAGE_CHANNEL_TYPE_UNORM_INT24 = 4
HSA_EXT_IMAGE_CHANNEL_TYPE_UNORM_SHORT_555 = 5
HSA_EXT_IMAGE_CHANNEL_TYPE_UNORM_SHORT_565 = 6
HSA_EXT_IMAGE_CHANNEL_TYPE_UNORM_SHORT_101010 = 7
HSA_EXT_IMAGE_CHANNEL_TYPE_SIGNED_INT8 = 8
HSA_EXT_IMAGE_CHANNEL_TYPE_SIGNED_INT16 = 9
HSA_EXT_IMAGE_CHANNEL_TYPE_SIGNED_INT32 = 10
HSA_EXT_IMAGE_CHANNEL_TYPE_UNSIGNED_INT8 = 11
HSA_EXT_IMAGE_CHANNEL_TYPE_UNSIGNED_INT16 = 12
HSA_EXT_IMAGE_CHANNEL_TYPE_UNSIGNED_INT32 = 13
HSA_EXT_IMAGE_CHANNEL_TYPE_HALF_FLOAT = 14
HSA_EXT_IMAGE_CHANNEL_TYPE_FLOAT = 15
```

3.3.1.7 hsa_ext_image_channel_order_t

```
enum hsa_ext_image_channel_order_t
```

Image component order associated with the image. See Image section in HSA Programming Reference Manual for definitions on each component order. The enumeration values match the HSAIL BRIG type BrigImageChannelOrder.

Values

```
HSA_EXT_IMAGE_CHANNEL_ORDER_A = 0
HSA_EXT_IMAGE_CHANNEL_ORDER_R = 1
HSA_EXT_IMAGE_CHANNEL_ORDER_RX = 2
HSA_EXT_IMAGE_CHANNEL_ORDER_RG = 3
HSA_EXT_IMAGE_CHANNEL_ORDER_RGX = 4
HSA_EXT_IMAGE_CHANNEL_ORDER_RA = 5
HSA_EXT_IMAGE_CHANNEL_ORDER_RGB = 6
HSA_EXT_IMAGE_CHANNEL_ORDER_RGBX = 7
```

```

HSA_EXT_IMAGE_CHANNEL_ORDER_RGBA = 8
HSA_EXT_IMAGE_CHANNEL_ORDER_BGRA = 9
HSA_EXT_IMAGE_CHANNEL_ORDER_ARGB = 10
HSA_EXT_IMAGE_CHANNEL_ORDER_ABGR = 11
HSA_EXT_IMAGE_CHANNEL_ORDER_SRGB = 12
HSA_EXT_IMAGE_CHANNEL_ORDER_SRGBX = 13
HSA_EXT_IMAGE_CHANNEL_ORDER_SRGBA = 14
HSA_EXT_IMAGE_CHANNEL_ORDER_SBGRA = 15
HSA_EXT_IMAGE_CHANNEL_ORDER_INTENSITY = 16
HSA_EXT_IMAGE_CHANNEL_ORDER_LUMINANCE = 17
HSA_EXT_IMAGE_CHANNEL_ORDER_DEPTH = 18
HSA_EXT_IMAGE_CHANNEL_ORDER_DEPTH_STENCIL = 19

```

3.3.1.8 hsa_ext_image_format_t

```

typedef struct hsa_ext_image_format_s {
    hsa_ext_image_channel_type_t channel_type;
    hsa_ext_image_channel_order_t channel_order;
} hsa_ext_image_format_t

```

Image format descriptor (attributes of the image format).

Data Fields

channel_type

Channel type of the image.

channel_order

Channel order of the image.

3.3.1.9 hsa_ext_image_descriptor_t

```

typedef struct hsa_ext_image_descriptor_s {
    hsa_ext_image_geometry_t geometry;
    size_t width;
    size_t height;
    size_t depth;
    size_t array_size;
    hsa_ext_image_format_t format;
} hsa_ext_image_descriptor_t

```

Implementation-independent HSA Image descriptor.

Data Fields

geometry

Geometry of the image.

width

Width of the image in components.

height

Height of the image in components, only used if geometry is 2D or higher.

depth

Depth of the image in slices, only used if geometry is 3D depth = 0 is same as depth = 1.

array_size

Number of images in the image array, only used if geometry is 1DArray and 2DArray.

format

Format of the image.

3.3.1.10 hsa_ext_image_range_t

```
typedef struct hsa_ext_image_range_s {
    uint32_t width;
    uint32_t height;
    uint32_t depth;
} hsa_ext_image_range_t
```

Three-dimensional image range description.

Data Fields*width*

The width for an image range (in coordinates).

height

The height for an image range (in coordinates).

depth

The depth for an image range (in coordinates).

3.3.1.11 hsa_ext_image_region_t

```
typedef struct hsa_ext_image_region_s {
    hsa_dim3_t image_offset;
    hsa_ext_image_range_t image_range;
} hsa_ext_image_region_t
```

Image region description. Used by image operations such as import, export, copy, and clear.

Data Fields*image_offset*

Offset in the image (in coordinates).

image_range

Dimensions of the image range (in coordinates).

3.3.1.12 hsa_ext_sampler_handle_t

```
typedef struct hsa_ext_sampler_handle_s {
    uint64_t handle;
} hsa_ext_sampler_handle_t
```

Sampler handle. Samplers are populated by **hsa_ext_sampler_create_handle**. Sampler handles are only unique within an agent, not across agents.

Data Fields

handle

Component-specific HSA sampler.

3.3.1.13 hsa_ext_sampler_addressing_mode_t

```
enum hsa_ext_sampler_addressing_mode_t
```

Sampler address modes. The sampler address mode describes the processing of out-of-range image coordinates. The values match the HSAIL BRIG type BrigSamplerAddressing.

Values

HSA_EXT_SAMPLER_ADDRESSING_UNDEFINED = 0

Out-of-range coordinates are not handled.

HSA_EXT_SAMPLER_ADDRESSING_CLAMP_TO_EDGE = 1

Clamp out-of-range coordinates to the image edge.

HSA_EXT_SAMPLER_ADDRESSING_CLAMP_TO_BORDER = 2

Clamp out-of-range coordinates to the image border.

HSA_EXT_SAMPLER_ADDRESSING_REPEAT = 3

Wrap out-of-range coordinates back into the valid coordinate range.

HSA_EXT_SAMPLER_ADDRESSING_MIRRORED_REPEAT = 4

Mirror out-of-range coordinates back into the valid coordinate range.

3.3.1.14 hsa_ext_sampler_coordinate_mode_t

```
enum hsa_ext_sampler_coordinate_mode_t
```

Sampler coordinate modes. The enumeration values match the HSAIL BRIG BRIG_SAMPLER_COORD bit in the type BrigSamplerModifier.

Values

HSA_EXT_SAMPLER_COORD_NORMALIZED = 0

Coordinates are all in the range of 0.0 to 1.0.

HSA_EXT_SAMPLER_COORD_UNNORMALIZED = 1

Coordinates are all in the range of 0 to (dimension-1).

3.3.1.15 hsa_ext_sampler_filter_mode_t

```
enum hsa_ext_sampler_filter_mode_t
```

Sampler filter modes. The enumeration values match the HSAIL BRIG type `BrigSamplerFilter`.

Values

`HSA_EXT_SAMPLER_FILTER_NEAREST = 0`

Filter to the image element nearest (in Manhattan distance) to the specified coordinate.

`HSA_EXT_SAMPLER_FILTER_LINEAR = 1`

Filter to the image element calculated by combining the elements in a 2x2 square block or 2x2x2 cube block around the specified coordinate. The elements are combined using linear interpolation.

3.3.1.16 `hsa_ext_sampler_descriptor_t`

```
typedef struct hsa_ext_sampler_descriptor_s {
    hsa_ext_sampler_coordinate_mode_t coordinate_mode;
    hsa_ext_sampler_filter_mode_t filter_mode;
    hsa_ext_sampler_addressing_mode_t address_mode;
} hsa_ext_sampler_descriptor_t
```

Implementation-independent sampler descriptor.

Data Fields

coordinate_mode

Sampler coordinate mode describes the normalization of image coordinates.

filter_mode

Sampler filter type describes the type of sampling performed.

address_mode

Sampler address mode describes the processing of out-of-range image coordinates.

3.3.1.17 `hsa_ext_image_get_format_capability`

```
hsa_status_t hsa_ext_image_get_format_capability(
    hsa_agent_t agent,
    const hsa_ext_image_format_t * image_format,
    hsa_ext_image_geometry_t image_geometry,
    uint32_t * capability_mask);
```

Retrieve image format capabilities for the specified image format on the specified HSA component.

Parameters

agent

(in) HSA agent to be associated with the image.

image_format

(in) Image format.

image_geometry

(in) Geometry of the image.

capability_mask

(out) Image format capability bit-mask.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *agent*, *image_format*, or *capability_mask* are NULL.

Description

If successful, the queried image format's capabilities bit-mask is written to the location specified by *capability_mask*. See `hsa_ext_image_format_capability_t` to determine all possible capabilities that can be reported in the bit mask.

3.3.1.18 hsa_ext_image_get_info

```
hsa_status_t hsa_ext_image_get_info(
    hsa_agent_t agent,
    const hsa_ext_image_descriptor_t * image_descriptor,
    hsa_ext_image_access_permission_t access_permission,
    hsa_ext_image_info_t * image_info);
```

Inquires the required HSA component-specific image data details from a implementation independent image descriptor.

Parameters

agent

(in) HSA agent to be associated with the image.

image_descriptor

(in) Implementation-independent image descriptor describing the image.

access_permission

(in) Access permission of the image by the HSA agent.

image_info

(out) Image info size and alignment requirements that the HSA agent requires.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If any of the arguments is NULL.

HSA_EXT_STATUS_ERROR_IMAGE_FORMAT_UNSUPPORTED

If the HSA agent does not support the image format specified by the descriptor.

HSA_EXT_STATUS_ERROR_IMAGE_SIZE_UNSUPPORTED

If the HSA agent does not support the image dimensions specified by the format descriptor.

Description

If successful, the queried HSA agent-specific image data info is written to the location specified by *image_info*. Based on the implementation the optimal image data size and alignment requirements could vary depending on the image attributes specified in *image_descriptor*.

The implementation must return the same image info requirements for different access permissions with exactly the same image descriptor as long as **hsa_ext_image_get_format_capability** reports HSA_EXT_IMAGE_FORMAT_ACCESS_INVARIANT_IMAGE_DATA for the image format specified in the image descriptor.

3.3.1.19 hsa_ext_image_create_handle

```
hsa_status_t hsa_ext_image_create_handle(
    hsa_agent_t agent,
    const hsa_ext_image_descriptor_t * image_descriptor,
    const void * image_data,
    hsa_ext_image_access_permission_t access_permission,
    hsa_ext_image_handle_t * image_handle);
```

Creates a agent-defined image handle from an implementation-independent image descriptor and a agent-specific image data. The image access defines how the HSA agent expects to use the image and must match the HSAIL image handle type used by the agent.

Parameters

agent

(in) HSA agent to be associated with the image.

image_descriptor

(in) Implementation-independent image descriptor describing the image.

image_data

(in) Address of the component-specific image data.

access_permission

(in) Access permission of the image by the HSA agent.

image_handle

(out) Agent-specific image handle.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If any of the arguments is NULL.

HSA_EXT_STATUS_ERROR_IMAGE_FORMAT_UNSUPPORTED

If the HSA agent does not have the capability to support the image format using the specified *agent_access*.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

If the HSA agent cannot create the specified handle because it is out of resources.

Description

If successful, the image handle is written to the location specified by *image_handle*. The image data memory must be allocated using the previously queried **hsa_ext_image_get_info** memory requirements with the same HSA agent and implementation-independent image descriptor.

The image data is not initialized and any previous memory contents is preserved. The memory management of image data is the application's responsibility and can only be freed until the memory is no longer needed and any image handles using it are destroyed.

access_permission defines how the HSA agent expects to use the image handle. The image format specified in the image descriptor must be capable by the HSA agent for the intended permission.

Image handles with different permissions can be created using the same image data with exactly the same image descriptor as long as **HSA_EXT_IMAGE_FORMAT_ACCESS_INVARIANT_IMAGE_DATA** is reported by **hsa_ext_image_get_format_capability** for the image format specified in the image descriptor. Images of non-linear s-form channel order can share the same image data with its equivalent linear non-s form channel order, provided the rest of the image descriptor parameters are identical.

If necessary, an application can use image operations (import, export, copy, clear) to prepare the image for the intended use regardless of the access permissions.

3.3.1.20 hsa_ext_image_import

```
hsa_status_t hsa_ext_image_import(
    hsa_agent_t agent,
    const void * src_memory,
    size_t src_row_pitch,
    size_t src_slice_pitch,
    hsa_ext_image_handle_t dst_image_handle,
    const hsa_ext_image_region_t * image_region,
    const hsa_signal_t * completion_signal);
```

Imports a linearly organized image data from memory directly to an image handle.

Parameters

agent

(in) HSA agent to be associated with the image.

src_memory

(in) Source memory.

src_row_pitch

(in) Number of bytes in one row of the source memory.

src_slice_pitch

(in) Number of bytes in one slice of the source memory.

dst_image_handle

(in) Destination Image handle.

image_region

(in) Image region to be updated.

completion_signal

(in) Signal to set when the operation is completed.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *agent*, *src_memory* or *image_region* are NULL.

Description

This operation updates the image data referenced by the image handle from the source memory. The size of the data imported from memory is implicitly derived from the image region.

If *completion_signal* is NULL, the operation occurs synchronously. Otherwise the function returns immediately and the completion signal is signaled when the operation completes.

If *src_row_pitch* is smaller than the destination region width (in bytes), then *src_row_pitch* = region width.

If *src_slice_pitch* is smaller than the destination region width * region height (in bytes), then *src_slice_pitch* = region width * region height.

It is the application's responsibility to avoid out of bounds memory access.

None of the source memory or image data memory in the previously created **hsa_ext_image_create_handle** image handle can overlap. Overlapping of any of the source and destination memory within the import operation produces undefined results.

3.3.1.21 hsa_ext_image_export

```
hsa_status_t hsa_ext_image_export(
    hsa_agent_t agent,
    hsa_ext_image_handle_t src_image_handle,
    void * dst_memory,
    size_t dst_row_pitch,
    size_t dst_slice_pitch,
    const hsa_ext_image_region_t * image_region,
    const hsa_signal_t * completion_signal);
```

Export image data from the image handle directly to memory organized linearly.

Parameters

agent

(in) HSA agent to be associated with the image.

src_image_handle

(in) Source image handle.

dst_memory

(in) Destination memory.

dst_row_pitch

(in) Number of bytes in one row of the destination memory.

dst_slice_pitch

(in) Number of bytes in one slice of the destination memory.

image_region

(in) Image region to be exported.

completion_signal

(in) Signal to set when the operation is completed.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENTIf *agent*, *dst_memory* or *image_region* are NULL.**Description**

The operation updates the destination memory with the image data in the image handle. The size of the data exported to memory is implicitly derived from the image region.

If *completion_signal* is NULL, the operation occurs synchronously. Otherwise the function returns immediately and the completion signal is signaled when the operation completes.

If *dst_row_pitch* is smaller than the source region width (in bytes), then *dst_row_pitch* = region width.

If *dst_slice_pitch* is smaller than the source region width * region height (in bytes), then *dst_slice_pitch* = region width * region height.

It is the application's responsibility to avoid out of bounds memory access.

None of the destination memory or image data memory in the previously created **hsa_ext_image_create_handle** image handle can overlap. Overlapping of any of the source and destination memory within the export operation produces undefined results.

3.3.1.22 hsa_ext_image_copy

```
hsa_status_t hsa_ext_image_copy(
    hsa_agent_t agent,
    hsa_ext_image_handle_t src_image_handle,
    hsa_ext_image_handle_t dst_image_handle,
    const hsa_ext_image_region_t * image_region,
    const hsa_signal_t * completion_signal);
```

Copies a region from one image to another.

Parameters*agent*

(in) HSA agent to be associated with the image.

src_image_handle

(in) Source image handle.

dst_image_handle

(in) Destination image handle.

image_region

(in) Image region to be copied.

completion_signal

(in) Signal to set when the operation is completed.

Return Values**HSA_STATUS_SUCCESS**

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *agent* or *image_region* are NULL.

Description

The operation copies the image data from the source image handle to the destination image handle. The size of the image data copied is implicitly derived from the image region.

If *completion_signal* is NULL, the operation occurs synchronously. Otherwise the function returns immediately and the completion signal is signaled when the operation completes.

It is the application's responsibility to avoid out of bounds memory access.

The source and destination handles must have been previously created using **hsa_ext_image_create_handle**. The source and destination image data memory are not allowed to be the same. Overlapping any of the source and destination memory produces undefined results.

The source and destination image formats don't have to match; appropriate format conversion is performed automatically. The source and destination images must be of the same geometry.

3.3.1.23 hsa_ext_image_clear

```
hsa_status_t hsa_ext_image_clear(
    hsa_agent_t agent,
    hsa_ext_image_handle_t image_handle,
    const float data[4],
    const hsa_ext_image_region_t * image_region,
    const hsa_signal_t * completion_signal);
```

Clears the image to a specified 4-component floating point data.

Parameters

agent

(in) HSA agent to be associated with the image.

image_handle

(in) Image to be cleared.

data

(in) 4-component clear value in floating point format.

image_region

(in) Image region to clear.

completion_signal

(in) Signal to set when the operation is completed.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *agent* or *image_region* are NULL.

Description

The operation clears the elements of the image with the data specified. The lowest bits of the data (number of bits depending on the image component type) are stored in the cleared image are based on the image component order. The size of the image data cleared is implicitly derived from the image region.

If *completion_signal* is NULL, the operation occurs synchronously. Otherwise the function returns immediately and the completion signal is signaled when the operation completes.

It is the application's responsibility to avoid out of bounds memory access.

Clearing an image automatically performs value conversion on the provided floating point values as is appropriate for the image format used.

For images of UNORM types, the floating point values must be in the [0..1] range. For images of SNORM types, the floating point values must be in the [-1..1] range. For images of UINT types, the floating point values are rounded down to an integer value. For images of SRGB types, the clear data is specified in a linear space, which is appropriately converted by the Runtime to sRGB color space.

Specifying clear value outside of the range representable by an image format produces undefined results.

3.3.1.24 hsa_ext_image_destroy_handle

```
hsa_status_t hsa_ext_image_destroy_handle(
    hsa_agent_t agent,
    hsa_ext_image_handle_t * image_handle);
```

Destroys the specified image handle.

Parameters

agent

(in) HSA agent to be associated with the image.

image_handle

(in) Image handle.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If *agent* or *image_handle* is NULL.

Description

If successful, the image handle previously created using **hsa_ext_image_create_handle** is destroyed.

Destroying the image handle does not free the associated image data.

The image handle should not be destroyed while there are references to it queued for execution or currently being used in a dispatch. Failure to properly track image data lifetime causes undefined results due to premature image handle deletion.

3.3.1.25 hsa_ext_sampler_create_handle

```

hsa_status_t hsa_ext_sampler_create_handle(
    hsa_agent_t agent,
    const hsa_ext_sampler_descriptor_t * sampler_descriptor,
    hsa_ext_sampler_handle_t * sampler_handle);

```

Create an HSA component-defined sampler handle from a component-independent sampler descriptor.

Parameters

agent

(in) HSA agent to be associated with the image.

sampler_descriptor

(in) Implementation-independent sampler descriptor.

sampler_handle

(out) Component-specific sampler handle.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If any of the arguments is NULL.

HSA_STATUS_ERROR_OUT_OF_RESOURCES

If the HSA agent cannot create the specified handle because it is out of resources.

Description

If successful, the sampler handle is written to the location specified by the sampler handle.

3.3.1.26 hsa_ext_sampler_destroy_handle

```

hsa_status_t hsa_ext_sampler_destroy_handle(
    hsa_agent_t agent,
    hsa_ext_sampler_handle_t * sampler_handle);

```

Destroys the specified sampler handle.

Parameters

agent

(in) HSA agent to be associated with the image.

sampler_handle

(in) Sampler handle.

Return Values

HSA_STATUS_SUCCESS

The function has been executed successfully.

HSA_STATUS_ERROR_NOT_INITIALIZED

The runtime has not been initialized.

HSA_STATUS_ERROR_INVALID_ARGUMENT

If any of the arguments is NULL.

Description

If successful, the sampler handle previously created using **hsa_ext_sampler_create_handle** is destroyed.

The sampler handle should not be destroyed while there are references to it queued for execution or currently being used in a dispatch.

3.4 Component Initiated Dispatches

Due to architected support for a queue and design of AQL, HSA supports component-initiated dispatch, which is the ability for a kernel to dispatch a new kernel by writing an AQL packet directly to a user queue. In simple use cases, the AQL packet can be created on the host and passed as a parameter to the kernel. This eliminates the need to do dynamic memory allocation on the component, but has the limitation that the problem fanout must be known at the time the first kernel is launched (so that the AQL packets can be preallocated). HSA also supports more advanced use cases where the AQL packet is dynamically allocated (including the memory space for kernel arguments and spill/arg/private space) on the component. This usage model obviously requires dynamic component-side memory allocation, for both host and component memory.

Some requirements to do component-initiated dispatch:

- Ability to dynamically choose a kernel to dispatch: Let us assume for example that there are three kernels (A, B and C). If the host launches A, then the user has the choice of launching B or C, or even A in case of recursion. So, the user should be able to get the ISA and segment size (HsaAqlKernel) from the corresponding BRIG dynamically. [caveat: The code sample here does not show how we can do this. It assumes that the HsaAqlKernel is being passed as an argument to the parent kernel (A in this case)]
- Ability to dynamically allocate memory from the shader: We need to allocate memory for AQLPacket, different kernel segments in the AQLPacket, kernel arguments, and so forth.
- Ability for a finalizer to identify a default HSA queue to write AQLPacket: The HSA queue information resides in the runtime layer of the stack. This needs to be exchanged with the compiler so it can be stored in the global space. This way, when the compiler sees the queue, it knows where to pick the HSA queue information to write the AQL-Packet.
- Ability to notify the completion of all the component-initiated dispatches on the host:
 - The beginning of execution of the child kernel may or may not wait for the parent kernel's completion. This is determined by the user and could be algorithm dependent.
 - If the parent (initiated from host) kernel finishes successfully, it means all kernels it initiated also finished successfully.
 - To implement this, we need to track the list of kernels launched from the parent. Change the status of parent to complete, only if parent and all its child kernels have completed successfully.

Implementations that support component initiated dispatches will need to support these requirements. If the implementation supports the stated requirements, the following actions will allow a component to initiate a dispatch:

- The queue and `hsa_ext_code_descriptor_t` (describing the kernel to launch) can be passed as arguments to the parent (the one launched from the host) kernel. If the dispatch is to the same queue, it is accessible via an HSAIL instruction.
- If not, get the HsaAqlKernel from the BRIG for the kernel that is chosen to be dynamically dispatched.
- When new work is to be created, the HSAIL code would:
 - Use the kernel dynamic memory allocator to allocate a new AQLPacket.
 - Use inline HSAIL to replicate the functionality of the `HsaInitAQLPacket` function. We could perhaps provide an HSAIL library to implement this functionality. Recall this function:
 - * Copies some fields from the HsaAqlKernel structure (for example, the kernel ISA) to the AQLPacket

- * Uses a host allocator to allocate memory for the kernel arguments
 - * Uses a component allocator to allocate memory for spill, private, and arg segments
- The HSAIL knows the signature of the called function and can fill in the AQL packet with regular HSAIL global store instructions.
- The HSA queue is architected, so the HSAIL can use memory store instructions to dispatch the kernel for dispatch. Depending how the user queues are configured, atomic accesses might be necessary to handle contention with other writers. Note that, if the queue information is not passed in as an argument, the default queue can be chosen by the finalizer as it was exchanged earlier from the runtime layer.
- We also need to handle deallocation of the kernel arguments and spill/private/arg space after the kernel completes.
- On the host, check if the parent has finished. If the parent has finished successfully, then it means that all the child kernels have finished successfully too. If the parent or any of the child kernels failed, an error code will be returned.

Index - Core APIs

hsa_agent_get_info, 13
 hsa_agent_iterate_regions, 42
 hsa_extension_query, 47
 hsa_init, 7
 hsa_iterate_agents, 14
 hsa_memory_allocate, 42
 hsa_memory_copy, 43
 hsa_memory_deregister, 45
 hsa_memory_free, 43
 hsa_memory_register, 44
 hsa_queue_add_write_index_acq_rel, 32
 hsa_queue_add_write_index_acquire, 32
 hsa_queue_add_write_index_relaxed, 32
 hsa_queue_add_write_index_release, 32
 hsa_queue_cas_write_index_acq_rel, 32
 hsa_queue_cas_write_index_acquire, 32
 hsa_queue_cas_write_index_relaxed, 32
 hsa_queue_cas_write_index_release, 32
 hsa_queue_create, 29
 hsa_queue_destroy, 29
 hsa_queue_inactivate, 30
 hsa_queue_load_read_index_acquire, 30
 hsa_queue_load_read_index_relaxed, 30
 hsa_queue_load_write_index_acquire, 31
 hsa_queue_load_write_index_relaxed, 31
 hsa_queue_store_read_index_relaxed, 33
 hsa_queue_store_read_index_release, 33
 hsa_queue_store_write_index_relaxed, 31
 hsa_queue_store_write_index_release, 31
 hsa_region_get_info, 41
 hsa_shut_down, 8
 hsa_signal_add_acq_rel, 20
 hsa_signal_add_acquire, 20
 hsa_signal_add_relaxed, 20
 hsa_signal_add_release, 20
 hsa_signal_and_acq_rel, 22
 hsa_signal_and_acquire, 22
 hsa_signal_and_relaxed, 22
 hsa_signal_and_release, 22
 hsa_signal_cas_acq_rel, 19
 hsa_signal_cas_acquire, 19
 hsa_signal_cas_relaxed, 19
 hsa_signal_cas_release, 19

- hsa_signal_create, 17
- hsa_signal_destroy, 17
- hsa_signal_exchange_acq_rel, 19
- hsa_signal_exchange_acquire, 19
- hsa_signal_exchange_relaxed, 19
- hsa_signal_exchange_release, 19
- hsa_signal_load_acquire, 18
- hsa_signal_load_relaxed, 18
- hsa_signal_or_acq_rel, 22
- hsa_signal_or_acquire, 22
- hsa_signal_or_relaxed, 22
- hsa_signal_or_release, 22
- hsa_signal_store_relaxed, 18
- hsa_signal_store_release, 18
- hsa_signal_subtract_acq_rel, 21
- hsa_signal_subtract_acquire, 21
- hsa_signal_subtract_relaxed, 21
- hsa_signal_subtract_release, 21
- hsa_signal_wait_acquire, 24
- hsa_signal_wait_relaxed, 24
- hsa_signal_xor_acq_rel, 23
- hsa_signal_xor_acquire, 23
- hsa_signal_xor_relaxed, 23
- hsa_signal_xor_release, 23
- hsa_status_string, 10
- hsa_vendor_extension_query, 46

Index - Extension APIs

hsa_ext_add_module, 70
 hsa_ext_define_agent_allocation_global_variable_address, 78
 hsa_ext_define_program_allocation_global_variable_address, 77
 hsa_ext_define_readonly_variable_address, 79
 hsa_ext_deserialize_finalization, 67
 hsa_ext_deserialize_program, 84
 hsa_ext_destroy_finalization, 66
 hsa_ext_finalize, 64
 hsa_ext_finalize_program, 71
 hsa_ext_image_clear, 99
 hsa_ext_image_copy, 98
 hsa_ext_image_create_handle, 95
 hsa_ext_image_destroy_handle, 100
 hsa_ext_image_export, 97
 hsa_ext_image_get_format_capability, 93
 hsa_ext_image_get_info, 94
 hsa_ext_image_import, 96
 hsa_ext_program_create, 69
 hsa_ext_program_destroy, 70
 hsa_ext_query_agent_global_variable_address, 79
 hsa_ext_query_call_convention, 76
 hsa_ext_query_indirect_function_descriptor_address, 81
 hsa_ext_query_kernel_descriptor_address, 81
 hsa_ext_query_program_agent_count, 73
 hsa_ext_query_program_agent_id, 73
 hsa_ext_query_program_agents, 74
 hsa_ext_query_program_allocation_global_variable_address, 78
 hsa_ext_query_program_brig_module, 75
 hsa_ext_query_program_module_count, 74
 hsa_ext_query_program_modules, 75
 hsa_ext_query_readonly_variable_address, 80
 hsa_ext_query_symbol_definition, 76
 hsa_ext_sampler_create_handle, 100
 hsa_ext_sampler_destroy_handle, 101
 hsa_ext_serialize_finalization, 66
 hsa_ext_serialize_program, 83
 hsa_ext_validate_program, 82
 hsa_ext_validate_program_module, 82

Bibliography

- [1] HSA Foundation. The HSA Programmer's Reference Manual. Technical report, HSA Foundation, 2013. 1, 16, 86
- [2] HSA Foundation. The HSA Platform System Architecture Specification. Technical report, HSA Foundation, 2014. 3, 4